

HW 1

Problem 1:

Part A:

- The code below is used to generate the make moons dataset:

```
X, y = generate_data()
plt.scatter(X[:, 0], X[:, 1], s=40, c=y, cmap=plt.cm.Spectral)
plt.show()

model = NeuralNetwork(nn_input_dim=2, nn_hidden_dim=3, nn_output_dim=2,
actFun_type='tanh')
model.fit_model(X,y)
model.visualize_decision_boundary(X,y)
```

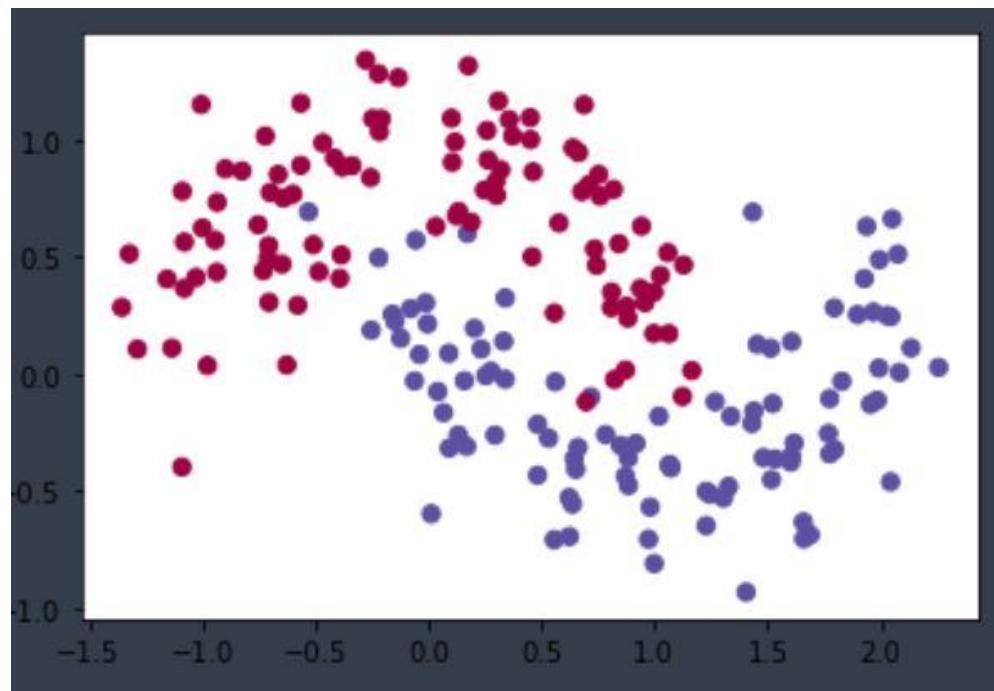


Figure 1: Make moons dataset figure.

Part B:

- The function below computes the activation function in *actFun(self, z, type)*:

```
def actFun(self, z, type):  
  
    if type == 'tanh':  
        z = (2/(1+np.exp(-2*z)))-1  
        #return np.tanh(z)  
        return z  
    elif type == 'sigmoid':  
        z = 1/(1+np.exp(-1*z))  
    elif type == 'relu':  
        return np.maximum(0, z)  
  
    else:  
        print("Incorrect function detected,\n" +  
              "please enter one of the following functions:\n" +  
              "ReLU, Sigmoid or Tanh")  
    return z
```

- The function below computes the activation function in `diff_actFun(self, z, type)`:

```
def diff_actFun(self, z, type):
    """
    diff_actFun compute the derivatives of the activation functions wrt the
    net input
    :param z: net input
    :param type: Tanh, Sigmoid, or ReLU
    :return: the derivatives of the activation functions wrt the net input
    """

    if type == 'tanh':
        z = (1/np.square(np.cosh(z))) #formula created using
        #tanh^2+sech^2=1
        #return 1 - np.square(np.tanh(z))
        #f = (2/(1+np.exp(-2*z)))-1
        #z = 1 - np.square(f)
        return z
    elif type == 'sigmoid':
        z = np.exp(-z)/(1 + 2*np.exp(-z) + np.exp(-2*z))
    elif type == 'relu':
        return np.where(z > 0, 1, 0)
```

- Derivative of Sigmoid is function is:

$$\frac{d}{dx} \left(\frac{1}{1 + e^{-x}} \right) = \frac{e^{-x}}{1 + 2e^{-x} + e^{-2x}}$$

- Derivative of Tanh is function is:

$$\frac{d}{dx} \left(\frac{2}{1 + e^{-x}} \right) = 1/\cosh^2(x)$$

- Derivative of ReLU is function is:

$$f(x) = 0 \text{ for } x < 0, x \text{ for } x > \text{ or equal to } 0$$

$$\frac{d}{dx}(f(x)) = 0 \text{ for } x < 0, 1 \text{ for } x > \text{ or equal to } 0$$

Part C:

Part C1:

- o The feedforward function is created using 4 formulas:
 - $z_1 \rightarrow \text{self.z1}$
 - $a_1 \rightarrow \text{self.a1}$
 - $z_2 \rightarrow \text{self.z2}$
 - $a_2 \rightarrow \text{self.prob}$

```
def feedforward(self, X, actFun):
    """
    feedforward builds a 3-layer neural network and computes the two
    probabilities,
        one for class 0 and one for class 1
    :param X: input data
    :param actFun: activation function
    :return:
    """

    # YOU IMPLEMENT YOUR feedforward HERE

    self.z1 = np.dot(X, self.W1) + self.b1
    self.a1 = actFun(self.z1)
    self.z2 = np.dot(self.a1, self.W2) + self.b2
```

```

#Version 1 of a2 = ^y = softmax(z2)
#self.probs = np.exp(self.z2) / np.sum(np.exp(self.z2),axis = 1,
keepdims = True)

#version 2 of a2 = ^y = softmax(z2)
A = np.exp(self.z2 - np.max(self.z2))
B = np.sum(A, axis = 1, keepdims = True)
result = A / B
self.probs = result

return None

```

Part C2:

- The function below calculates the loss function using the given **input data X** and **data labels Y** :

```

def calculate_loss(self, X, y):
    """
    calculate_loss compute the loss for prediction
    :param X: input data
    :param y: given labels
    :return: the loss for prediction
    """

    num_examples = len(X)
    self.feedforward(X, lambda x: self.actFun(x, type=self.actFun_type))

    # Calculating the loss
    yhat = self.probs #This is the same equation found feedforward function
    data_loss = np.sum(-np.log(yhat[range(num_examples), y]))

    # Add regularization term to loss (optional)
    data_loss += self.reg_lambda / 2 * (np.sum(np.square(self.W1)) +
    np.sum(np.square(self.W2)))
    return (1. / num_examples) * data_loss

```

Part D:

Part D1:

$$Delta3 = \frac{dL}{dz2} = \frac{dL}{dyhat} * \frac{dyhat}{dz2}$$

$$Delta 2 = differentiation [Delta3 * Transpose(W2)]$$

$$\frac{dL}{dW1} = \text{Transpose}(X) * \text{Delta2}$$

$$\frac{dL}{db1} = \sum \text{Delta2}$$

$$\frac{dL}{dW2} = \text{Transpose}(a1) * \text{Delta3}$$

$$\frac{dL}{db2} = \sum \text{Delta3}$$

Part D2:

- The code below implements the backpropagation:

```

def backprop(self, X, y):
    """
    backprop run backpropagation to compute the gradients used to update
    the parameters in the backward step
    :param X: input data
    :param y: given labels
    :return: dL/dW1, dL/b1, dL/dW2, dL/db2
    """

    # IMPLEMENT YOUR BACKPROP HERE

    num_examples = len(X)
    delta3 = self.probs
    delta3[range(num_examples), y] -= 1
    dW2 = (self.a1.T).dot(delta3)
    db2 = np.sum(delta3, axis=0, keepdims=True)
    delta2 = delta3.dot(self.W2.T) * self.diff_actFun(self.z1,
self.actFun_type)
    dW1 = np.dot(X.T, delta2)
    db1 = np.sum(delta2, axis=0)

    return dW1, dW2, db1, db2

```

Part E:

Part E1:

- The code below is using **tanh** function and the results are shown after the code below:

```
# # generate and visualize Make-Moons dataset
X, y = generate_data()
# plt.scatter(X[:, 0], X[:, 1], s=40, c=y, cmap=plt.cm.Spectral)
# plt.show()

model = NeuralNetwork(nn_input_dim=2, nn_hidden_dim=3,
nn_output_dim=2, actFun_type='tanh')
model.fit_model(X,y)
model.visualize_decision_boundary(X,y)
```

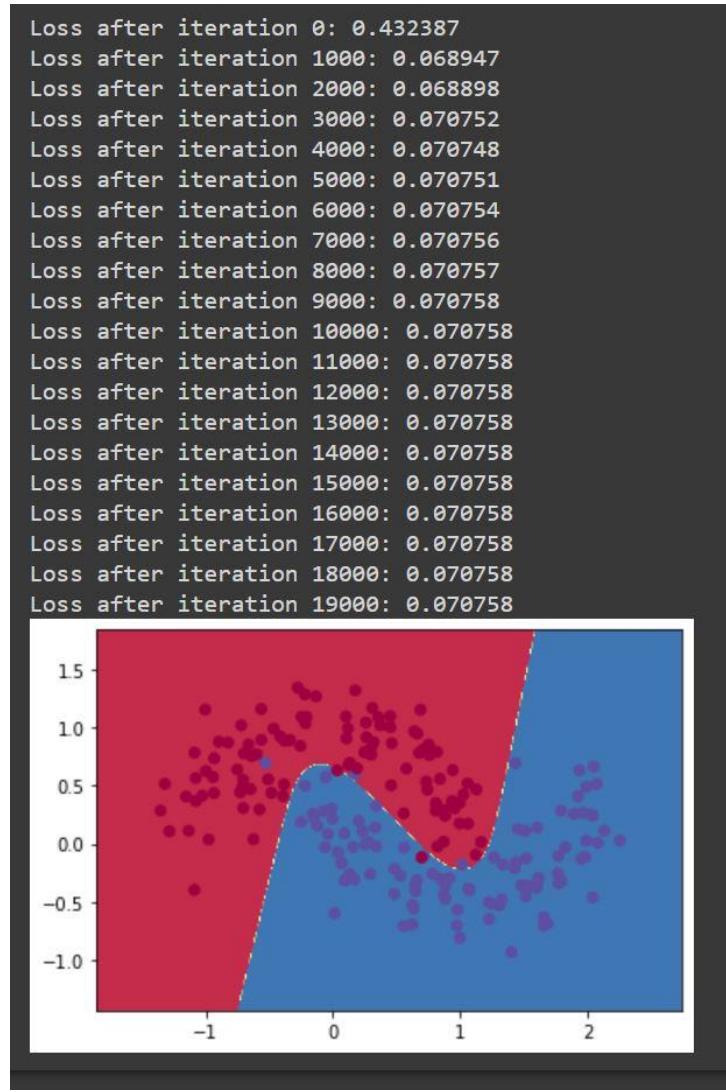


Figure 2: Training Results using **tanh** activation function.

- The code below is using **Sigmoid** function and the results are shown after the code below:

```
# # generate and visualize Make-Moons dataset
X, y = generate_data()
# plt.scatter(X[:, 0], X[:, 1], s=40, c=y, cmap=plt.cm.Spectral)
# plt.show()

model = NeuralNetwork(nn_input_dim=2, nn_hidden_dim=3 ,
nn_output_dim=2, actFun_type='sigmoid')
model.fit_model(X,y)
model.visualize_decision_boundary(X,y)
```

```

Loss after iteration 0: 0.628571
Loss after iteration 1000: 0.088431
Loss after iteration 2000: 0.079598
Loss after iteration 3000: 0.078604
Loss after iteration 4000: 0.078330
Loss after iteration 5000: 0.078233
Loss after iteration 6000: 0.078192
Loss after iteration 7000: 0.078174
Loss after iteration 8000: 0.078166
Loss after iteration 9000: 0.078161
Loss after iteration 10000: 0.078159
Loss after iteration 11000: 0.078158
Loss after iteration 12000: 0.078157
Loss after iteration 13000: 0.078156
Loss after iteration 14000: 0.078156
Loss after iteration 15000: 0.078156
Loss after iteration 16000: 0.078156
Loss after iteration 17000: 0.078156
Loss after iteration 18000: 0.078156
Loss after iteration 19000: 0.078155

```

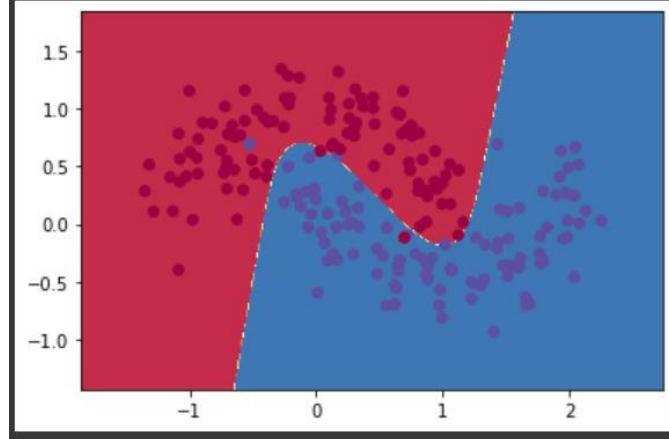


Figure 3: Training Results using **sigmoid** activation function.

- The code below is using **ReLU** function and the results are shown after the code below:

```

# # generate and visualize Make-Moons dataset
X, y = generate_data()
# plt.scatter(X[:, 0], X[:, 1], s=40, c=y, cmap=plt.cm.Spectral)
# plt.show()

model = NeuralNetwork(nn_input_dim=2, nn_hidden_dim=3 ,
nn_output_dim=2, actFun_type='relu')
model.fit_model(X,y)
model.visualize_decision_boundary(X,y)

```

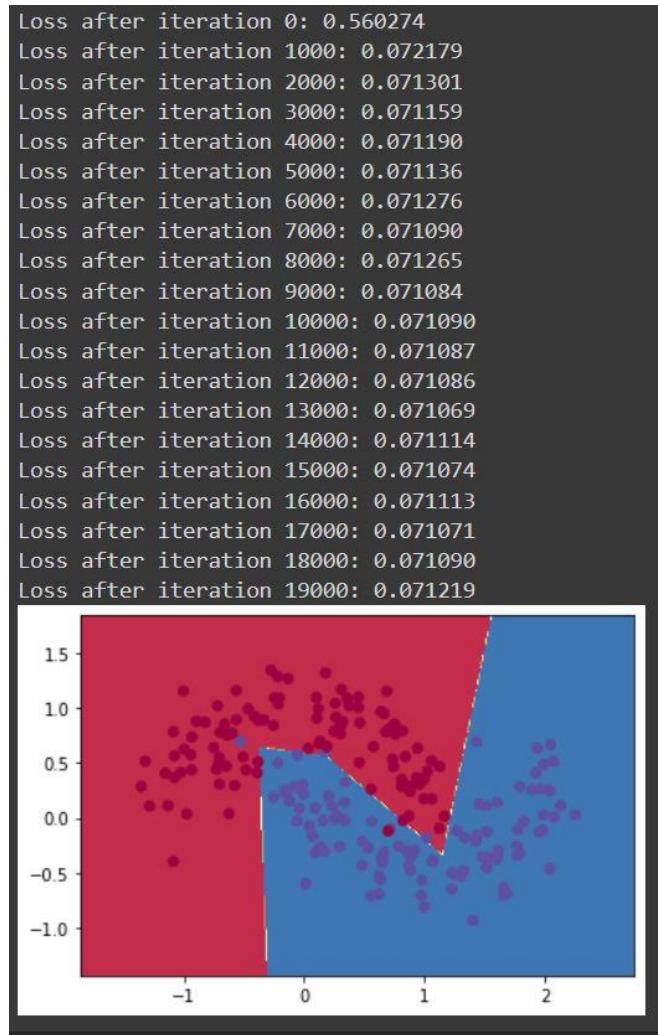


Figure 4: Training Results using **ReLU** activation function.

Part E2:

- Figures 6 to 19 shows training of the 3-layer neural network done using **tanh** activation function with various hidden number of layers.
- Increasing the hidden number of layers to 10 as shown in figure 6 our loss is 3.05% which is a significant decrease in loss compared to the 7.07% shown in figure 2.
- Increasing the hidden number of layers from 10 till 32 as shown in figures 7 till 13 it can be seen that the training loss plateaus at around 3.08%. If the hidden number of layers is increased from 32 to 50 it can be seen in figures 14 and 15

that the training loss actually increases to 3.1% which tells us that after a certain hidden number of layers is added the training loss effects our training model negatively.

- If we further increase the hidden number of layers from 100 to 400 it can be seen in figures 16 to 19 our loss is gradually increasing which confirms that adding more hidden number of layers after 10 actually effects our model negatively since the training loss increases rather than decreasing.

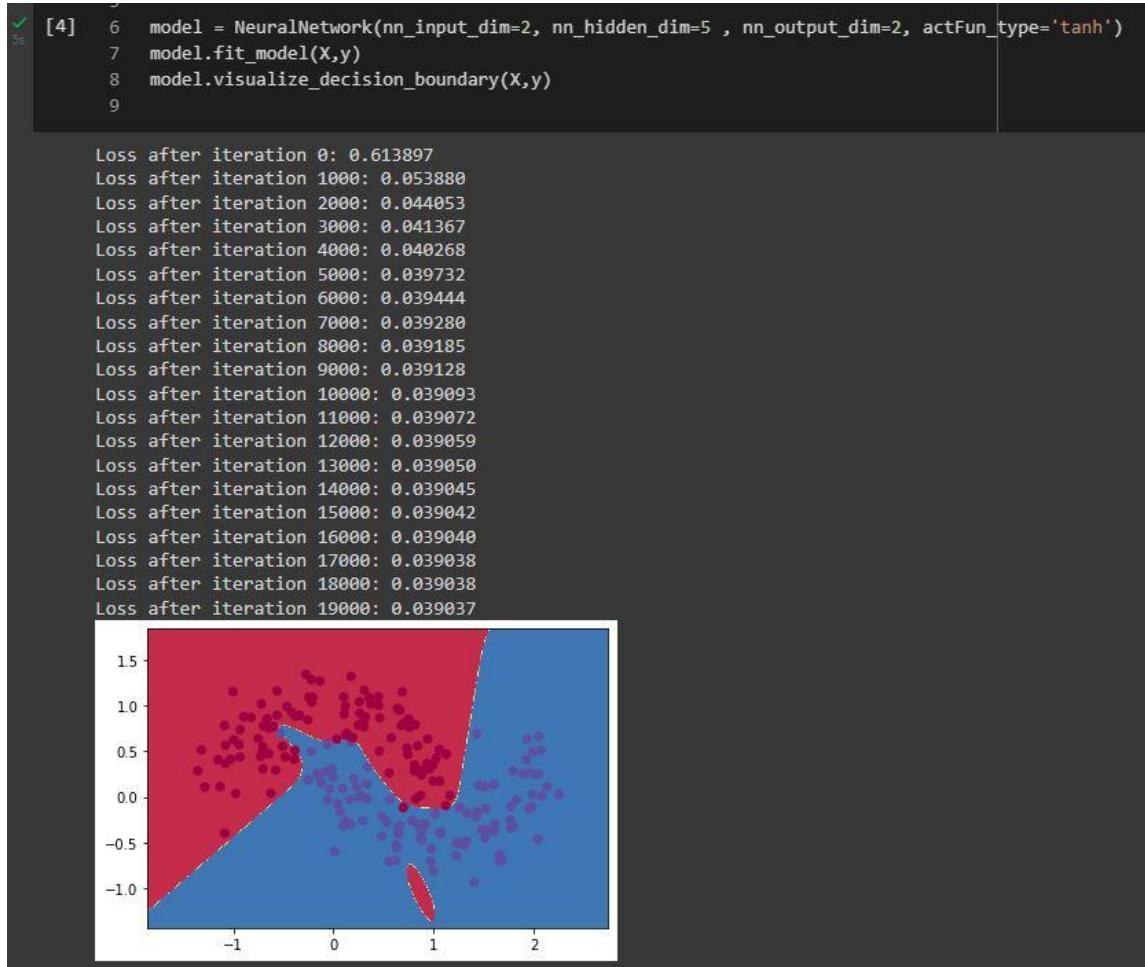


Figure 5: Training Results using **tanh** activation function after setting number of hidden units to 5.

```

6 model = NeuralNetwork(nn_input_dim=2, nn_hidden_dim=10 , nn_output_dim=2, actFun_type='tanh')
7 model.fit_model(X,y)
8 model.visualize_decision_boundary(X,y)
9

⇒ Loss after iteration 0: 0.634102
Loss after iteration 1000: 0.047377
Loss after iteration 2000: 0.038216
Loss after iteration 3000: 0.035216
Loss after iteration 4000: 0.033740
Loss after iteration 5000: 0.032800
Loss after iteration 6000: 0.032126
Loss after iteration 7000: 0.031624
Loss after iteration 8000: 0.031264
Loss after iteration 9000: 0.031044
Loss after iteration 10000: 0.030910
Loss after iteration 11000: 0.030821
Loss after iteration 12000: 0.030758
Loss after iteration 13000: 0.030711
Loss after iteration 14000: 0.030675
Loss after iteration 15000: 0.030646
Loss after iteration 16000: 0.030622
Loss after iteration 17000: 0.030601
Loss after iteration 18000: 0.030582
Loss after iteration 19000: 0.030565

```

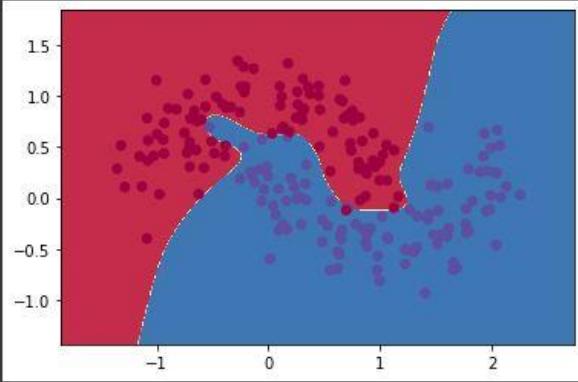


Figure 6: Training Results using **tanh** activation function after setting number of hidden units to **10**.

```

5
[25] 6 model = NeuralNetwork(nn_input_dim=2, nn_hidden_dim=12, nn_output_dim=2, actFun_type='tanh')
7 model.fit_model(X,y)
8 model.visualize_decision_boundary(X,y)
9

Loss after iteration 0: 0.677208
Loss after iteration 1000: 0.050805
Loss after iteration 2000: 0.043095
Loss after iteration 3000: 0.039945
Loss after iteration 4000: 0.038503
Loss after iteration 5000: 0.037486
Loss after iteration 6000: 0.036667
Loss after iteration 7000: 0.036000
Loss after iteration 8000: 0.035415
Loss after iteration 9000: 0.034876
Loss after iteration 10000: 0.034362
Loss after iteration 11000: 0.033855
Loss after iteration 12000: 0.033323
Loss after iteration 13000: 0.032796
Loss after iteration 14000: 0.032343
Loss after iteration 15000: 0.031943
Loss after iteration 16000: 0.031595
Loss after iteration 17000: 0.031299
Loss after iteration 18000: 0.031059
Loss after iteration 19000: 0.030888

```

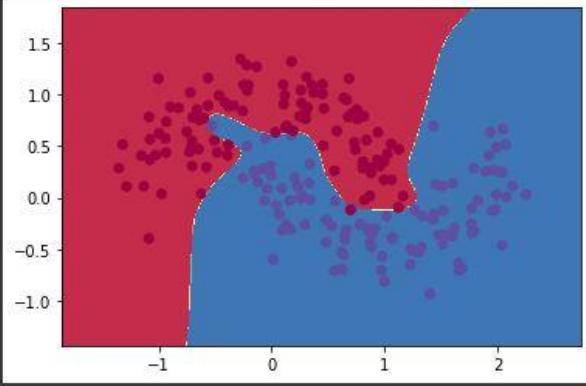


Figure 7: Training Results using **tanh** activation function after setting number of hidden units to **12**.

```

5
6 model = NeuralNetwork(nn_input_dim=2, nn_hidden_dim=16, nn_output_dim=2, actFun_type='tanh')
7 model.fit_model(X,y)
8 model.visualize_decision_boundary(X,y)
9

Loss after iteration 0: 0.671185
Loss after iteration 1000: 0.064950
Loss after iteration 2000: 0.041072
Loss after iteration 3000: 0.035949
Loss after iteration 4000: 0.034260
Loss after iteration 5000: 0.033236
Loss after iteration 6000: 0.032480
Loss after iteration 7000: 0.031927
Loss after iteration 8000: 0.031574
Loss after iteration 9000: 0.031353
Loss after iteration 10000: 0.031204
Loss after iteration 11000: 0.031094
Loss after iteration 12000: 0.031003
Loss after iteration 13000: 0.030922
Loss after iteration 14000: 0.030846
Loss after iteration 15000: 0.030772
Loss after iteration 16000: 0.030702
Loss after iteration 17000: 0.030637
Loss after iteration 18000: 0.030580
Loss after iteration 19000: 0.030530

```

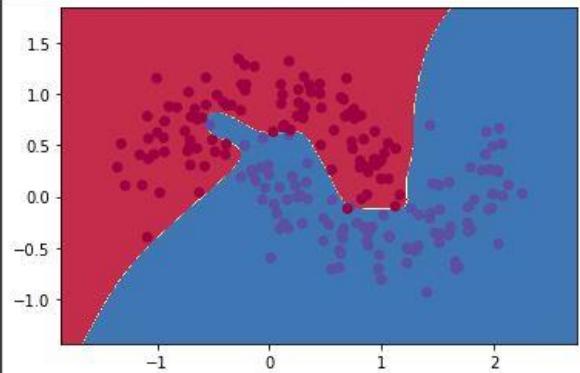


Figure 8: Training Results using **tanh** activation function after setting number of hidden units to **16**.

```

5
6 model = NeuralNetwork(nn_input_dim=2, nn_hidden_dim=18, nn_output_dim=2, actFun_type='tanh')
7 model.fit_model(X,y)
8 model.visualize_decision_boundary(X,y)
9

↳ Loss after iteration 0: 0.364404
Loss after iteration 1000: 0.052611
Loss after iteration 2000: 0.043724
Loss after iteration 3000: 0.038608
Loss after iteration 4000: 0.036049
Loss after iteration 5000: 0.034500
Loss after iteration 6000: 0.033374
Loss after iteration 7000: 0.032571
Loss after iteration 8000: 0.032102
Loss after iteration 9000: 0.031800
Loss after iteration 10000: 0.031587
Loss after iteration 11000: 0.031429
Loss after iteration 12000: 0.031308
Loss after iteration 13000: 0.031217
Loss after iteration 14000: 0.031144
Loss after iteration 15000: 0.031083
Loss after iteration 16000: 0.031030
Loss after iteration 17000: 0.030984
Loss after iteration 18000: 0.030943
Loss after iteration 19000: 0.030908



```

Figure 9: Training Results using **tanh** activation function after setting number of hidden units to **18**.

```
5
[28] 6   model = NeuralNetwork(nn_input_dim=2, nn_hidden_dim=20, nn_output_dim=2, actFun_type='tanh')
7   model.fit_model(X,y)
8   model.visualize_decision_boundary(X,y)
9

Loss after iteration 0: 0.498538
Loss after iteration 1000: 0.047806
Loss after iteration 2000: 0.039152
Loss after iteration 3000: 0.036454
Loss after iteration 4000: 0.035112
Loss after iteration 5000: 0.034247
Loss after iteration 6000: 0.033601
Loss after iteration 7000: 0.033072
Loss after iteration 8000: 0.032624
Loss after iteration 9000: 0.032252
Loss after iteration 10000: 0.031943
Loss after iteration 11000: 0.031683
Loss after iteration 12000: 0.031463
Loss after iteration 13000: 0.031286
Loss after iteration 14000: 0.031150
Loss after iteration 15000: 0.031039
Loss after iteration 16000: 0.030948
Loss after iteration 17000: 0.030877
Loss after iteration 18000: 0.030821
Loss after iteration 19000: 0.030774
```

Figure 10: Training Results using **tanh** activation function after setting number of hidden units to **20**.

```

5
[29] 6 model = NeuralNetwork(nn_input_dim=2, nn_hidden_dim=25, nn_output_dim=2, actFun_type='tanh')
7 model.fit_model(X,y)
8 model.visualize_decision_boundary(X,y)
9

Loss after iteration 0: 1.421211
Loss after iteration 1000: 0.052361
Loss after iteration 2000: 0.042873
Loss after iteration 3000: 0.038661
Loss after iteration 4000: 0.036657
Loss after iteration 5000: 0.035436
Loss after iteration 6000: 0.034610
Loss after iteration 7000: 0.033989
Loss after iteration 8000: 0.033462
Loss after iteration 9000: 0.032994
Loss after iteration 10000: 0.032576
Loss after iteration 11000: 0.032205
Loss after iteration 12000: 0.031881
Loss after iteration 13000: 0.031603
Loss after iteration 14000: 0.031374
Loss after iteration 15000: 0.031191
Loss after iteration 16000: 0.031044
Loss after iteration 17000: 0.030925
Loss after iteration 18000: 0.030827
Loss after iteration 19000: 0.030746


```

Figure 11: Training Results using **tanh** activation function after setting number of hidden units to **25**.

```
6 model = NeuralNetwork(nn_input_dim=2, nn_hidden_dim=30, nn_output_dim=2, actFun_type='tanh')
7 model.fit_model(X,y)
8 model.visualize_decision_boundary(X,y)
9
[1]: Loss after iteration 0: 0.345044
Loss after iteration 1000: 0.057101
Loss after iteration 2000: 0.050564
Loss after iteration 3000: 0.045179
Loss after iteration 4000: 0.040010
Loss after iteration 5000: 0.037333
Loss after iteration 6000: 0.035688
Loss after iteration 7000: 0.034540
Loss after iteration 8000: 0.033705
Loss after iteration 9000: 0.033048
Loss after iteration 10000: 0.032519
Loss after iteration 11000: 0.032088
Loss after iteration 12000: 0.031746
Loss after iteration 13000: 0.031495
Loss after iteration 14000: 0.031308
Loss after iteration 15000: 0.031168
Loss after iteration 16000: 0.031067
Loss after iteration 17000: 0.030991
Loss after iteration 18000: 0.030932
Loss after iteration 19000: 0.030884
```

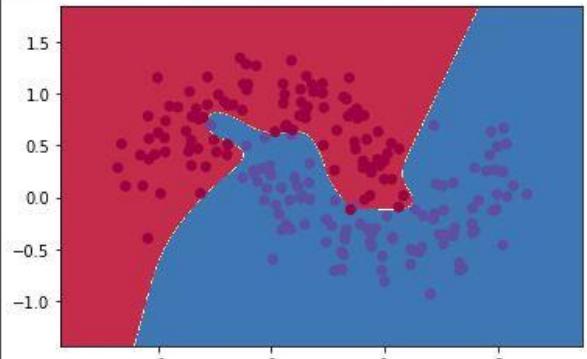


Figure 12: Training Results using **tanh** activation function after setting number of hidden units to **30**.

```
5
[31] 6 model = NeuralNetwork(nn_input_dim=2, nn_hidden_dim=32, nn_output_dim=2, actFun_type='tanh')
7 model.fit_model(X,y)
8 model.visualize_decision_boundary(X,y)
9
```

```
Loss after iteration 0: 1.349165
Loss after iteration 1000: 0.056592
Loss after iteration 2000: 0.050824
Loss after iteration 3000: 0.047725
Loss after iteration 4000: 0.043147
Loss after iteration 5000: 0.039042
Loss after iteration 6000: 0.036523
Loss after iteration 7000: 0.034945
Loss after iteration 8000: 0.033909
Loss after iteration 9000: 0.033184
Loss after iteration 10000: 0.032642
Loss after iteration 11000: 0.032215
Loss after iteration 12000: 0.031870
Loss after iteration 13000: 0.031600
Loss after iteration 14000: 0.031394
Loss after iteration 15000: 0.031237
Loss after iteration 16000: 0.031113
Loss after iteration 17000: 0.031012
Loss after iteration 18000: 0.030930
Loss after iteration 19000: 0.030862
```

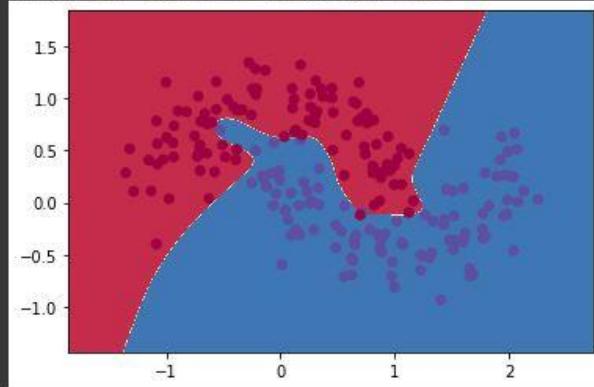


Figure 13: Training Results using **tanh** activation function after setting number of hidden units to **32**.

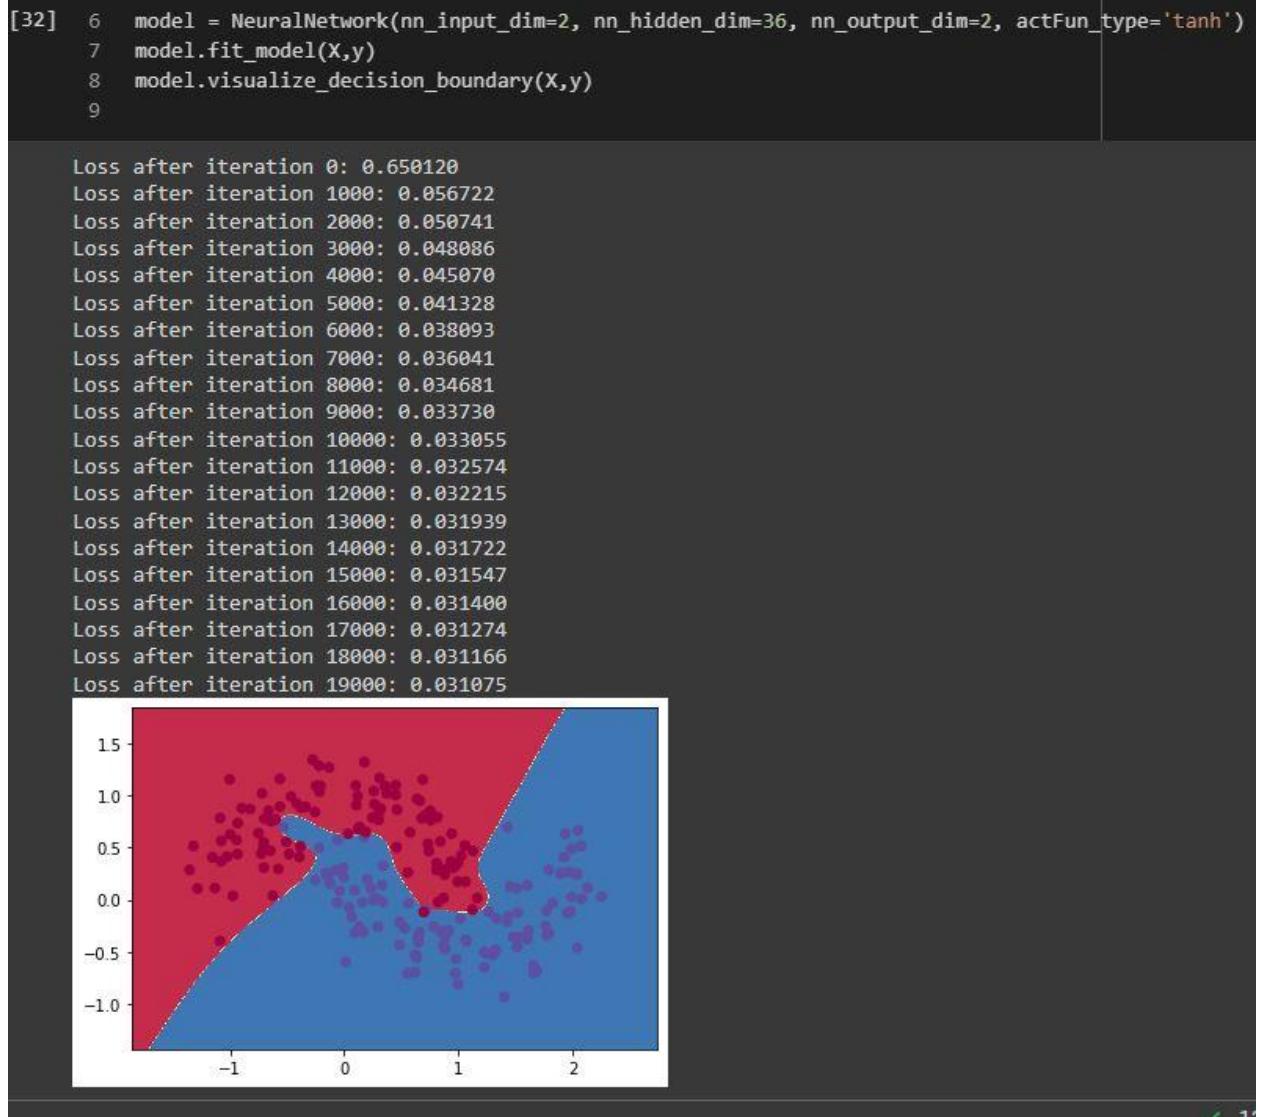


Figure 14: Training Results using **tanh** activation function after setting number of hidden units to **36**.

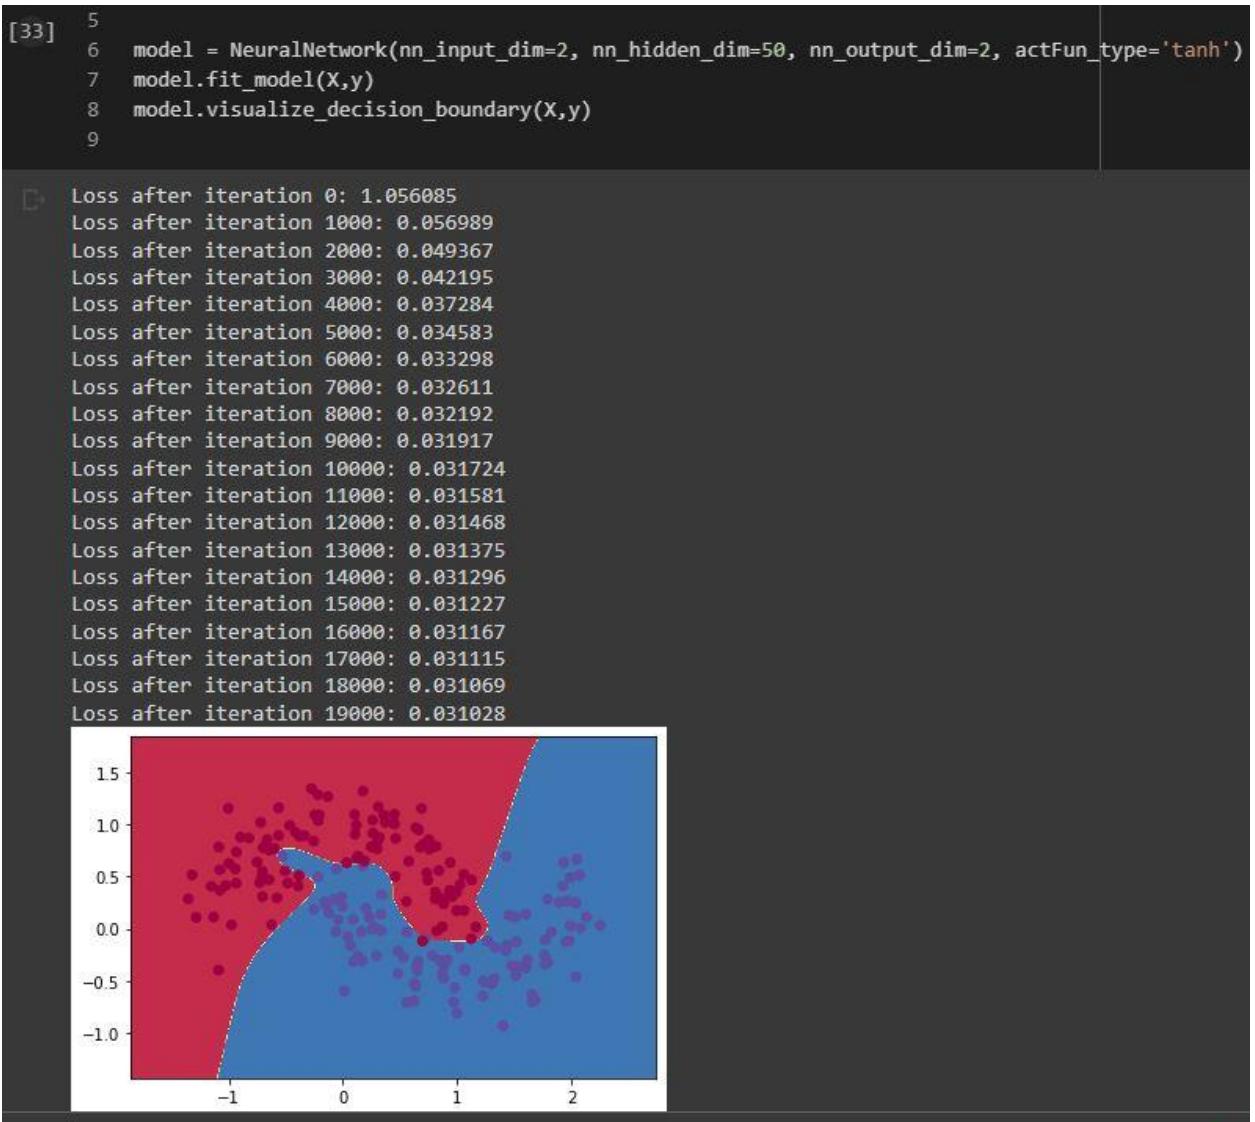


Figure 15: Training Results using **tanh** activation function after setting number of hidden units to **50**.

```

6 model = NeuralNetwork(nn_input_dim=2, nn_hidden_dim=100, nn_output_dim=2, actFun_type='tanh')
7 model.fit_model(X,y)
8 model.visualize_decision_boundary(X,y)
9
```

Loss after iteration 0: 3.337639
 Loss after iteration 1000: 0.055917
 Loss after iteration 2000: 0.051549
 Loss after iteration 3000: 0.049034
 Loss after iteration 4000: 0.044996
 Loss after iteration 5000: 0.041430
 Loss after iteration 6000: 0.039521
 Loss after iteration 7000: 0.038245
 Loss after iteration 8000: 0.037321
 Loss after iteration 9000: 0.036611
 Loss after iteration 10000: 0.036011
 Loss after iteration 11000: 0.035453
 Loss after iteration 12000: 0.034989
 Loss after iteration 13000: 0.034534
 Loss after iteration 14000: 0.034084
 Loss after iteration 15000: 0.033670
 Loss after iteration 16000: 0.033307
 Loss after iteration 17000: 0.032995
 Loss after iteration 18000: 0.032727
 Loss after iteration 19000: 0.032493

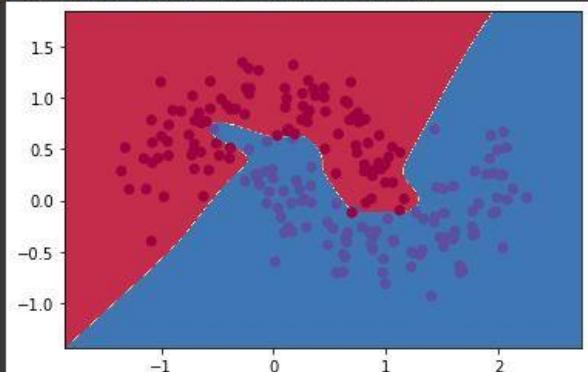


Figure 16: Training Results using **tanh** activation function after setting number of hidden units to **100**.

```
[35] 6 model = NeuralNetwork(nn_input_dim=2, nn_hidden_dim=200, nn_output_dim=2, actFun_type='tanh')
7 model.fit_model(x,y)
8 model.visualize_decision_boundary(x,y)
9
Loss after iteration 0: 5.105435
Loss after iteration 1000: 0.057131
Loss after iteration 2000: 0.052637
Loss after iteration 3000: 0.049063
Loss after iteration 4000: 0.045402
Loss after iteration 5000: 0.042689
Loss after iteration 6000: 0.040019
Loss after iteration 7000: 0.038934
Loss after iteration 8000: 0.038202
Loss after iteration 9000: 0.037683
Loss after iteration 10000: 0.037293
Loss after iteration 11000: 0.036985
Loss after iteration 12000: 0.036732
Loss after iteration 13000: 0.036520
Loss after iteration 14000: 0.036340
Loss after iteration 15000: 0.036186
Loss after iteration 16000: 0.036052
Loss after iteration 17000: 0.035935
Loss after iteration 18000: 0.035831
Loss after iteration 19000: 0.035737
```

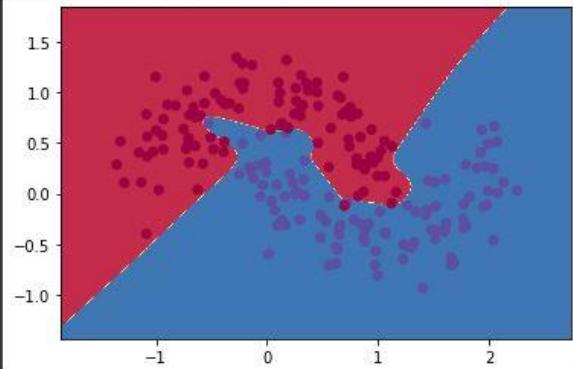


Figure 17: Training Results using **tanh** activation function after setting number of hidden units to **200**.

```

[36] 5
      model = NeuralNetwork(nn_input_dim=2, nn_hidden_dim=300, nn_output_dim=2, actFun_type='tanh')
      model.fit_model(X,y)
      model.visualize_decision_boundary(X,y)
      9

      Loss after iteration 0: 6.761681
      Loss after iteration 1000: 0.060224
      Loss after iteration 2000: 0.054406
      Loss after iteration 3000: 0.051136
      Loss after iteration 4000: 0.045940
      Loss after iteration 5000: 0.041287
      Loss after iteration 6000: 0.038355
      Loss after iteration 7000: 0.036679
      Loss after iteration 8000: 0.035665
      Loss after iteration 9000: 0.035061
      Loss after iteration 10000: 0.034640
      Loss after iteration 11000: 0.034312
      Loss after iteration 12000: 0.034022
      Loss after iteration 13000: 0.033738
      Loss after iteration 14000: 0.033453
      Loss after iteration 15000: 0.033173
      Loss after iteration 16000: 0.032914
      Loss after iteration 17000: 0.032686
      Loss after iteration 18000: 0.032496
      Loss after iteration 19000: 0.032340



```

Figure 18: Training Results using **tanh** activation function after setting number of hidden units to **300**.

```

[37] 5
      model = NeuralNetwork(nn_input_dim=2, nn_hidden_dim=400, nn_output_dim=2, actFun_type='tanh')
      6
      7
      8
      9
      model.fit_model(X,y)
      model.visualize_decision_boundary(X,y)

      ▾ Loss after iteration 0: 8.781713
      Loss after iteration 1000: 0.056969
      Loss after iteration 2000: 0.046617
      Loss after iteration 3000: 0.041765
      Loss after iteration 4000: 0.039756
      Loss after iteration 5000: 0.038676
      Loss after iteration 6000: 0.037889
      Loss after iteration 7000: 0.037243
      Loss after iteration 8000: 0.036733
      Loss after iteration 9000: 0.036258
      Loss after iteration 10000: 0.035815
      Loss after iteration 11000: 0.035457
      Loss after iteration 12000: 0.035185
      Loss after iteration 13000: 0.034978
      Loss after iteration 14000: 0.034820
      Loss after iteration 15000: 0.034698
      Loss after iteration 16000: 0.034601
      Loss after iteration 17000: 0.034530
      Loss after iteration 18000: 0.034470
      Loss after iteration 19000: 0.034385

```

Figure 19: Training Results using **tanh** activation function after setting number of hidden units to **400**.

Part F:

- The code below is the new class called **DeepNeuralNetwork** that inherits **Neural Network**:

```

class DeepNeuralNetwork(object):
    """
    This class builds and trains a deep neural network
    """

    def __init__(self, nn_dims, actFun_type='tanh', reg_lambda=0.01,
seed=0):

        self.nn_dims = nn_dims
        self.actFun_type = actFun_type
        self.reg_lambda = reg_lambda

        # initialize the weights and biases in the network
        self.W = []
        self.b = []

        np.random.seed(seed)
        for i in range(len(nn_dims)-1):
            self.W.append(np.random.randn(self.nn_dims[i],
self.nn_dims[i+1])
                           / np.sqrt(self.nn_dims[i]))
            self.b.append(np.zeros((1, self.nn_dims[i+1])))

    def actFun(self, z, type):

        if type == 'tanh':
            z = (2/(1+np.exp(-2*z)))-1
            #return np.tanh(z)
            return z
        elif type == 'sigmoid':
            z = 1/(1+np.exp(-1*z))
        elif type == 'relu':
            return np.maximum(0, z)

        else:
            print("Incorrect function detected,\n"
                  "please enter one of the following functions:\n"
                  "ReLU, Sigmoid or Tanh")
            return z

    def diff_actFun(self, z, type):

        if type == 'tanh':
            z = (1/np.square(np.cosh(z))) #formula created using
tanh^2+sech^2=1
            #return 1 - np.square(np.tanh(z))
            #f = (2/(1+np.exp(-2*z)))-1
            #z = 1 - np.square(f)
            return z
        elif type == 'sigmoid':
            z = np.exp(-z)/(1 + 2*np.exp(-z) + np.exp(-2*z))
        elif type == 'relu':
            return np.where(z > 0, 1, 0)

        else:
            print("Incorrect function detected,\n"
                  "please enter one of the following functions:\n"
                  "ReLU, Sigmoid or Tanh")
            return z

    def feedforward(self, X, actFun):

```

```

        self.a = []
        self.z = []
        for i in range(len(self.W)):
            if i == 0:
                self.z.append(np.dot(X, self.W[i]) + self.b[i])
            else:
                self.z.append(np.dot(self.a[i-1], self.W[i]) + self.b[i])
        if i != len(self.W) - 1:
            self.a.append(actFun(self.z[i]))
        C = np.exp(self.z[len(self.z)-1])
        self.probs = C / np.sum(C, axis=1, keepdims=True)
        return None

    def calculate_loss(self, X, y):
        num_examples = len(X)
        self.feedforward(X, lambda x: self.actFun(x,
type=self.actFun_type))
        # Calculating the loss

        probs = np.exp(self.z[len(self.z)-1]) / \
            np.sum(np.exp(self.z[len(self.z)-1])), axis=1,
keepdims=True)
        loss = -np.log(probs[range(num_examples), y])
        final_data_loss = np.sum(loss)

        # Add regularization term
        W_sum = 0
        for i in len(self.W):
            W_sum += np.sum(np.square(self.W[i]))
        final_data_loss += self.reg_lambda / 2 * W_sum
        return (1. / num_examples) * final_data_loss

    def predict(self, X):
        self.feedforward(X, lambda x: self.actFun(x,
type=self.actFun_type))
        return np.argmax(self.probs, axis=1)

    def backprop(self, X, y):

        num_examples = len(X)
        delta = self.probs
        delta[range(num_examples), y] -= 1

        db = []
        dW = []
        for num in range(len(self.z)):
            i = len(self.z) - num - 1
            if i != 0:
                dW.insert(0, np.dot(self.a[i - 1].T, delta))
                db.insert(0, np.sum(delta, axis=0, keepdims=True))
                delta = np.dot(delta, self.W[i].T) * \
                    self.diff_actFun(self.z[i-1],
type=self.actFun_type)
            else:
                db.insert(0, np.sum(delta, axis=0, keepdims=False))
                dW.insert(0, np.dot(X.T, delta))

        return dW, db

```

```

        def fit_model(self, X, y, epsilon=0.01, num_passes=20000,
print_loss=True):

            # Gradient descent.
            for i in range(0, num_passes):
                # Forward propagation
                self.feedforward(X, lambda x: self.actFun(x,
type=self.actFun_type))
                # Backpropagation
                dW, db = self.backprop(X, y)

                # Add regularization terms
                for i in range(len(dW)):
                    dW[i] += self.reg_lambda * self.W[i]

                # Gradient descent parameter update
                for i in range(len(self.W)):
                    self.W[i] += -epsilon * dW[i]
                    self.b[i] += -epsilon * db[i]

                # print the loss.
                if print_loss and i % 1000 == 0:
                    print("Loss after iteration %i: %f" % (i,
self.calculate_loss(X, y)))

            def visualize_decision_boundary(self, X, y):
                plot_decision_boundary(lambda x: self.predict(x), X, y)

```

- The code below is used to execute training for the **Make-Moons Dataset**:

```

def main():

    # # generate and visualize Make-Moons dataset
    X, y = generate_data()

    #Train the Neural Network Model
    model = DeepNeuralNetwork(nn_dims=[2,10,8], actFun_type='sigmoid')
    model.fit_model(X,y)
    model.visualize_decision_boundary(X,y)

if __name__ == "__main__":
    main()

```

- The code below is used to execute training for the **Make-Circles Dataset**:

```
def main():

    # # generate and visualize Make-Circles dataset
    X, y = datasets.make_circles(n_samples=100, shuffle=True, noise=None,
random_state=None, factor=0.8)

    #Train the Neural Network Model
    model = DeepNeuralNetwork(nn_dims=[2,10,8], actFun_type='sigmoid')
    model.fit_model(X,y)
    model.visualize_decision_boundary(X,y)

if __name__ == "__main__":
    main()
```

- Figures 20 to 28 show the **Make-Moons Dataset** using 3, 4 and 5 layer network. From figures 20 to 22 it can be seen that using relu activation function with 4 and 5 layer network makes a significant loss reduction and better data fit as compared to a 3 layer network. The sigmoid activation function in figures 23 to 25 don't show much of an improvement unless the hidden number of layers are increased significantly such as 100. The tanh function on the other hand shows much improvement and better fit even with 4 and 5 layer compared to relu and sigmoid functions and lower number of hidden layers as shown in figures 26 to 28.

```

✓ [75] 377      # # generate and visualize Make-Moons dataset
378      X, y = generate_data()
379
380      model = DeepNeuralNetwork(nn_dims=[2, 10, 2], actFun_type='relu')
381      model.fit_model(X,y)
382      model.visualize_decision_boundary(X,y)
383
384  if __name__ == "__main__":
385      main()

```

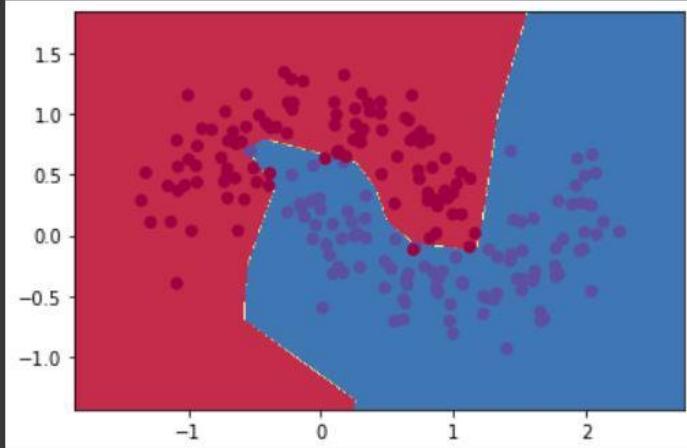


Figure 20: Training Results using **relu** activation function using 3 layer neural network.

```

375  def main():
376
377      # # generate and visualize Make-Moons dataset
378      X, y = generate_data()
379      # # generate and visualize Make-Circles dataset
380      #X, y = datasets.make_circles(n_samples=100, shuffle=True, noise=None, random_state=None, factor=0.8)
381      model = DeepNeuralNetwork(nn_dims=[2, 10, 4, 5], actFun_type='relu')
382      model.fit_model(X,y)
383      model.visualize_decision_boundary(X,y)
384
385  if __name__ == "__main__":
386      main()

```

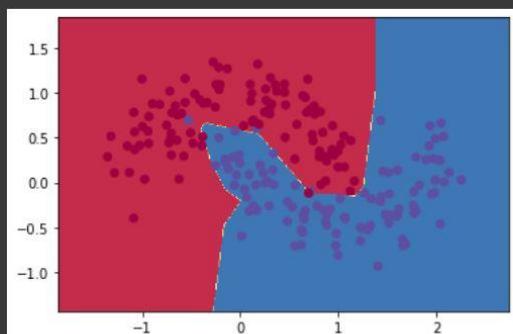


Figure 21: Training Results using **relu** activation function using 4 layer neural network.

```
381     model = DeepNeuralNetwork(nn_dims=[2, 10, 4, 8,8], actFun_type='relu')
382     model.fit_model(X,y)
383     model.visualize_decision_boundary(X,y)
384
385 if __name__ == "__main__":
386     main()
```

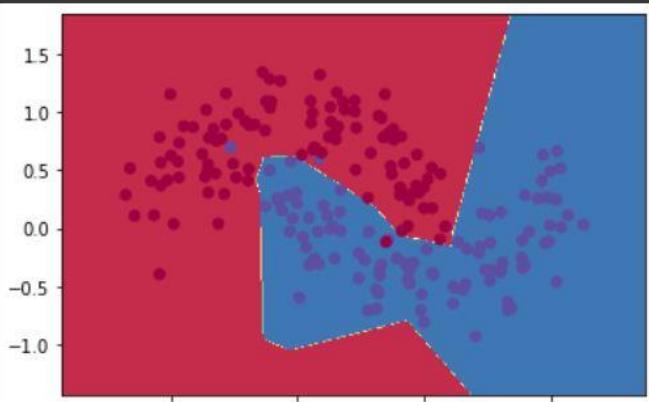


Figure 22: Training Results using **relu** activation function using 5 layer neural network.

```
375 < def main():
376
377     # # generate and visualize Make-Moons dataset
378     X, y = generate_data()
379
380     model = DeepNeuralNetwork(nn_dims=[2, 4, 2], actFun_type='sigmoid')
381     model.fit_model(X,y)
382     model.visualize_decision_boundary(X,y)
383
384 if __name__ == "__main__":
385     main()
```

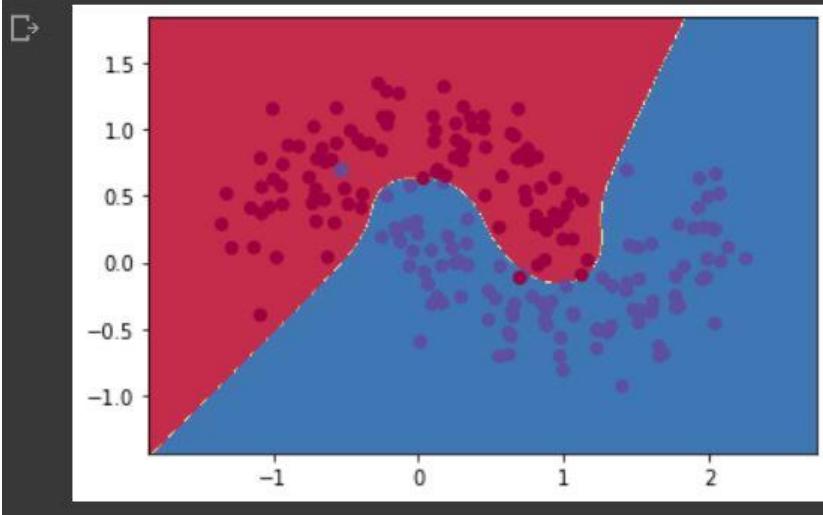


Figure 23: Training Results using **sigmoid** activation function using 3 layer neural network.

```

381     model = DeepNeuralNetwork(nn_dims=[2, 20, 8, 8], actFun_type='sigmoid')
382     model.fit_model(X,y)
383     model.visualize_decision_boundary(X,y)
384
385 if __name__ == "__main__":
386     main()

```

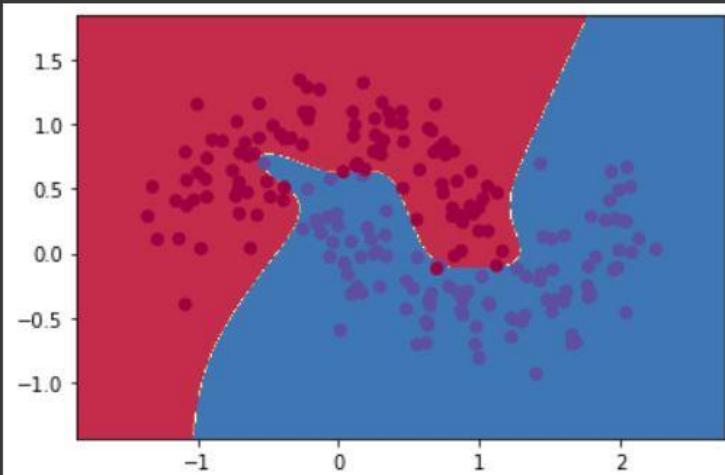


Figure 24: Training Results using **sigmoid** activation function using 4-layer neural network.

```

381     model = DeepNeuralNetwork(nn_dims=[2, 100, 8, 8, 10], actFun_type='sigmoid')
382     model.fit_model(X,y)
383     model.visualize_decision_boundary(X,y)
384
385 if __name__ == "__main__":
386     main()

```

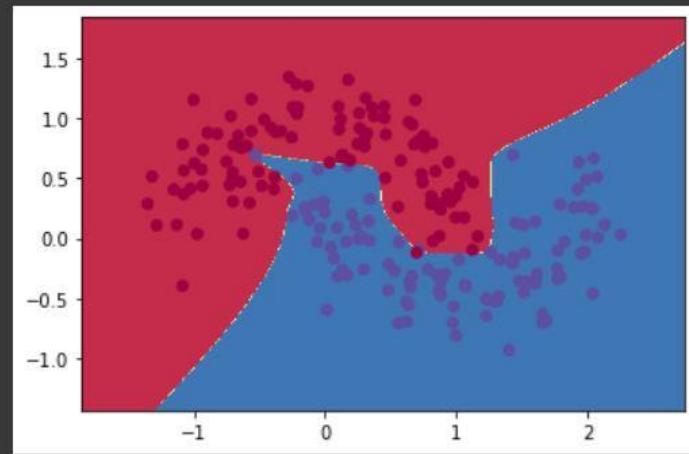


Figure 25: Training Results using **sigmoid** activation function using 5-layer neural network.

```
379 |     # ## generate and visualize Make-Moons dataset
380 |     X, y = generate_data()
381 |
382 |     model = DeepNeuralNetwork(nn_dims=[2, 10, 2], actFun_type='tanh')
383 |     model.fit_model(X,y)
384 |     model.visualize_decision_boundary(X,y)
385 |
386 | if __name__ == "__main__":
387 |     main()
```

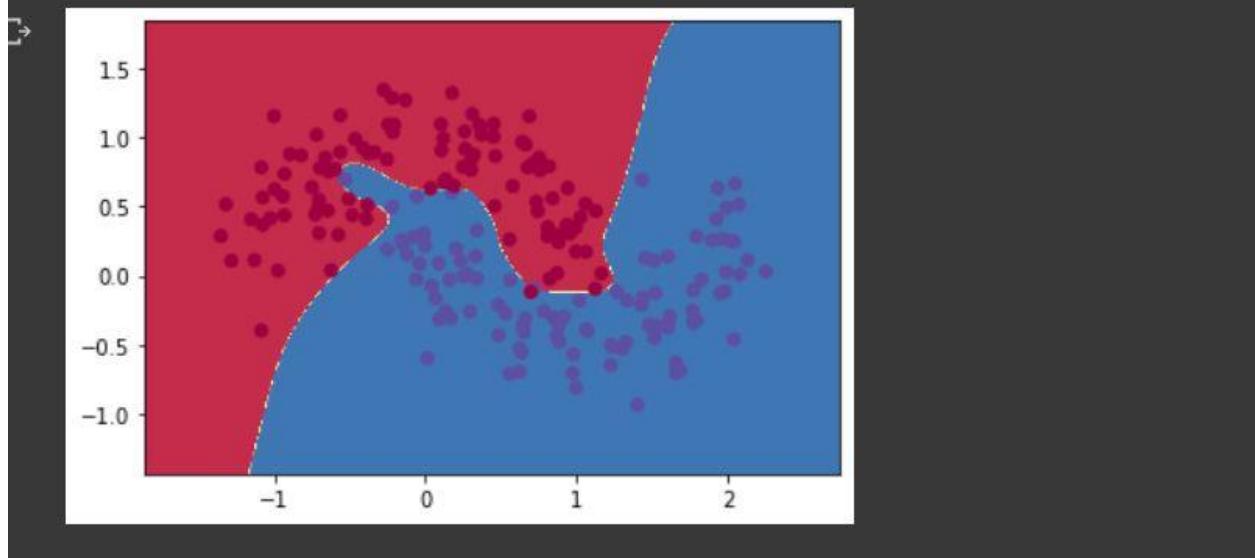


Figure 26: Training Results using **tanh** activation function using 3-layer neural network.

```
577     # # generate and visualize Make-Moons dataset
578     X, y = generate_data()
579
580     #model = DeepNeuralNetwork(nn_dims=[2, 100, 2], actFun_type='sigmoid')
581     model = DeepNeuralNetwork(nn_dims=[2, 4, 4, 4], actFun_type='tanh')
582     model.fit_model(X,y)
583     model.visualize_decision_boundary(X,y)
584
585 if __name__ == "__main__":
586     main()
```

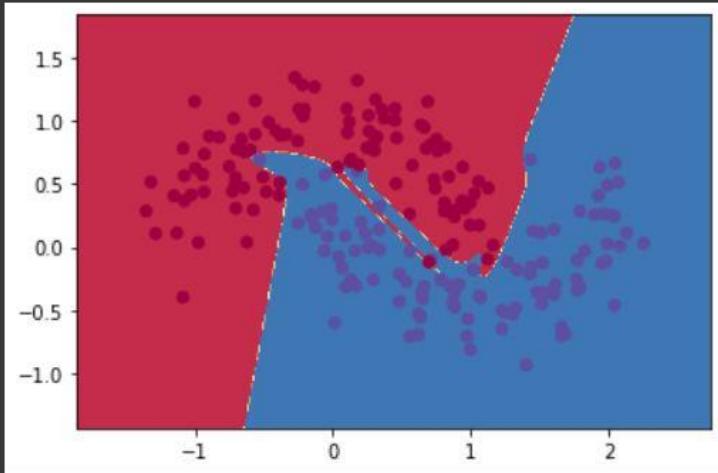


Figure 27: Training Results using **tanh** activation function using 4-layer neural network.

```

375  def main():
376
377      # # generate and visualize Make-Moons dataset
378      X, y = generate_data()
379
380      #model = DeepNeuralNetwork(nn_dims=[2, 100, 2], actFun_type='sigmoid')
381      model = DeepNeuralNetwork(nn_dims=[2, 10, 10, 10,10], actFun_type='tanh')
382      model.fit_model(X,y)
383      model.visualize_decision_boundary(X,y)
384
385  if __name__ == "__main__":
386      main()

```

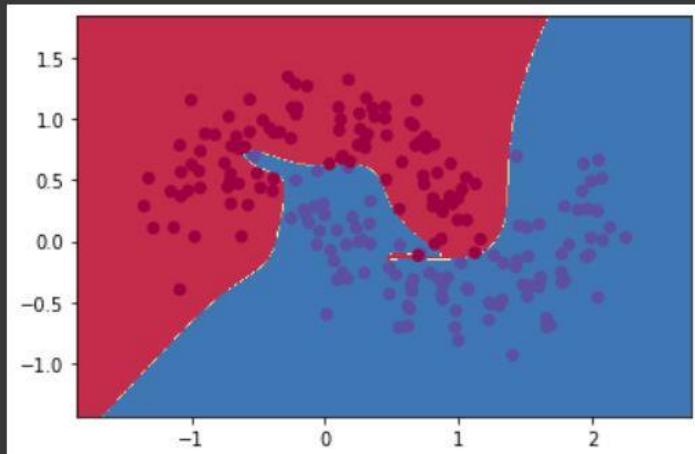


Figure 28: Training Results using **tanh** activation function using 5-layer neural network.

- Figures 29 to 34 show the **Make-Circles Dataset** using 3, 4 and 5 layer network. From figures 29 and 30 it can be seen that using relu activation function with 4 layer network has better data fit but doesn't do well when using a 5 layer network. The same behavior type can be observed with sigmoid activation function in figures 31 and 32 , with a 4 layer network the sigmoid function operates well but it doesn't do well with a 5 layer network. Figures 33 and 34 show the tanh function where using both 4 and 5 layer network not much improvement can be seen even when the hidden number of layers are increased to around 100 which is the same behavior also observed when tanh was used in a 3 layer network.

```

380 | X, y = datasets.make_circles(n_samples=100, shuffle=True, noise=None, random_state=None, factor=0.8)
381 | model = DeepNeuralNetwork(nn_dims=[2,100,150,150], actFun_type='relu')
382 | model.fit_model(X,y)
383 | model.visualize_decision_boundary(X,y)
384 |
385 ✓if __name__ == "__main__":
386     main()

```

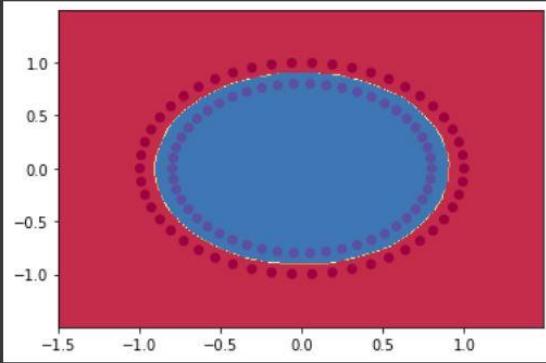


Figure 29: Training Results using **relu** activation function using 4-layer neural network.

```

380     X, y = datasets.make_circles(n_samples=100, shuffle=True, noise=None, random_state=None, factor=0.8)
381     model = DeepNeuralNetwork(nn_dims=[2,10,8,8,5], actFun_type='relu')
382     model.fit_model(X,y)
383     model.visualize_decision_boundary(X,y)
384
385 if __name__ == "__main__":
386     main()

```

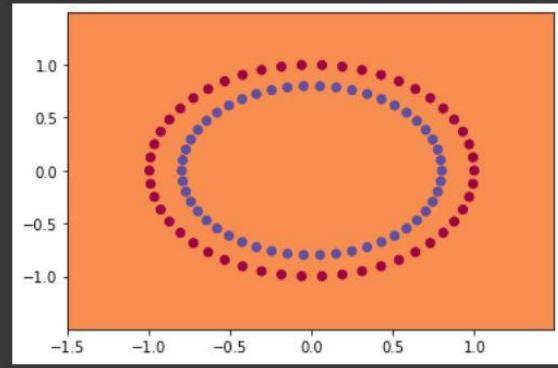


Figure 30: Training Results using **relu** activation function using 5-layer neural network.

```

380     X, y = datasets.make_circles(n_samples=100, shuffle=True, noise=None, random_state=None, factor=0.8)
381     model = DeepNeuralNetwork(nn_dims=[2,10,8,8], actFun_type='sigmoid')
382     model.fit_model(X,y)
383     model.visualize_decision_boundary(X,y)
384
385 if __name__ == "__main__":
386     main()

```

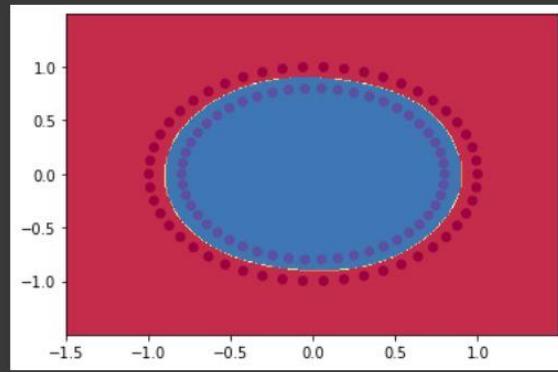


Figure 31: Training Results using **sigmoid** activation function using 4-layer neural network.

```

380     X, y = datasets.make_circles(n_samples=100, shuffle=True, noise=None, random_state=None, factor=0.8)
381     model = DeepNeuralNetwork(nn_dims=[2,10,8,8,5], actFun_type='sigmoid')
382     model.fit_model(X,y)
383     model.visualize_decision_boundary(X,y)
384
385 if __name__ == "__main__":
386     main()

```

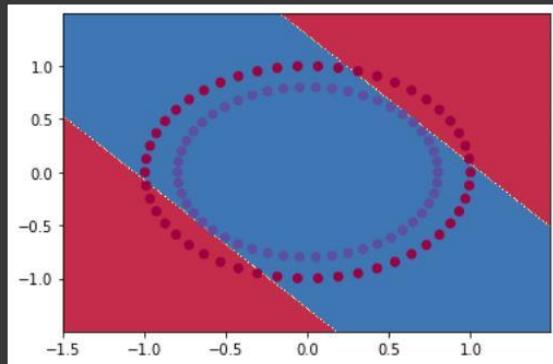


Figure 32: Training Results using **sigmoid** activation function using 5-layer neural network.

```

380     X, y = datasets.make_circles(n_samples=100, shuffle=True, noise=None, random_state=None, factor=0.8)
381     model = DeepNeuralNetwork(nn_dims=[2,100,20,20], actFun_type='tanh')
382     model.fit_model(X,y)
383     model.visualize_decision_boundary(X,y)
384
385 if __name__ == "__main__":
386     main()

```

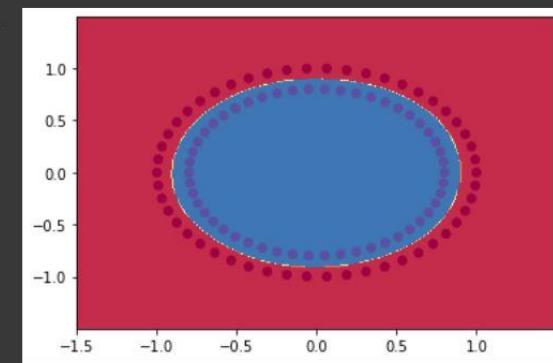


Figure 33: Training Results using **tanh** activation function using 4-layer neural network.

```

380     X, y = datasets.make_circles(n_samples=100, shuffle=True, noise=None, random_state=None, factor=0.8)
381     model = DeepNeuralNetwork(nn_dims=[2,10,8,8,8], actFun_type='tanh')
382     model.fit_model(X,y)
383     model.visualize_decision_boundary(X,y)
384
385 if __name__ == "__main__":
386     main()

```

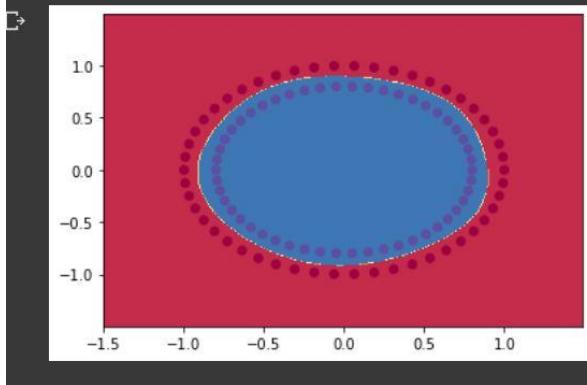


Figure 34: Training Results using **tanh** activation function using 5-layer neural network.

Problem 2:

Problem 2 Part A2:

- The code for ***weight_variable (shape)*** is shown below:

```

def weight_variable(shape):
    """
    Initialize weights
    :param shape: shape of weights, e.g. [w, h ,Cin, Cout] where
    w: width of the filters
    h: height of the filters
    Cin: the number of the channels of the filters
    Cout: the number of filters
    :return: a tensor variable for weights with initial values
    """

    # IMPLEMENT YOUR WEIGHT_VARIABLE HERE
    initial = tf.truncated_normal(shape, stddev=0.1)
    W = tf.Variable(initial)
    return W

```

- The code for ***bias_variable (shape)*** is shown below:

```

def bias_variable(shape):
    """
    Initialize biases
    """

```

```

:param shape: shape of biases, e.g. [Cout] where
Cout: the number of filters
:return: a tensor variable for biases with initial values
''

# IMPLEMENT YOUR BIAS VARIABLE HERE
initial = tf.constant(0.1, shape=shape)
b = tf.Variable(initial)
return b

```

- Th code for $\text{conv2d}(x, W)$ is shown below:

```

def conv2d(x, w):
    """
    Perform 2-D convolution
    :param x: input tensor of size [N, W, H, Cin] where
    N: the number of images
    W: width of images
    H: height of images
    Cin: the number of channels of images
    :param W: weight tensor [w, h, Cin, Cout]
    w: width of the filters
    h: height of the filters
    Cout: the number of the channels of the filters = the number of
    channels of images
    Cout: the number of filters
    :return: a tensor of features extracted by the filters, a.k.a. the
    results after convolution
    """

    # IMPLEMENT YOUR CONV2D HERE
    h_conv = tf.nn.conv2d(x, w, strides=[1, 1, 1, 1], padding='SAME')

    return h_conv

```

- Th code for $\text{max_pool_2x2}(x)$ is shown below:

```

def max_pool_2x2(x):
    """
    Perform non-overlapping 2-D maxpooling on 2x2 regions in the input
    data
    :param x: input data
    :return: the results of maxpooling (max-marginalized + downsampling)
    """

    # IMPLEMENT YOUR MAX_POOL_2X2 HERE
    h_max = tf.nn.max_pool(x, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1],
                           padding='SAME')

    return h_max

```

Problem 2 Part A3:

- Code for input data and input labels:

```
# placeholders for input data and input labels
x = tf.placeholder(tf.float32, [None, 784], name='x')
y_ = tf.placeholder(tf.float32, [None, 10], name='y_')

# Store Accuracies
val_accuracy_ = tf.placeholder(tf.float32, shape=())
test_accuracy_ = tf.placeholder(tf.float32, shape=())

# reshape the input image
x_image = tf.reshape(x, [-1, 28, 28, 1])
```

- Code for First Convolution layer:

```
# first convolutional layer
W_conv1 = weight_variable([5, 5, 1, 32])
b_conv1 = bias_variable([32])
input_1 = conv2d(x_image, W_conv1) + b_conv1
h_conv1 = tf.nn.relu(input_1)
h_pool1 = max_pool_2x2(h_conv1)

# second convolutional layer
W_conv2 = weight_variable([5, 5, 32, 64])
b_conv2 = bias_variable([64])
input_2 = conv2d(h_pool1, W_conv2) + b_conv2
h_conv2 = tf.nn.relu(input_2)
h_pool2 = max_pool_2x2(h_conv2)
```

- Code for densely connected layer:

```
# densely connected layer
W_fc1 = weight_variable([7 * 7 * 64, 1024])
b_fc1 = bias_variable([1024])
h_pool2_flat = tf.reshape(h_pool2, [-1, 7 * 7 * 64])
input_fc1 = tf.matmul(h_pool2_flat, W_fc1) + b_fc1
h_fc1 = tf.nn.relu(input_fc1)
```

- Code for Dropout Layer:

```
# dropout
keep_prob = tf.placeholder(tf.float32)
```

```
    h_fc1_drop = h_fc1_drop = tf.nn.dropout(h_fc1, keep_prob)
```

- Code for Softmax:

```
# softmax
W_fc2 = weight_variable([1024, 10])
b_fc2 = bias_variable([10])
y_conv = tf.nn.softmax(tf.matmul(h_fc1_drop, W_fc2) + b_fc2, name='y')
```

Problem 2 Part A4:

- Code for Training setup:

```
# setup training
cross_entropy = tf.reduce_mean(-tf.reduce_sum(y_ * tf.log(y_conv),
reduction_indices=[1]))
train_step = tf.train.AdamOptimizer(1e-4).minimize(cross_entropy)
correct_prediction = tf.equal(tf.argmax(y_conv, 1), tf.argmax(y_, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32),
name='accuracy')

# Add a scalar summary for the snapshot loss.
tf.summary.scalar(cross_entropy.op.name, cross_entropy)
# Build the summary operation based on the TF collection of Summaries.
# Layer-1
stats_summary("W_conv1", W_conv1)
stats_summary("b_conv1", b_conv1)
stats_summary("input_1", input_1)
stats_summary("h_conv1", h_conv1)
stats_summary("h_pool1", h_pool1)
# Layer-2
stats_summary("W_conv2", W_conv2)
stats_summary("b_conv2", b_conv2)
stats_summary("input_2", input_2)
stats_summary("h_conv2", h_conv2)
stats_summary("h_pool2", h_pool2)
# densely connected layer
stats_summary("W_fc1", W_fc1)
stats_summary("b_fc1", b_fc1)
stats_summary("h_pool2_flat", h_pool2_flat)
stats_summary("input_fc1", input_fc1)
stats_summary("h_fc1", h_fc1)
# Output layer - Softmax
stats_summary("W_fc2", W_fc2)
stats_summary("b_fc2", b_fc2)
stats_summary("y_conv", y_conv)

summary_op = tf.summary.merge_all()
summary_op_test = tf.summary.scalar('test_accuracy', test_accuracy_)
summary_op_val = tf.summary.scalar('val_accuracy', val_accuracy_)

# Add the variable initializer Op.
init = tf.initialize_all_variables()
```

```

# Create a saver for writing training checkpoints.
saver = tf.train.Saver()

# Instantiate a SummaryWriter to output summaries and the Graph.
summary_writer = tf.summary.FileWriter(result_dir, sess.graph)
# Val test
summary_writer_test = tf.summary.FileWriter(test_dir, sess.graph)
summary_writer_val = tf.summary.FileWriter(val_dir, sess.graph)

# Run the Op to initialize the variables.
sess.run(init)

```

Problem 2 Part A5:

- Code running the training:

```

# run the training
for i in range(max_step):
    batch = mnist.train.next_batch(50) # make the data batch, which is
    used in the training iteration.
    # the batch size is 50
    if i % 100 == 0:
        # output the training accuracy every 100 iterations
        train_accuracy = accuracy.eval(feed_dict={
            x: batch[0], y_: batch[1], keep_prob: 1.0})
        print("step %d, training accuracy %g" % (i, train_accuracy))

        # Update the events file which is used to monitor the training (in
        this case,
        # only the training loss is monitored)
        summary_str = sess.run(summary_op, feed_dict={x: batch[0], y_:
batch[1], keep_prob: 0.5})
        summary_writer.add_summary(summary_str, i)
        summary_writer.flush()

        # save the checkpoints every 1100 iterations
    if i % 1100 == 0 or i == max_step:
        checkpoint_file = os.path.join(result_dir, 'checkpoint')
        saver.save(sess, checkpoint_file, global_step=i)

        feed_dict_test = {x: mnist.test.images, y_: mnist.test.labels,
keep_prob: 1.0}
        test_accuracy = accuracy.eval(feed_dict=feed_dict_test)
        print('step %d, test accuracy %g' % (i, test_accuracy))
        summary_str_test = sess.run(summary_op_test,
feed_dict={test_accuracy_: test_accuracy})
        summary_writer_test.add_summary(summary_str_test, i)
        summary_writer_test.flush()

        feed_dict_val = {x: mnist.validation.images, y_:
mnist.validation.labels, keep_prob: 1.0}
        val_accuracy = accuracy.eval(feed_dict=feed_dict_val)
        print('step %d, validation accuracy %g' % (i, val_accuracy))
        summary_str_val = sess.run(summary_op_val,
feed_dict={val_accuracy_: val_accuracy})
        summary_writer_val.add_summary(summary_str_val, i)

```

```
summary_writer_val.flush()

train_step.run(feed_dict={x: batch[0], y_: batch[1], keep_prob: 0.5})
# run one train_step
```

- Below shows the results of the training:

The screenshot shows a Jupyter Notebook interface with a dark theme. On the left, there is a sidebar with various icons for file operations like running, saving, and deleting. The main area displays a series of log messages from a training process. The messages show training and validation accuracy over 3300 steps. The accuracy starts at 0.0928 and rises to 0.985, with test accuracy reaching 0.9779. The log ends with a warning about step 3300 being a test step.

```
Run: main X
step 0, validation accuracy 0.0928
step 100, training accuracy 0.82
step 200, training accuracy 0.94
step 300, training accuracy 0.94
step 400, training accuracy 0.88
step 500, training accuracy 0.96
step 600, training accuracy 1
step 700, training accuracy 0.96
step 800, training accuracy 0.96
step 900, training accuracy 0.98
step 1000, training accuracy 0.96
step 1100, training accuracy 0.98
step 1100, test accuracy 0.9697
step 1100, validation accuracy 0.9716
step 1200, training accuracy 1
step 1300, training accuracy 0.98
step 1400, training accuracy 0.98
step 1500, training accuracy 0.96
step 1600, training accuracy 1
step 1700, training accuracy 1
step 1800, training accuracy 0.94
step 1900, training accuracy 1
step 2000, training accuracy 0.98
step 2100, training accuracy 0.98
step 2200, training accuracy 1
step 2200, test accuracy 0.9779
step 2200, validation accuracy 0.9798
step 2300, training accuracy 0.94
step 2400, training accuracy 0.94
step 2500, training accuracy 1
step 2600, training accuracy 0.98
step 2700, training accuracy 0.98
step 2800, training accuracy 0.98
step 2900, training accuracy 0.98
step 3000, training accuracy 1
step 3100, training accuracy 0.98
step 3200, training accuracy 1
step 3300, training accuracy 1
step 3300 test accuracy 0.985
```

P Version Control Run Python Packages TODO Python Console

Packages installed successfully: Installed packages: 'tensorflow-datasets' (today 4:07 PM)

The screenshot shows a Jupyter Notebook interface with a purple header bar. The header bar has tabs for 'Run' (highlighted), 'main X', and other tabs that are mostly obscured by a dark overlay. On the far right of the header bar, there are icons for Version Control, Run, Python Packages, TODO, Python Console, and Help.

The main area of the interface is a dark-themed code editor. On the left side of the editor, there is a vertical toolbar with several icons: a green play button, an up arrow, a wrench, a down arrow, a square, a double-headed arrow, a printer, a target, and a trash can. The icon for the double-headed arrow is highlighted with a blue border.

The code editor contains the following text:

```
step 2600, training accuracy 0.98
step 2700, training accuracy 0.98
step 2800, training accuracy 0.98
step 2900, training accuracy 0.98
step 3000, training accuracy 1
step 3100, training accuracy 0.98
step 3200, training accuracy 1
step 3300, training accuracy 1
step 3300, test accuracy 0.985
step 3300, validation accuracy 0.9822
step 3400, training accuracy 1
step 3500, training accuracy 1
step 3600, training accuracy 0.98
step 3700, training accuracy 1
step 3800, training accuracy 0.98
step 3900, training accuracy 0.98
step 4000, training accuracy 0.92
step 4100, training accuracy 0.98
step 4200, training accuracy 0.98
step 4300, training accuracy 1
step 4400, training accuracy 1
step 4400, test accuracy 0.9849
step 4400, validation accuracy 0.986
step 4500, training accuracy 1
step 4600, training accuracy 1
step 4700, training accuracy 0.98
step 4800, training accuracy 0.98
step 4900, training accuracy 1
step 5000, training accuracy 1
step 5100, training accuracy 1
step 5200, training accuracy 1
step 5300, training accuracy 1
step 5400, training accuracy 0.98
validation accuracy 0.988
test accuracy 0.9874
The training takes 399.980971 second to finish

Process finished with exit code 0
```

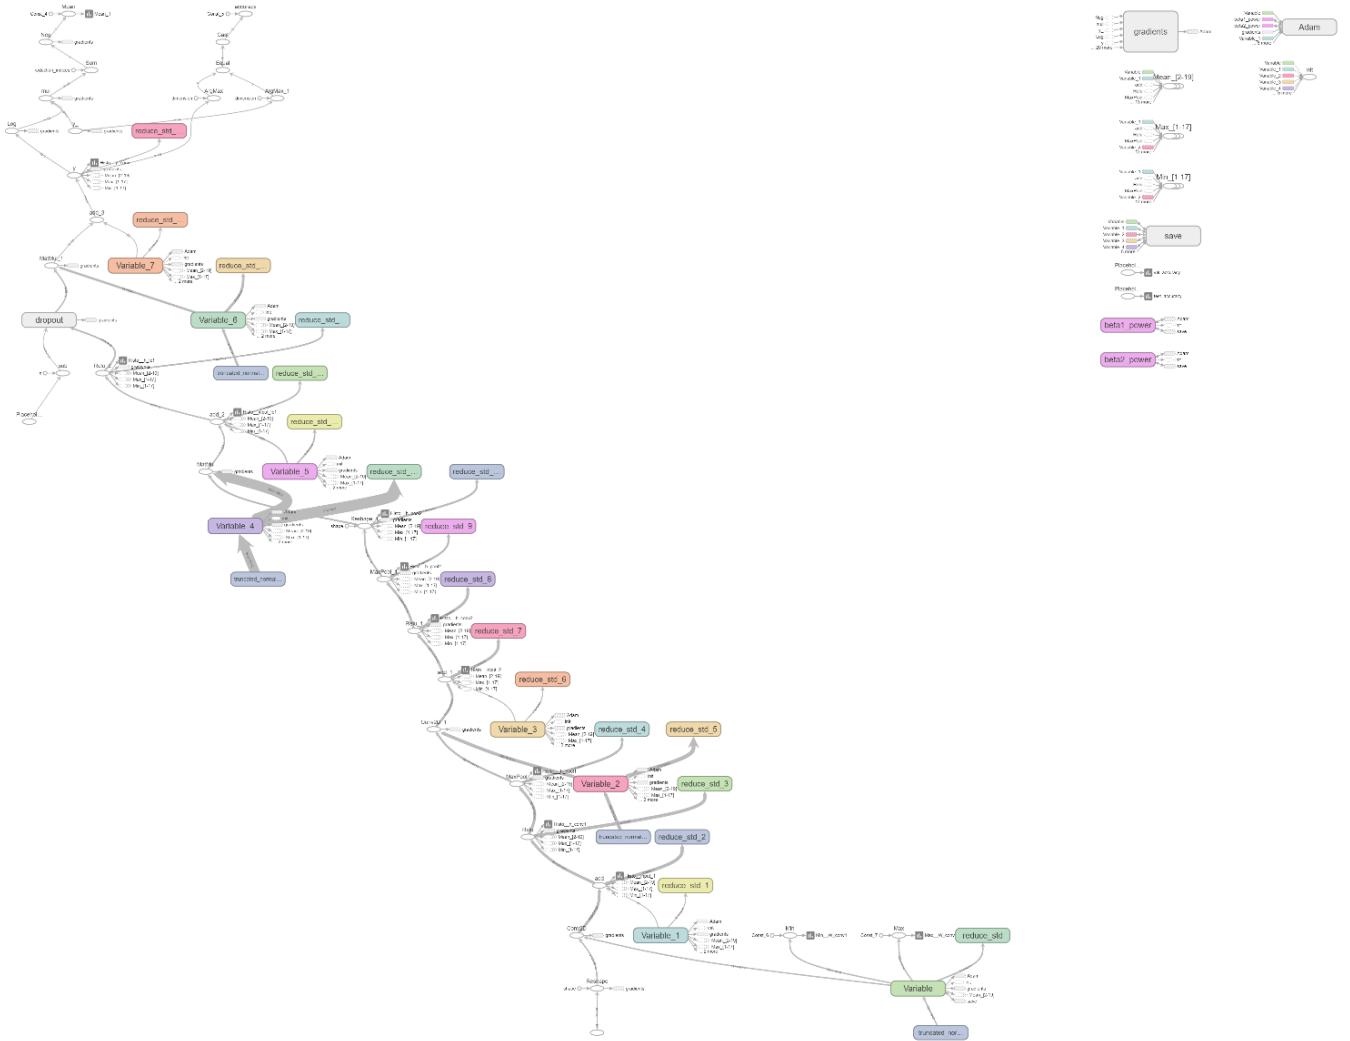
Below the code editor, there is a status bar with several icons: Version Control, Run, Python Packages, TODO, Python Console, and Help. There is also a message indicating that packages were installed successfully: 'Packages installed successfully: Installed packages: 'tensorflow-datasets' (today 4:07 PM)'.

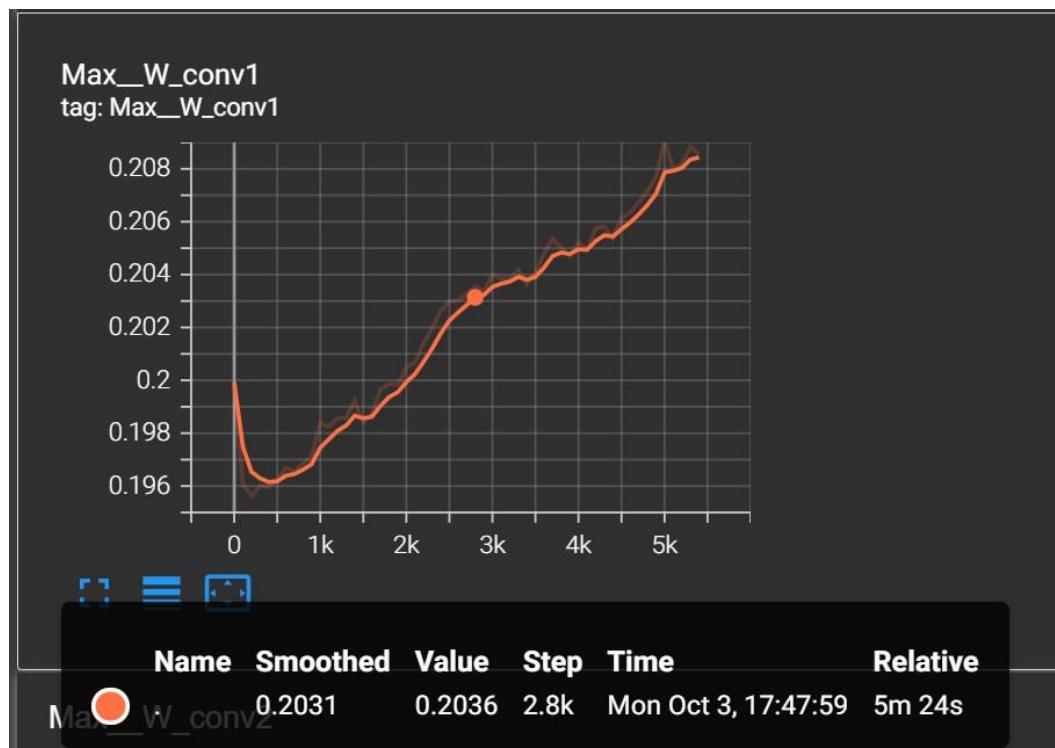
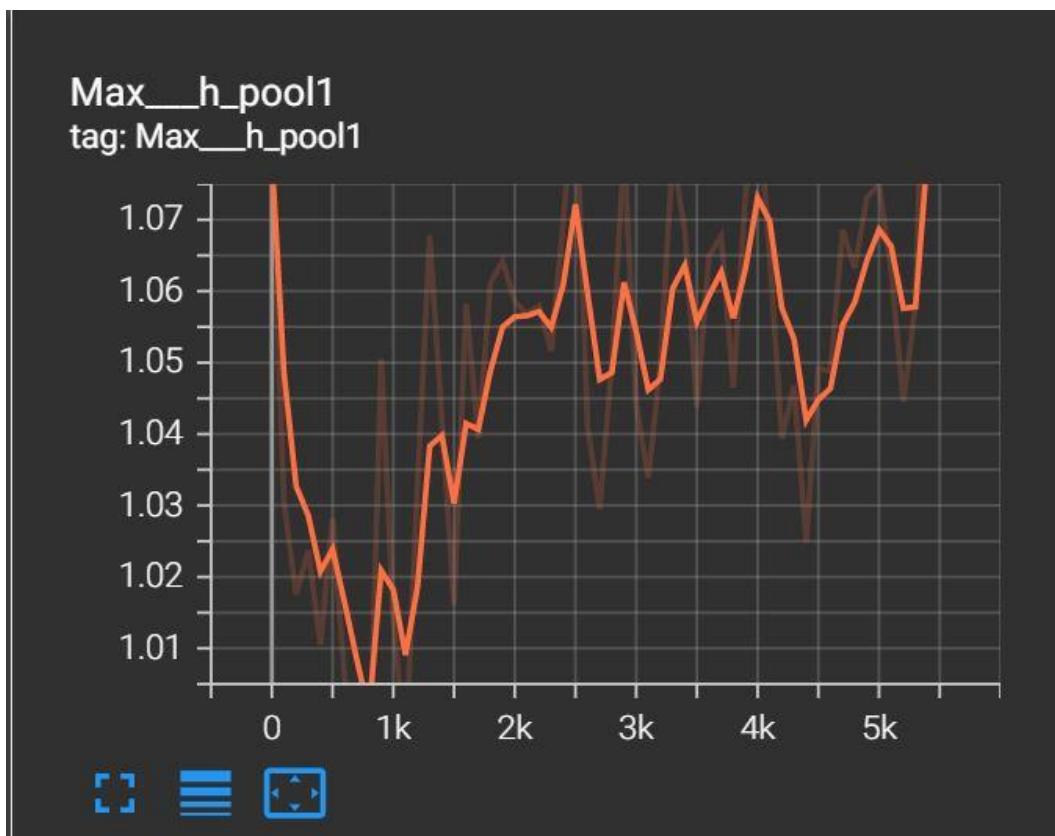
- Based on the figure above it can be seen that the training was completed after 5500 steps at **98.74% test accuracy** and it took 400 seconds to complete the training.

Problem 2 Part A6:

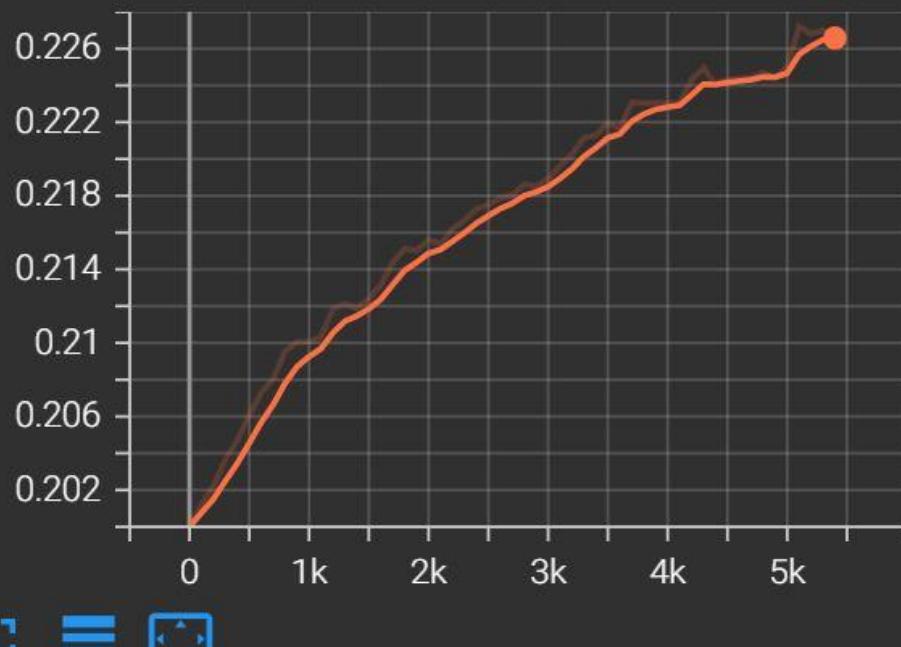
- The images below are results from training using **tensorboard**. The command used to run tensorboard and then accessed at <http://localhost:6006> :

```
tensorboard --logdir=results --bind_all
```

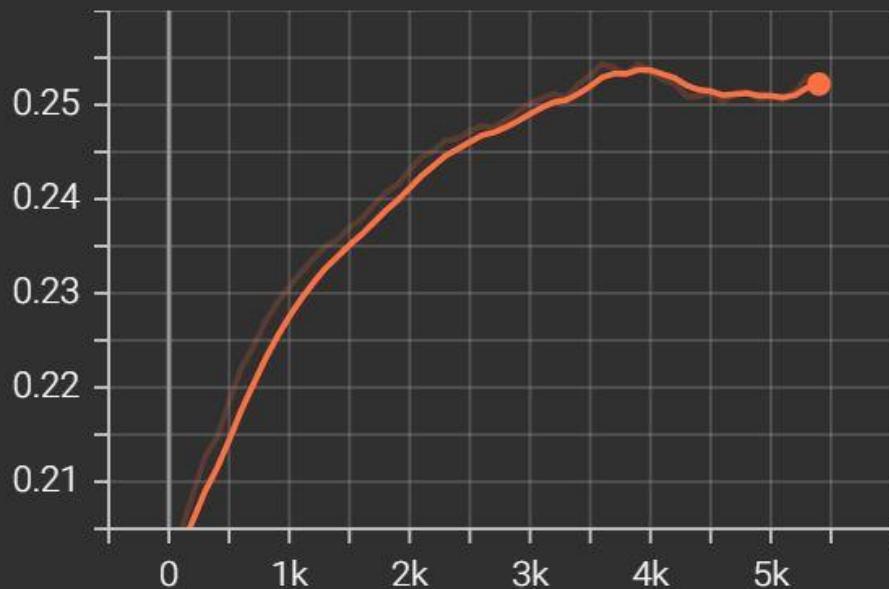




Max_W_conv2
tag: Max_W_conv2



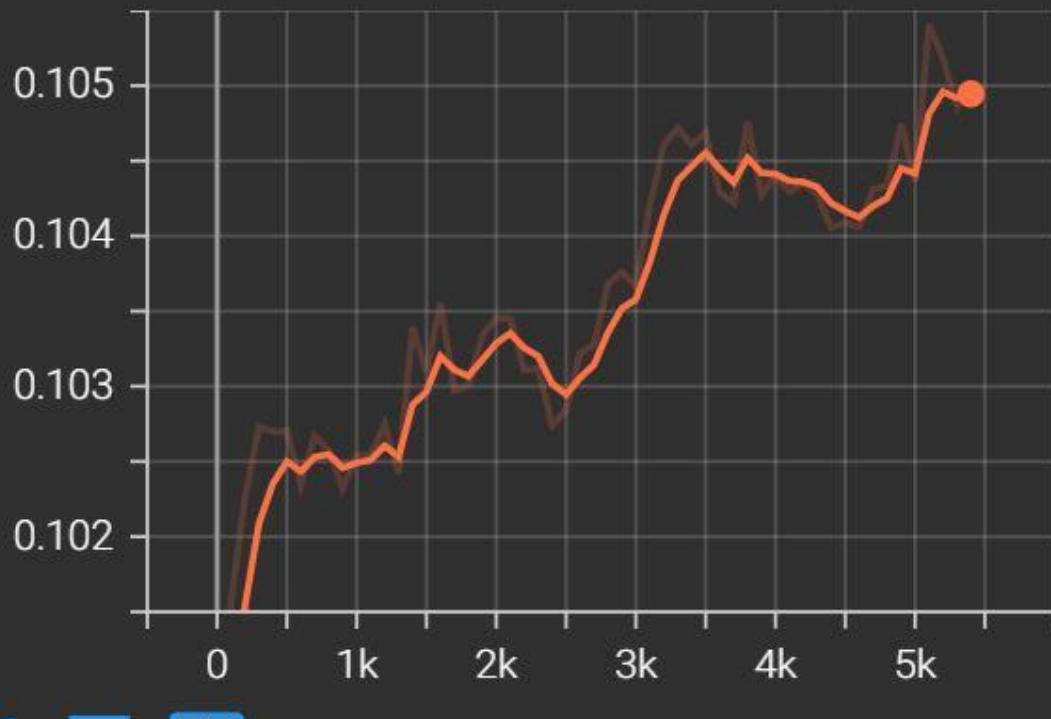
Max_W_fc1
tag: Max_W_fc1



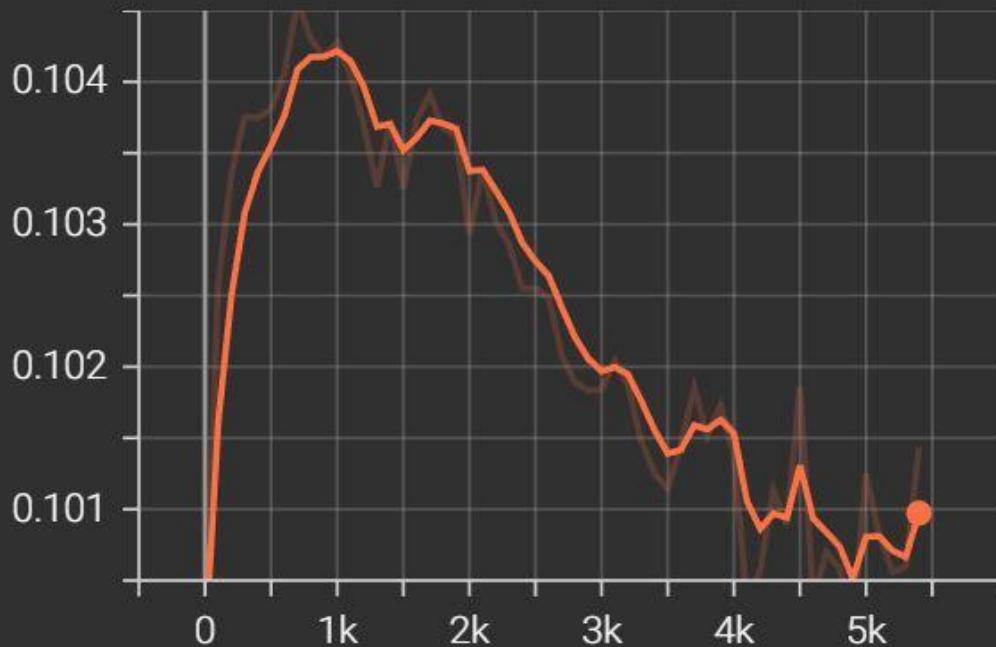
Max_W_fc2
tag: Max_W_fc2

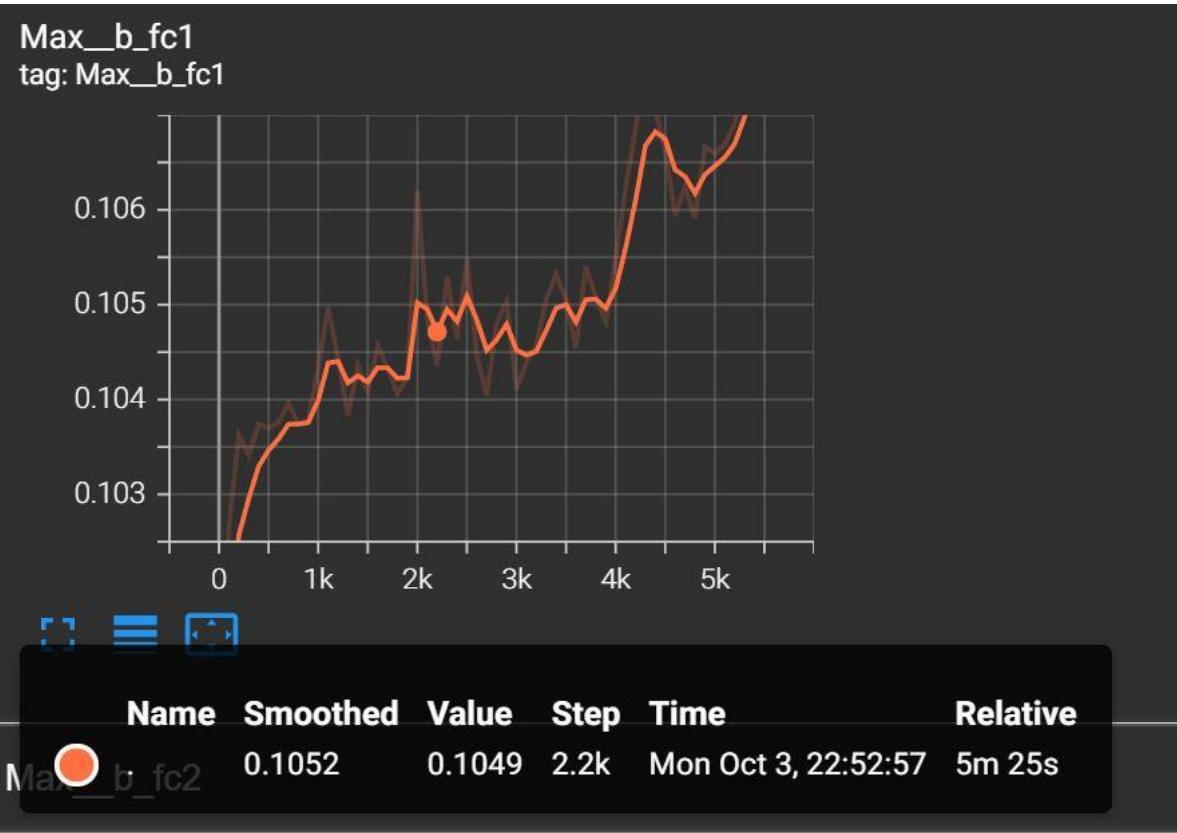


Max_b_conv1
tag: Max_b_conv1

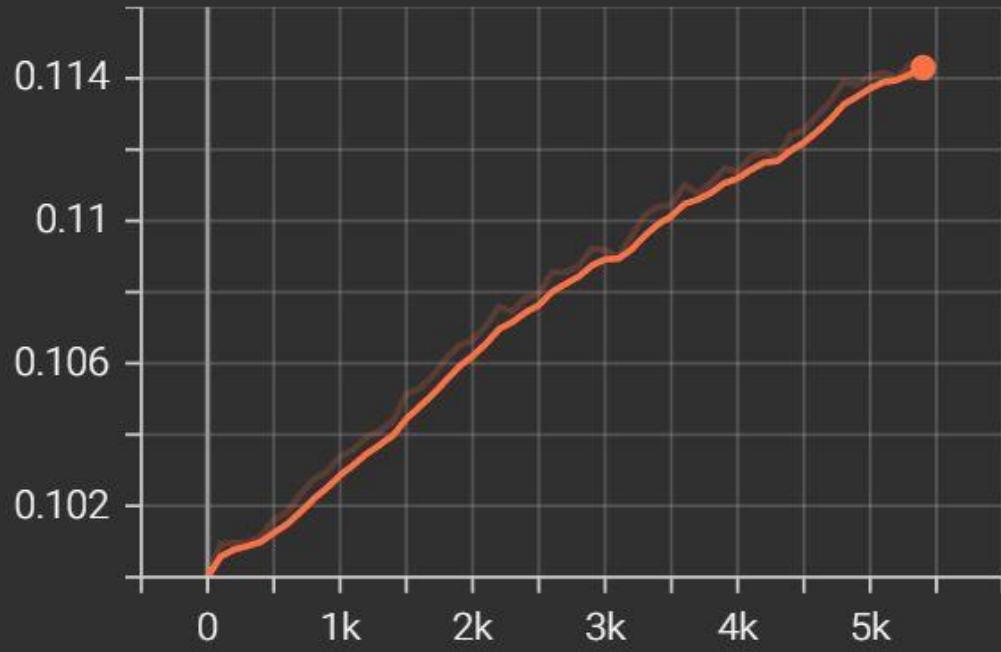


Max_b_conv2
tag: Max_b_conv2

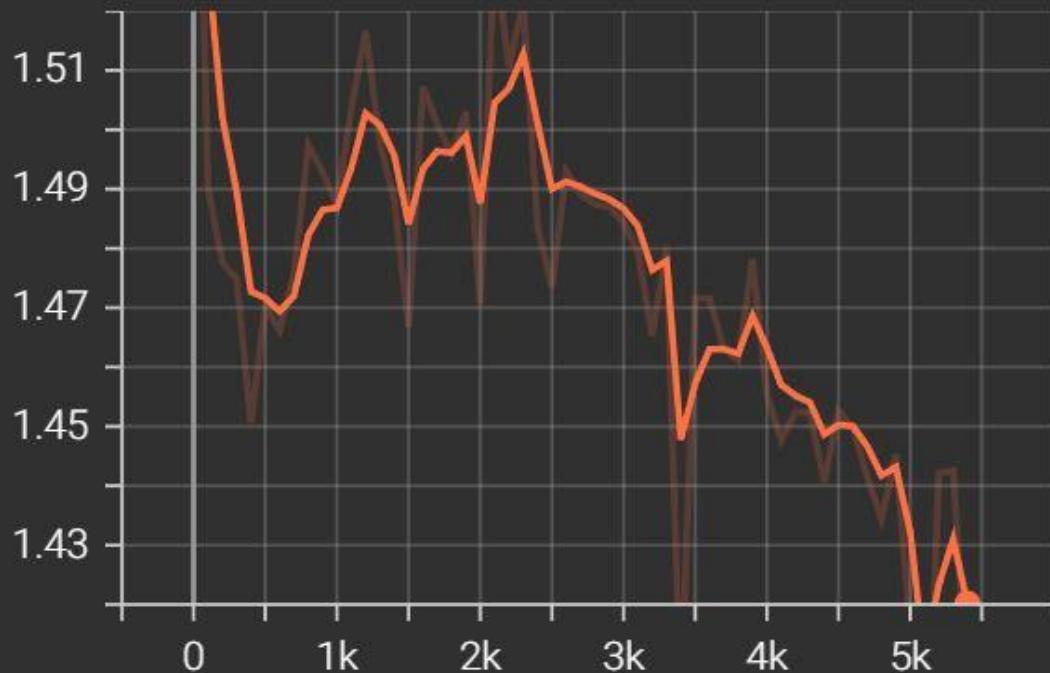




Max_b_fc2
tag: Max_b_fc2



Max_h_conv1
tag: Max_h_conv1

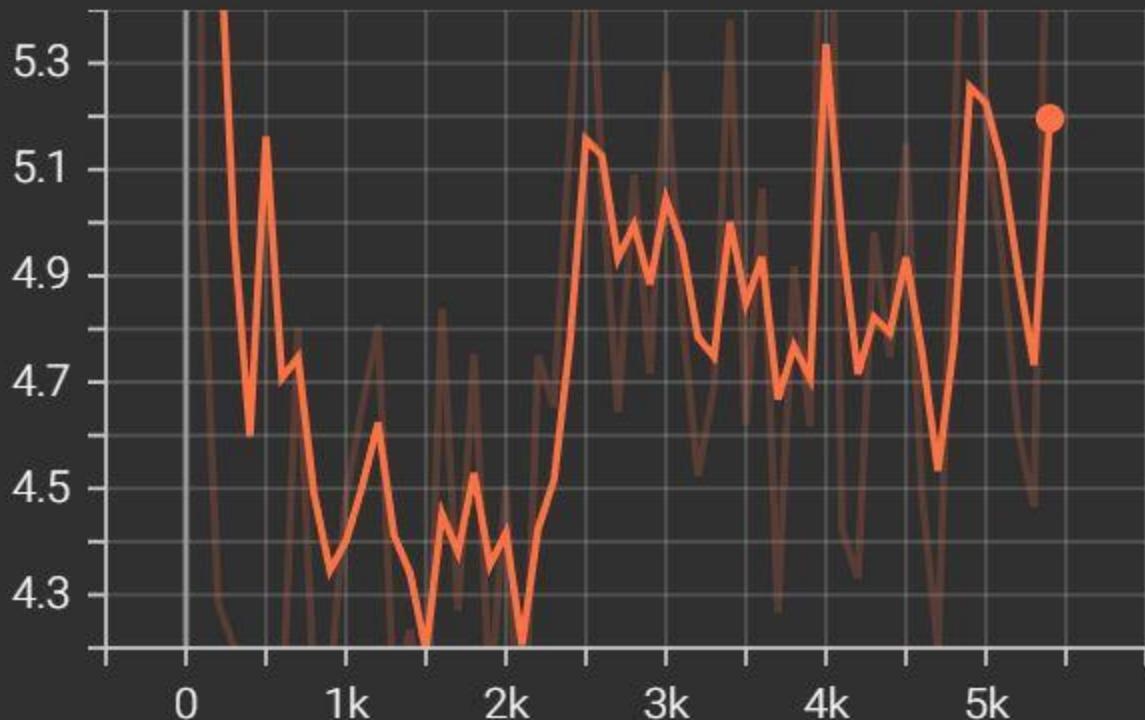


Max_h_conv2
tag: Max_h_conv2



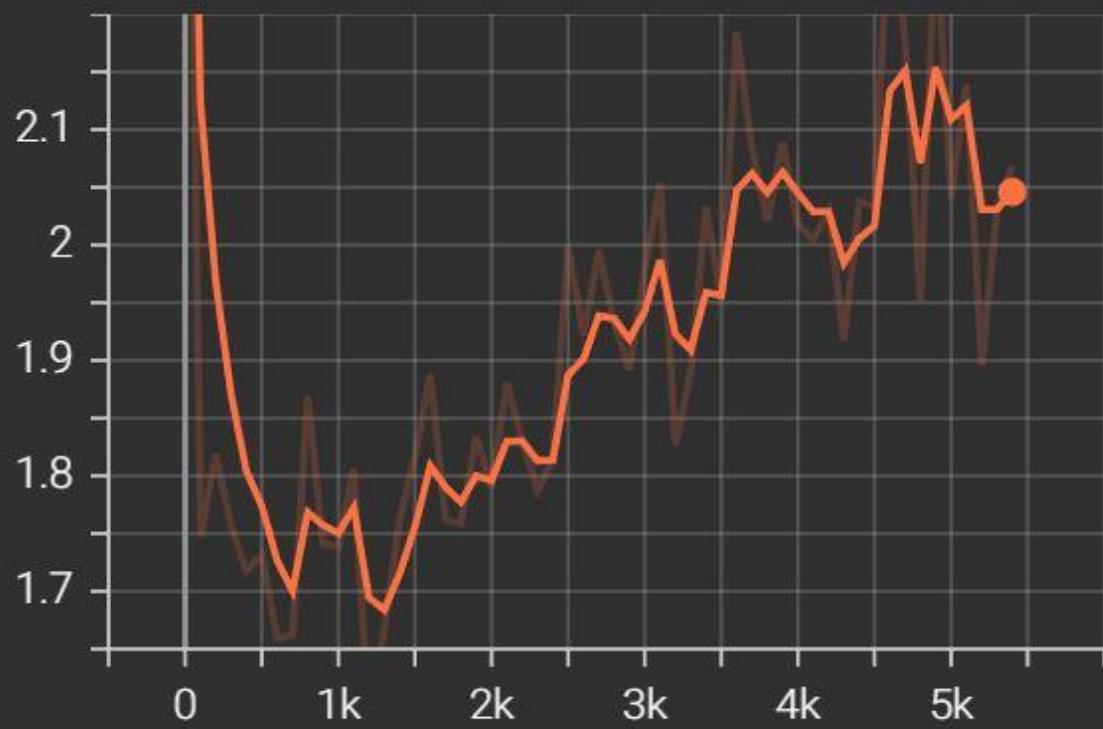
Max_h_fc1

tag: Max_h_fc1



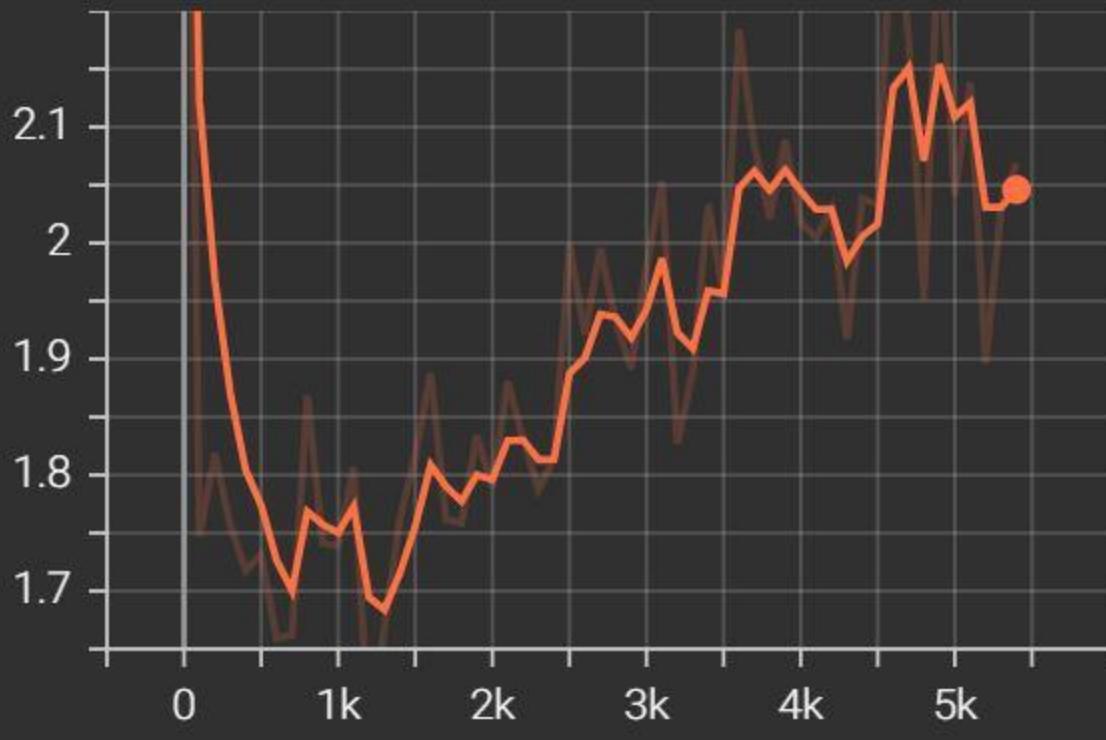
Max_h_pool2

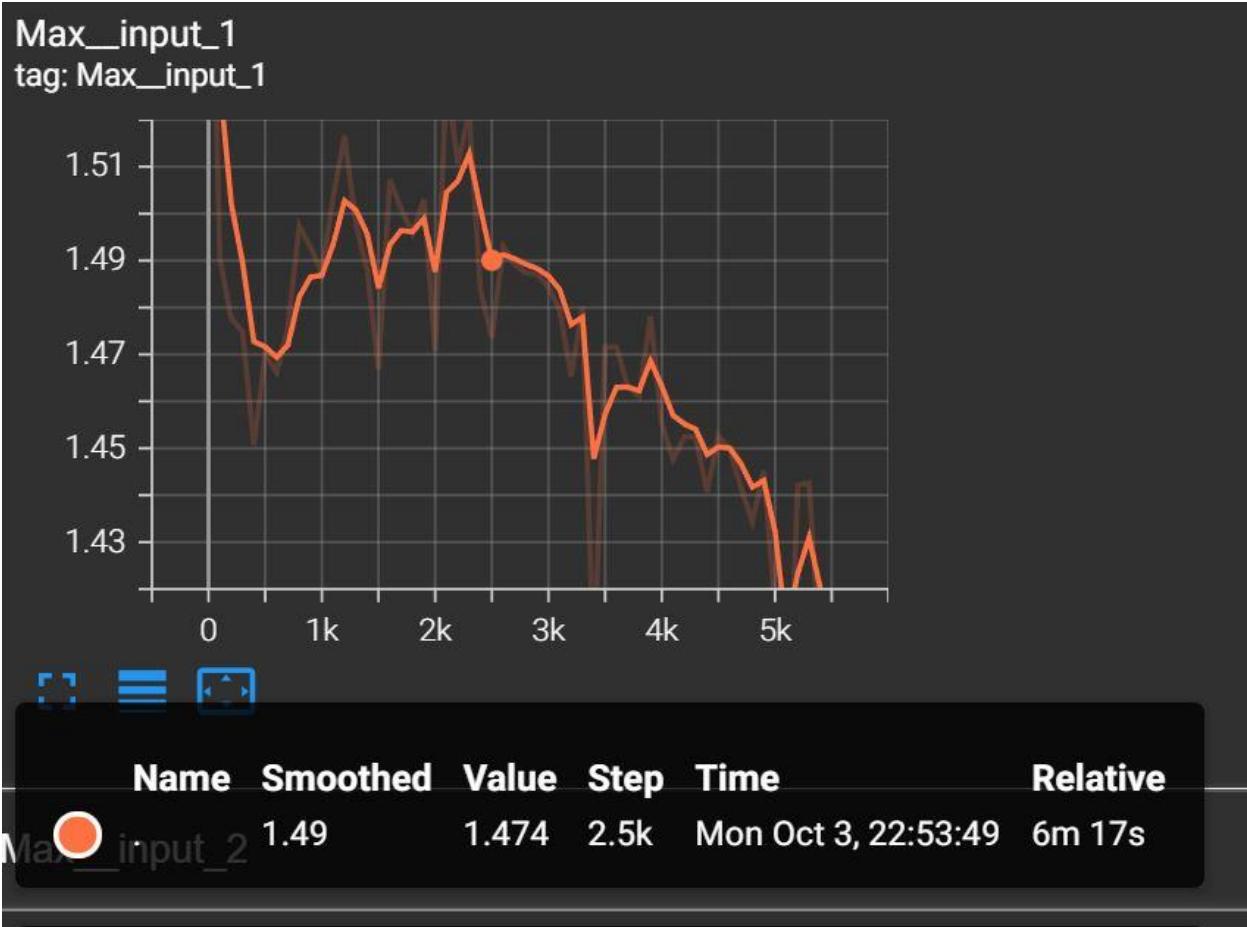
tag: Max_h_pool2



Max_h_pool2_flat

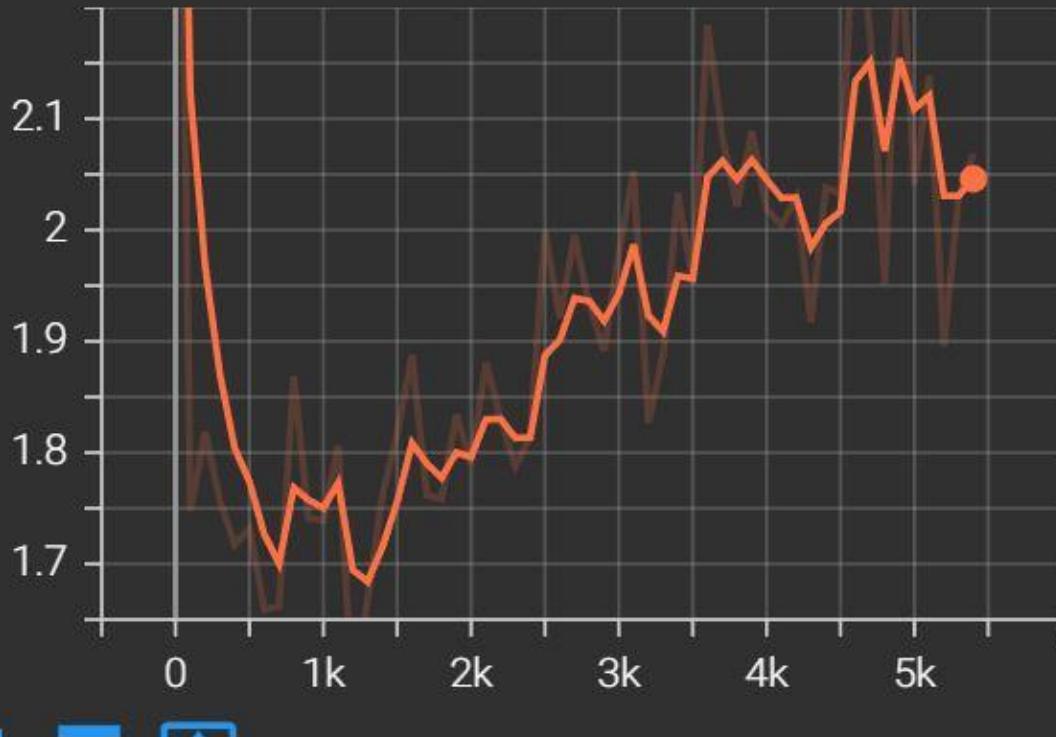
tag: Max_h_pool2_flat

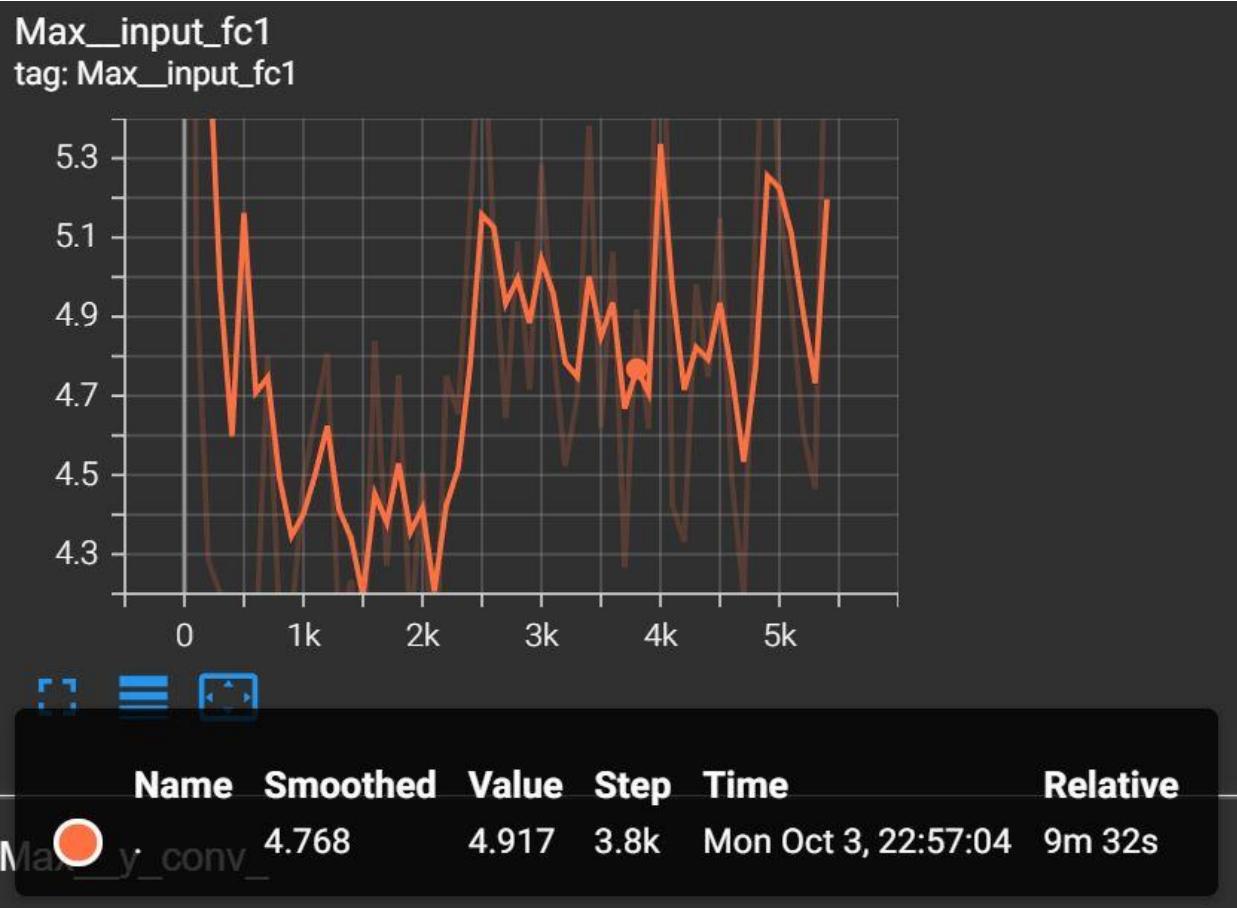




Max_input_2

tag: Max_input_2





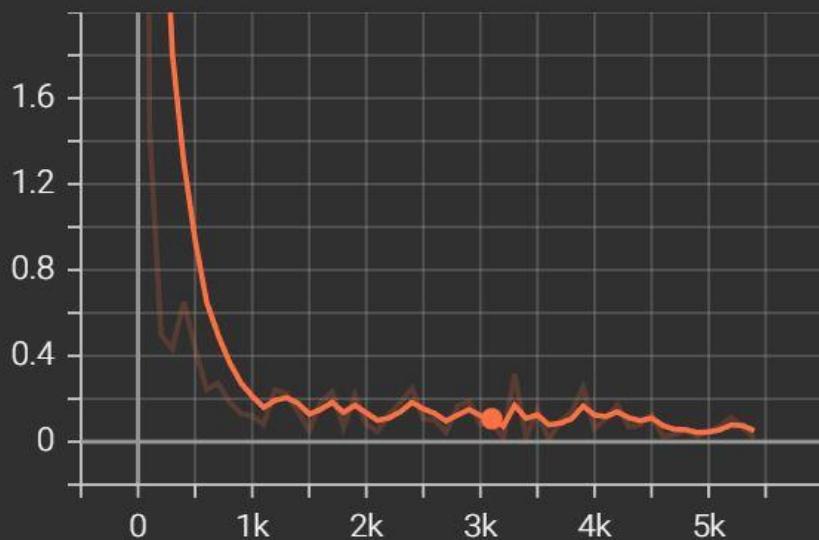
Max_y_conv_
tag: Max_y_conv_



Name	Smoothed	Value	Step	Time	Relative
Max_y_conv	1	1	1	2.2k	Mon Oct 3, 22:52:57

Mean_1

tag: Mean_1

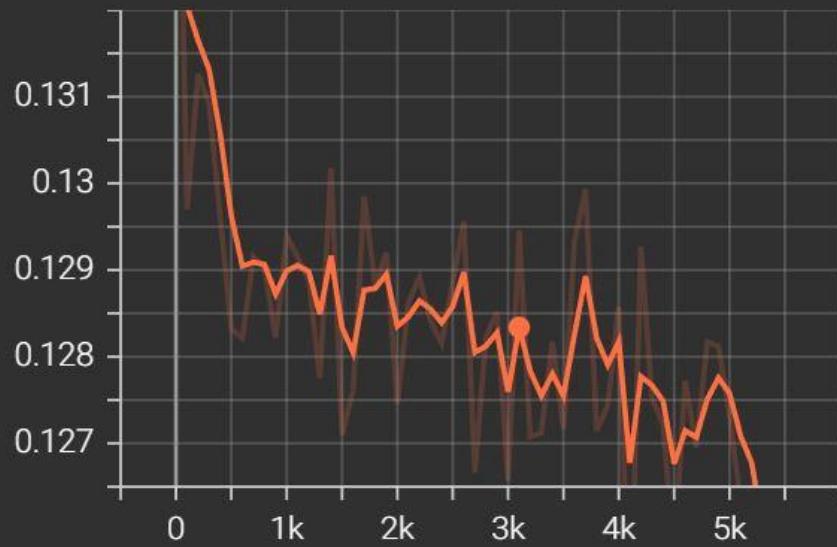


Name	Smoothed	Value	Step	Time	Relative
------	----------	-------	------	------	----------

Mean_1 · h_pool	0.107	0.08522	3.1k	Mon Oct 3, 22:55:13	7m 40s
-----------------	-------	---------	------	---------------------	--------

Mean__h_pool1

tag: Mean__h_pool1

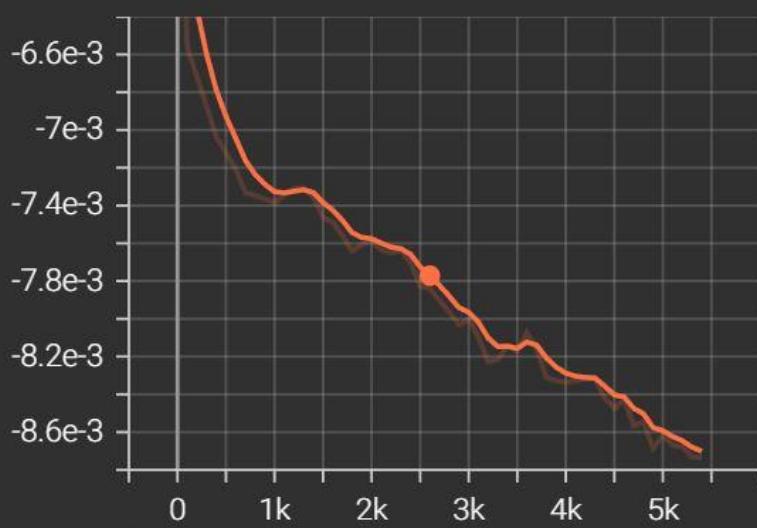


Name	Smoothed	Value	Step	Time	Relative
------	----------	-------	------	------	----------

Mean__h_pool1	0.1283	0.1295	3.1k	Mon Oct 3, 22:55:13	7m 40s
---------------	--------	--------	------	---------------------	--------

Mean_W_conv1

tag: Mean_W_conv1



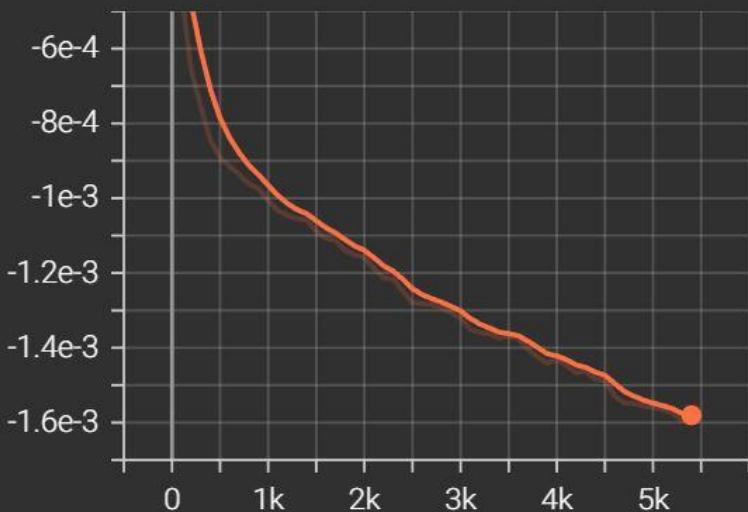
Name	Smoothed	Value	Step	Time	Relative
Mean_W_conv1	-7.7712e-3	-7.8403e-3	2.6k	Mon Oct 3, 22:54:03	6m 31s

Mean_W_conv2
tag: Mean_W_conv2



Mean_W_fc1

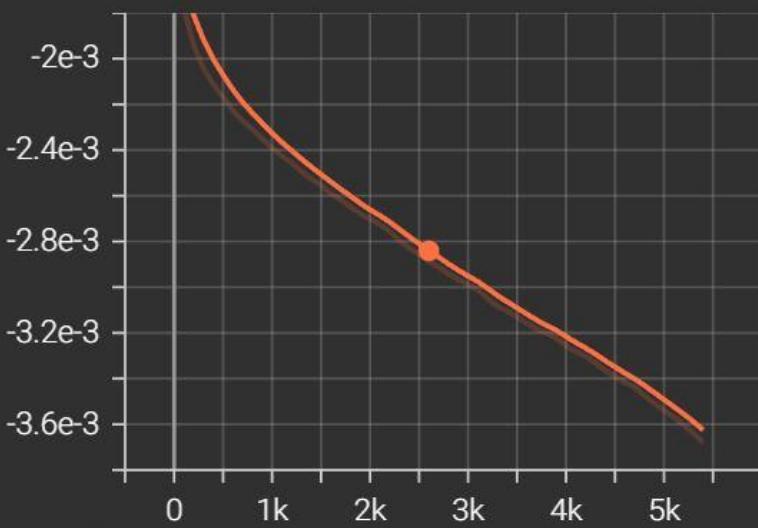
tag: Mean_W_fc1



Name	Smoothed	Value	Step	Time	Relative
Mean_W_fc2	-5.9061e-4	-6.2932e-4	1.3k	Mon Oct 3, 22:50:52	3m 19s

Mean_W_fc2

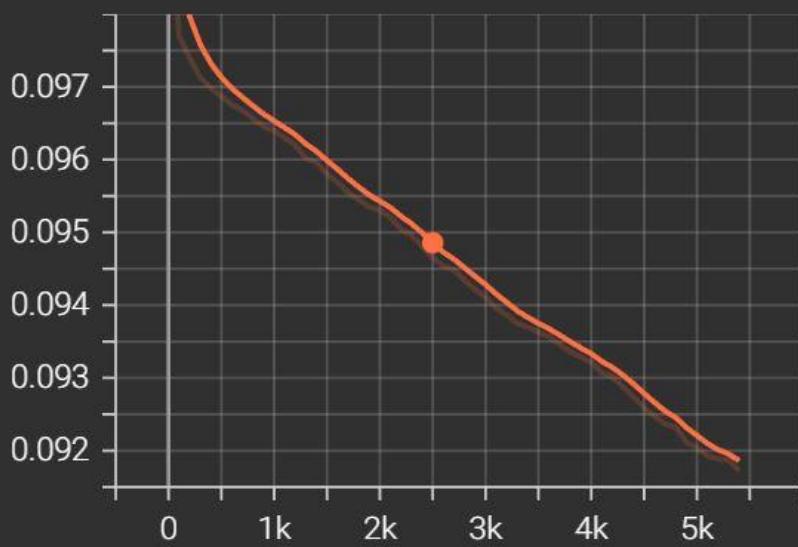
tag: Mean_W_fc2



Name	Smoothed	Value	Step	Time	Relative
Mean_b_conv	-2.8408e-3	-2.8869e-3	2.6k	Mon Oct 3, 22:54:03	6m 31s

Mean_b_conv1

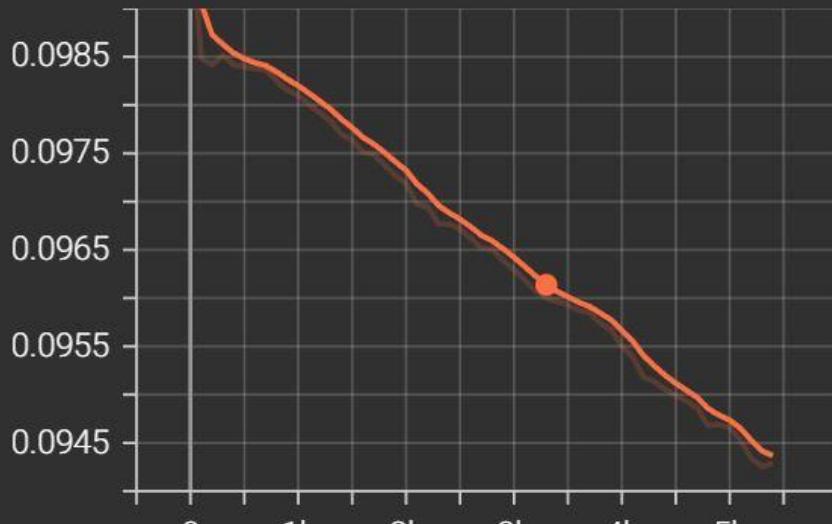
tag: Mean_b_conv1



Name	Smoothed	Value	Step	Time	Relative
Mean_b_conv1	0.09486	0.09465	2.5k	Mon Oct 3, 22:53:49	6m 17s

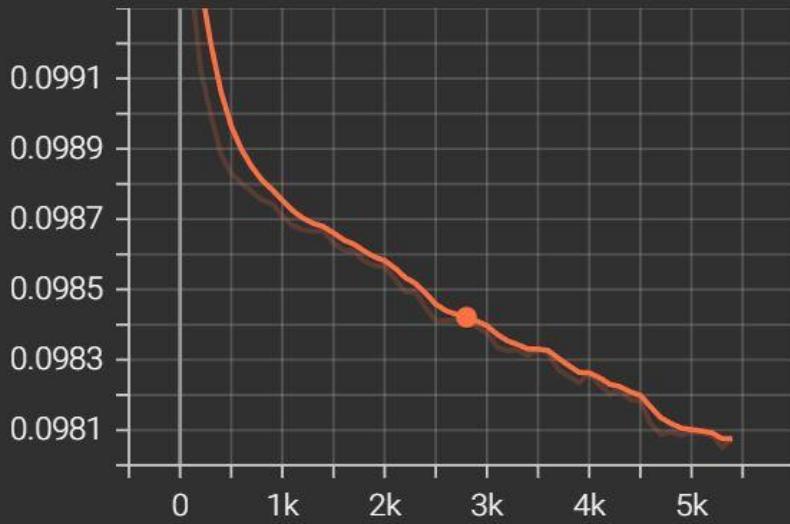
Mean_b_conv2

tag: Mean_b_conv2



Name	Smoothed	Value	Step	Time	Relative
Mean_b_fc1	0.09614	0.096	3.3k	Mon Oct 3, 22:55:40	8m 8s

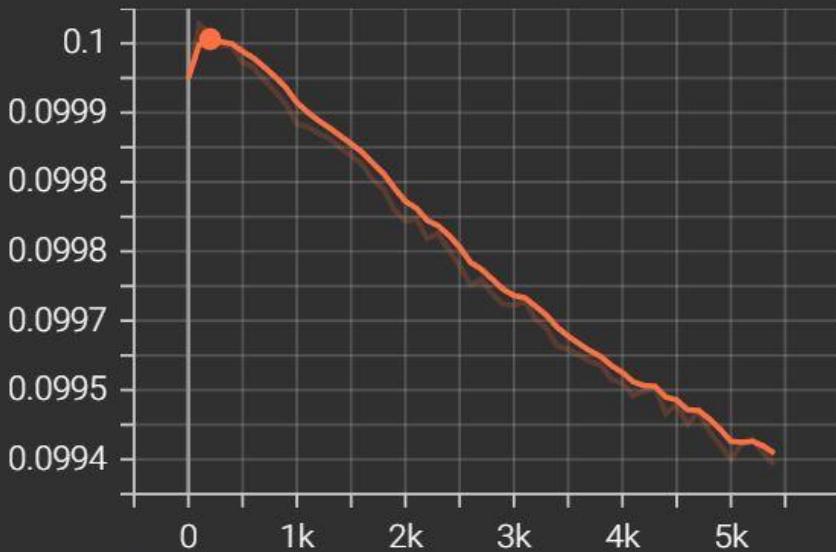
Mean_b_fc1
tag: Mean_b_fc1



Name	Smoothed	Value	Step	Time	Relative
Mean_b_fc2	0.09842	0.09841	2.8k	Mon Oct 3, 22:54:31	6m 59s

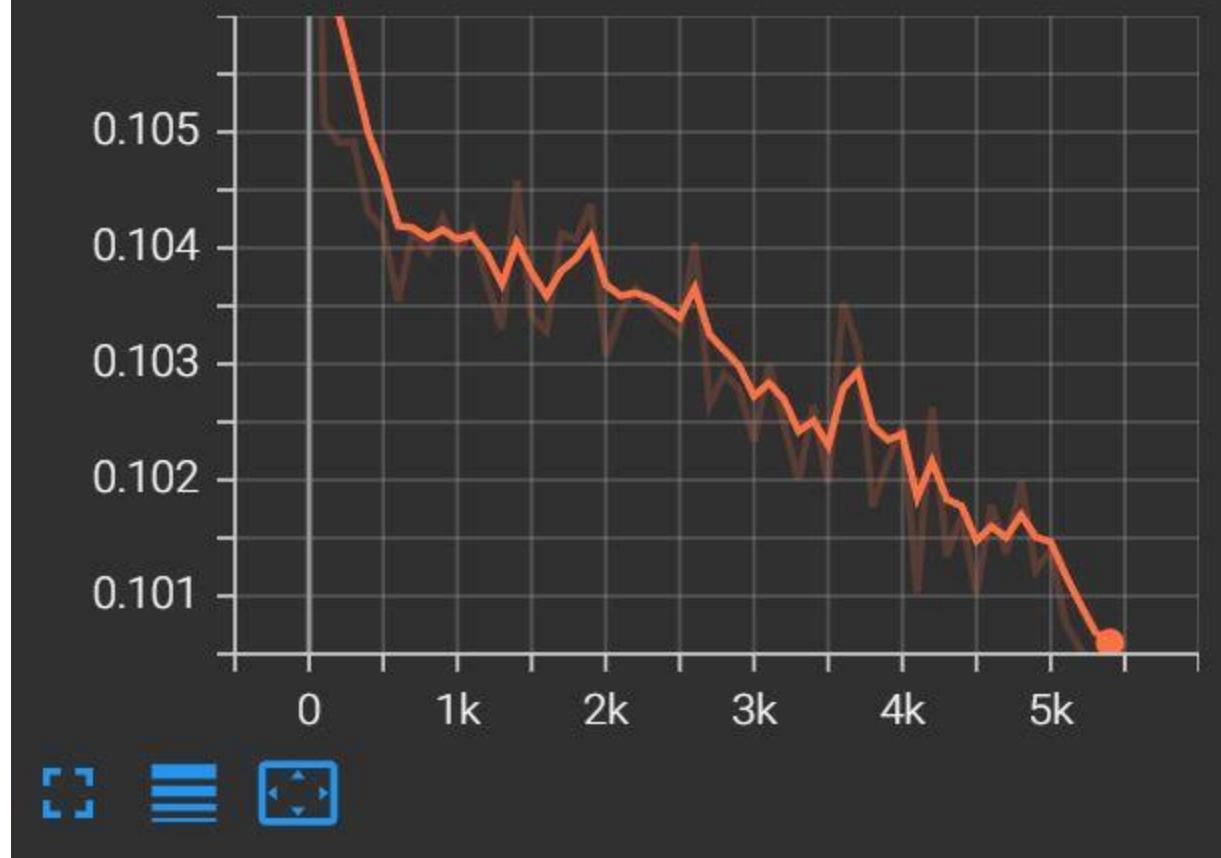
Mean_b_fc2

tag: Mean_b_fc2

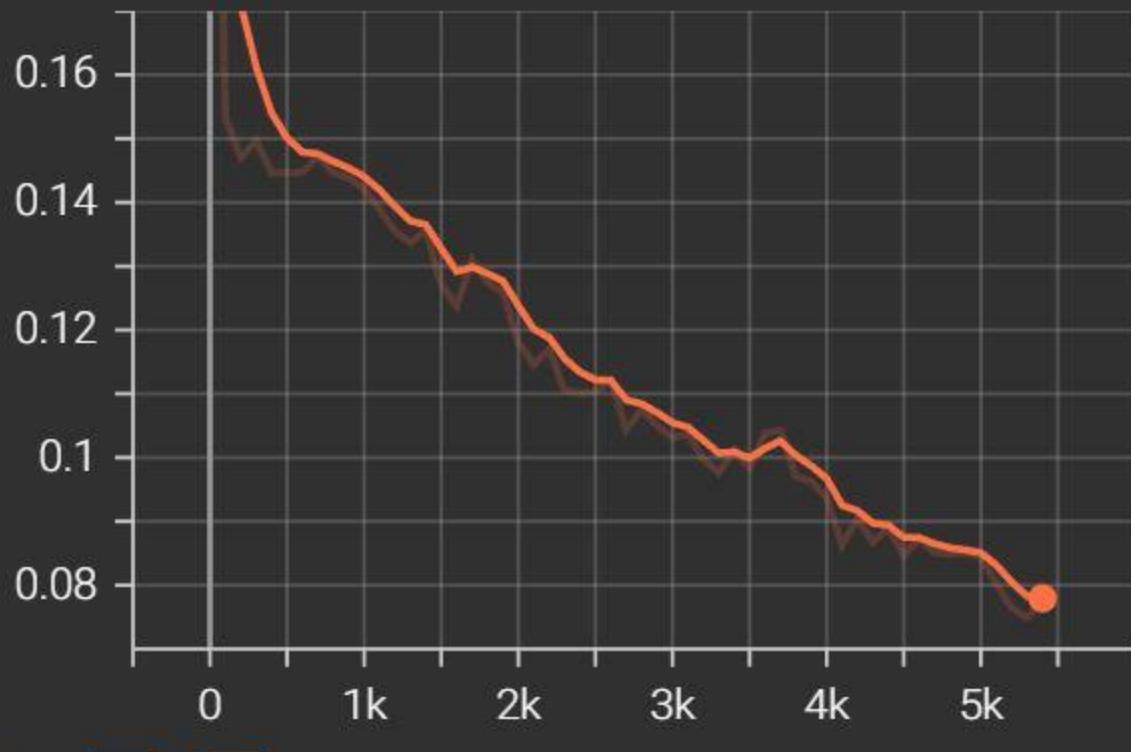


Name	Smoothed	Value	Step	Time	Relative
Name	Smoothed	Value	Step	Time	Relative
Mean_b_fc2	0.1001	0.1001	200	Mon Oct 3, 22:48:09	37s
Mean_b_fc2	0.09634	0.0962	1.2k	Mon Oct 3, 22:50:38	3m 6s

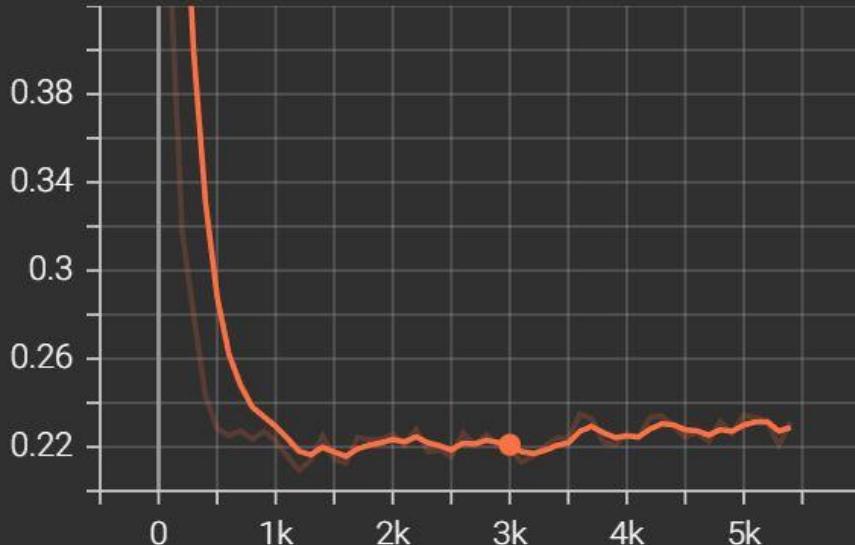
Mean_h_conv1
tag: Mean_h_conv1



Mean_h_conv2
tag: Mean_h_conv2



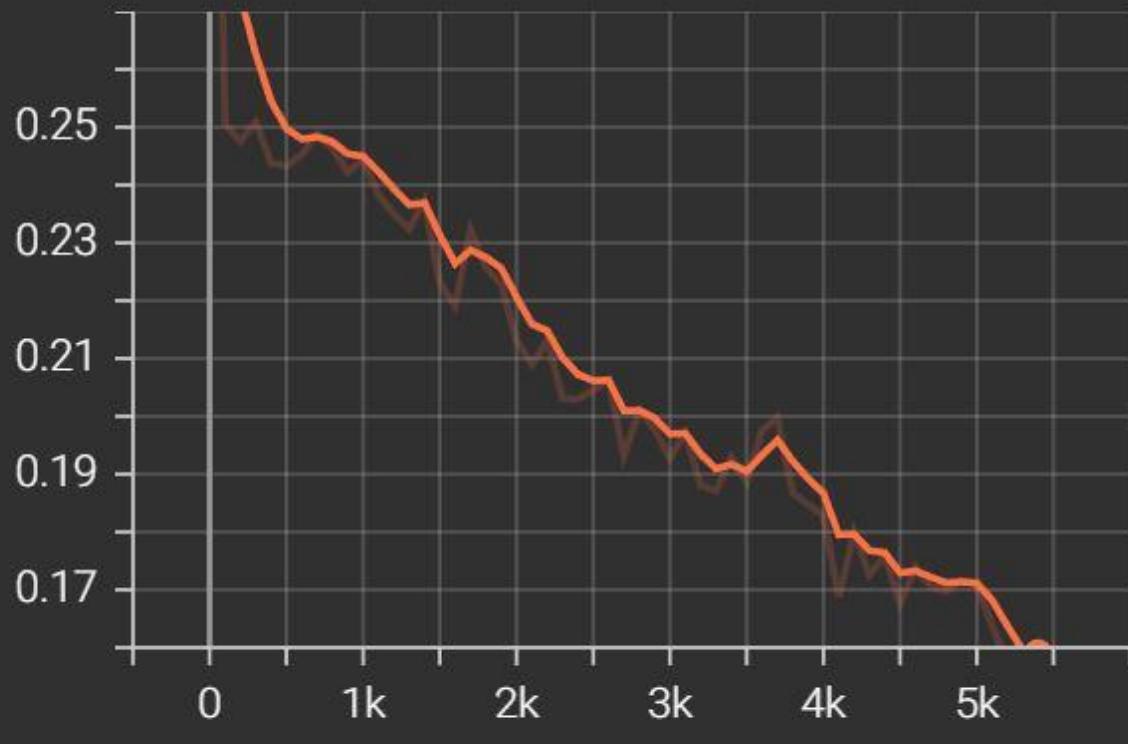
Mean_h_fc1
tag: Mean_h_fc1



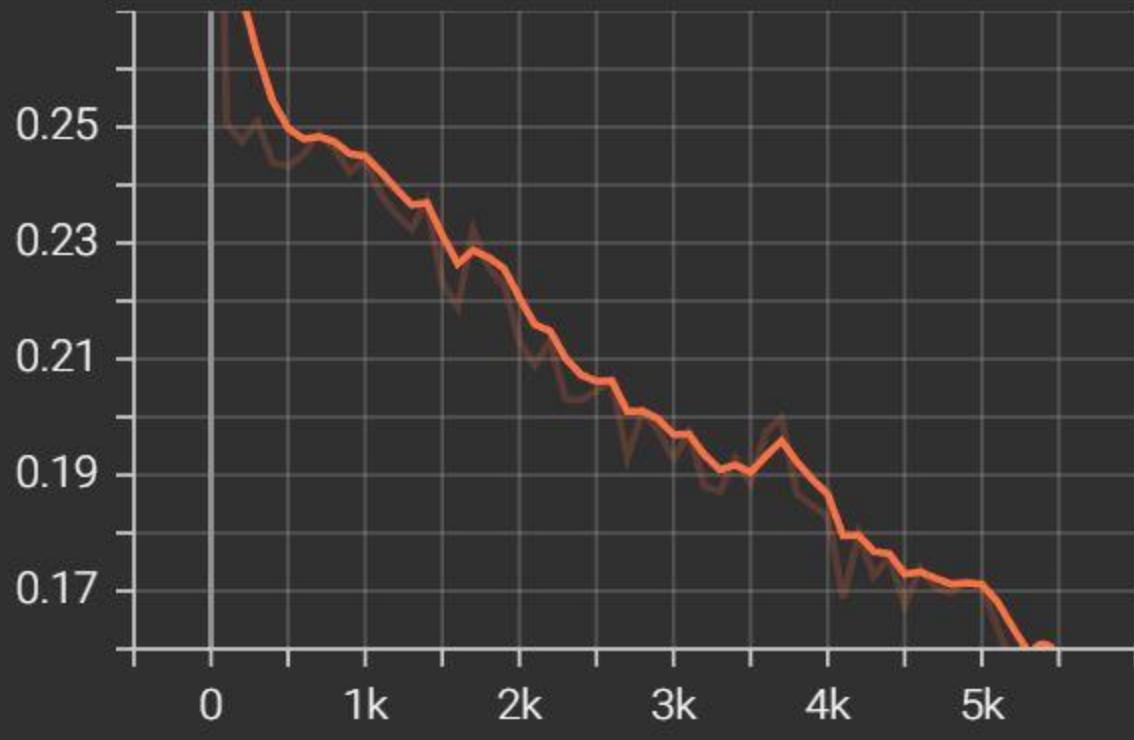
Name	Smoothed	Value	Step	Time	Relative
------	----------	-------	------	------	----------

Mean_h_pool	0.2209	0.2193	3k	Mon Oct 3, 22:54:59	7m 26s
-------------	--------	--------	----	---------------------	--------

Mean_h_pool2
tag: Mean_h_pool2

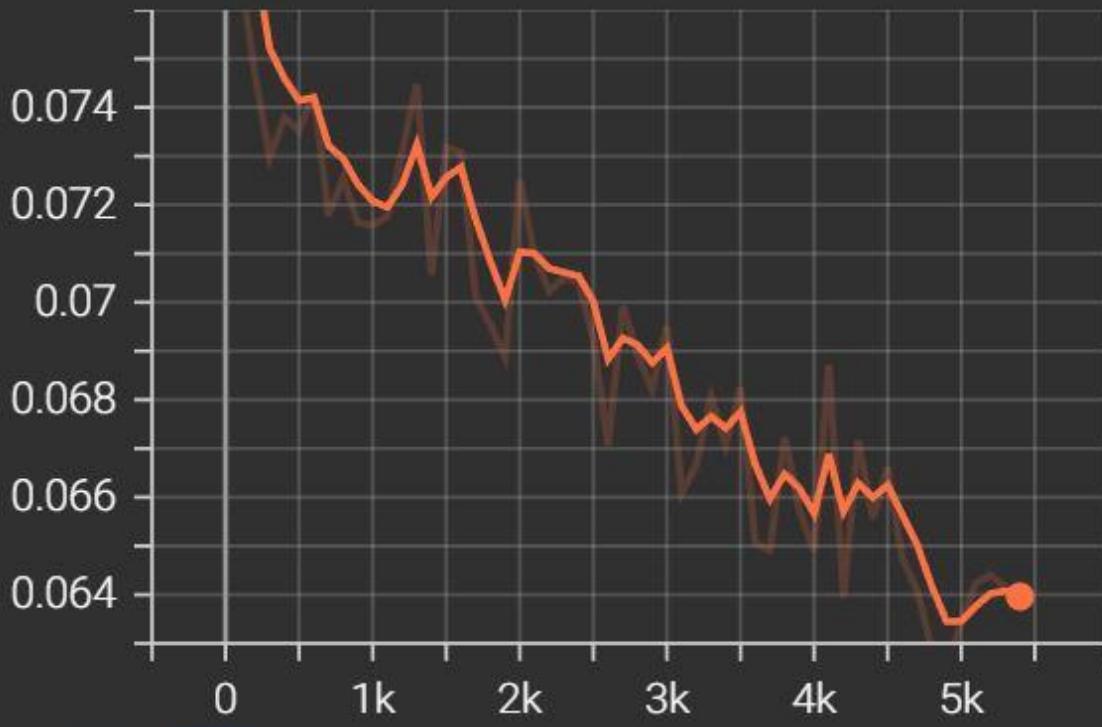


Mean_h_pool2_flat
tag: Mean_h_pool2_flat

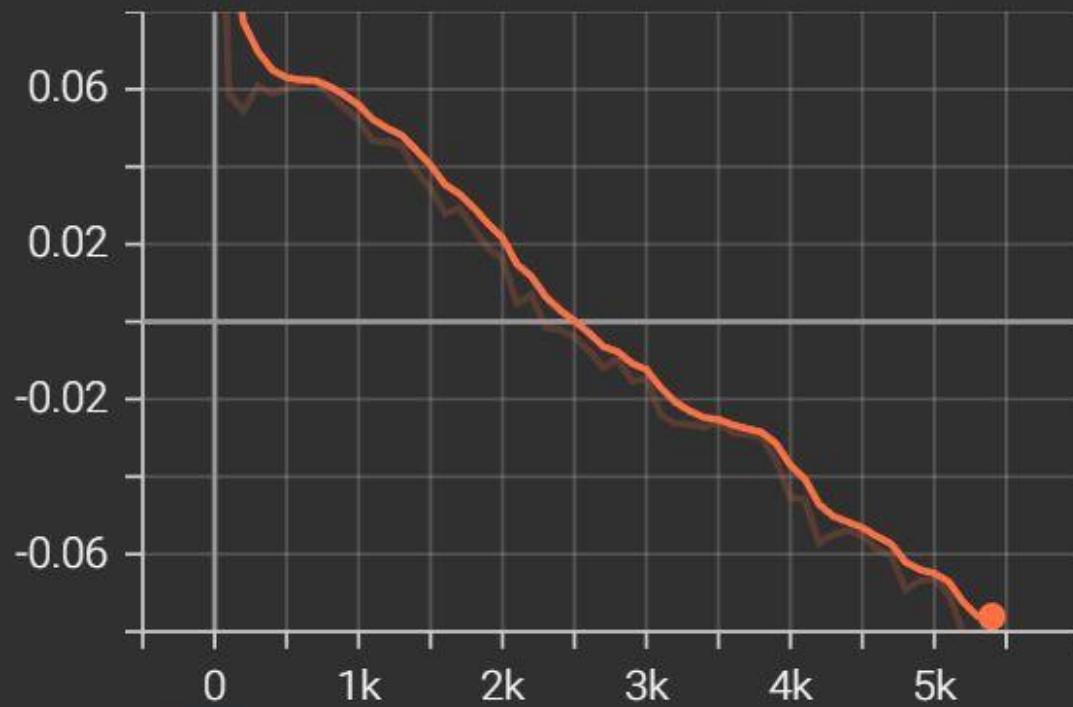


Mean_input_1

tag: Mean_input_1

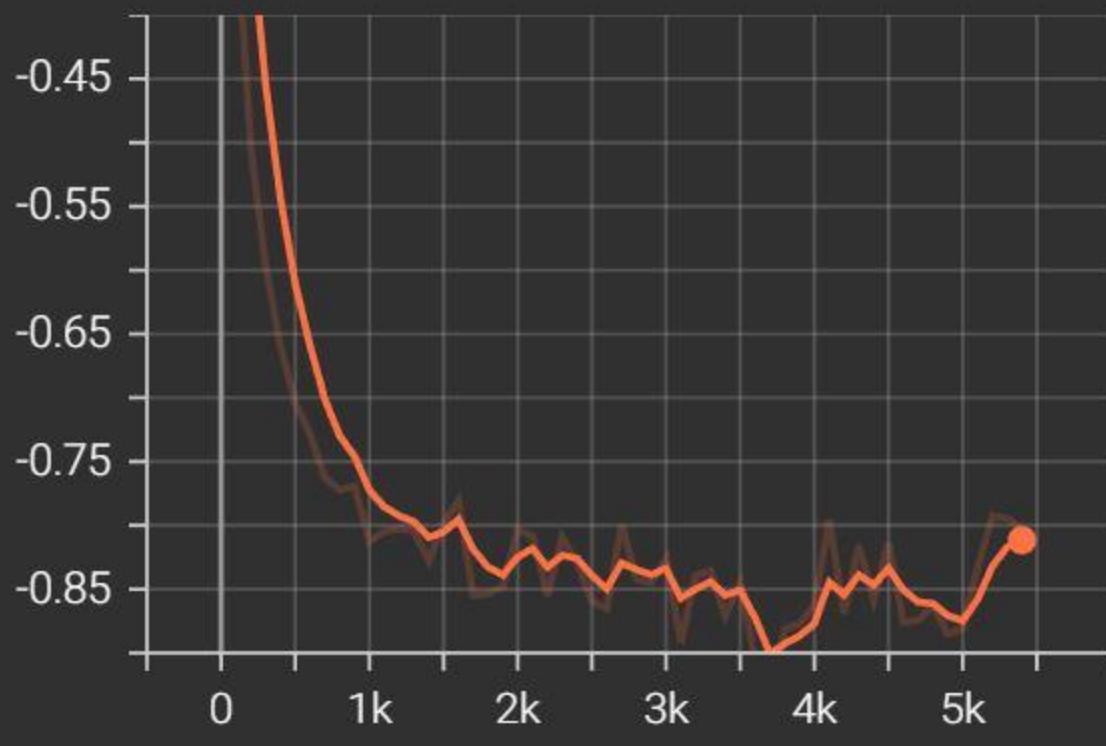


Mean_input_2
tag: Mean_input_2

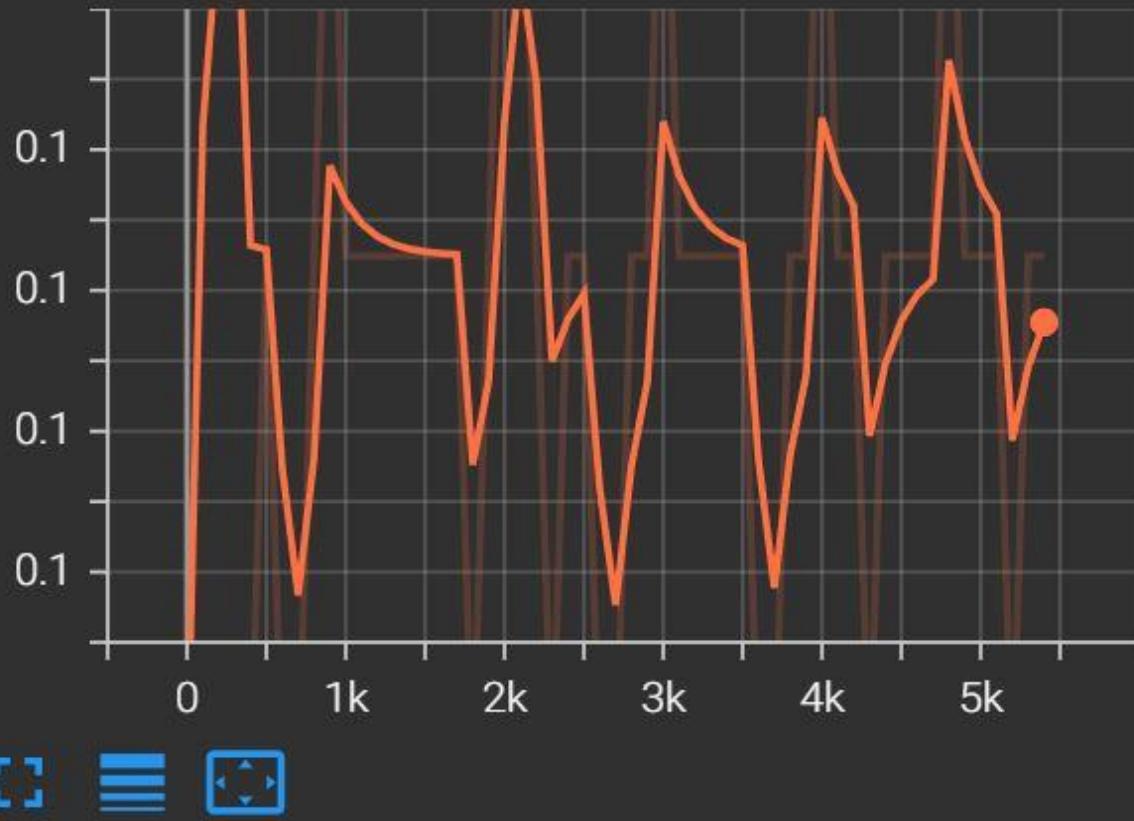


Mean_input_fc1

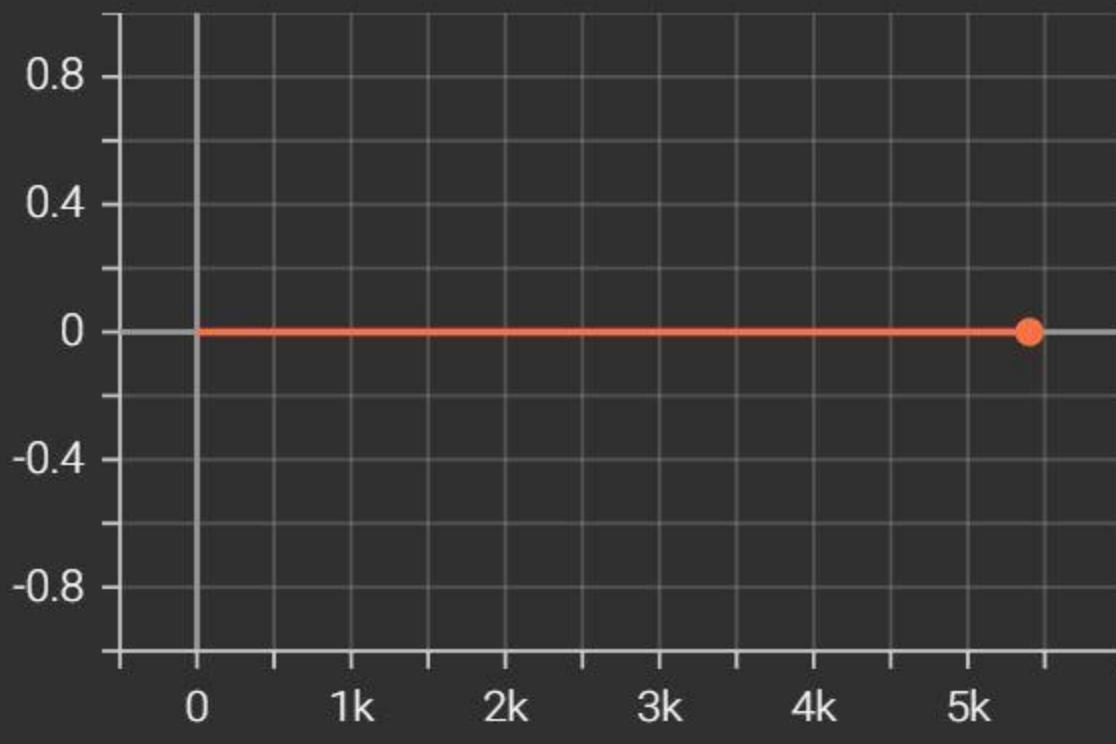
tag: Mean_input_fc1



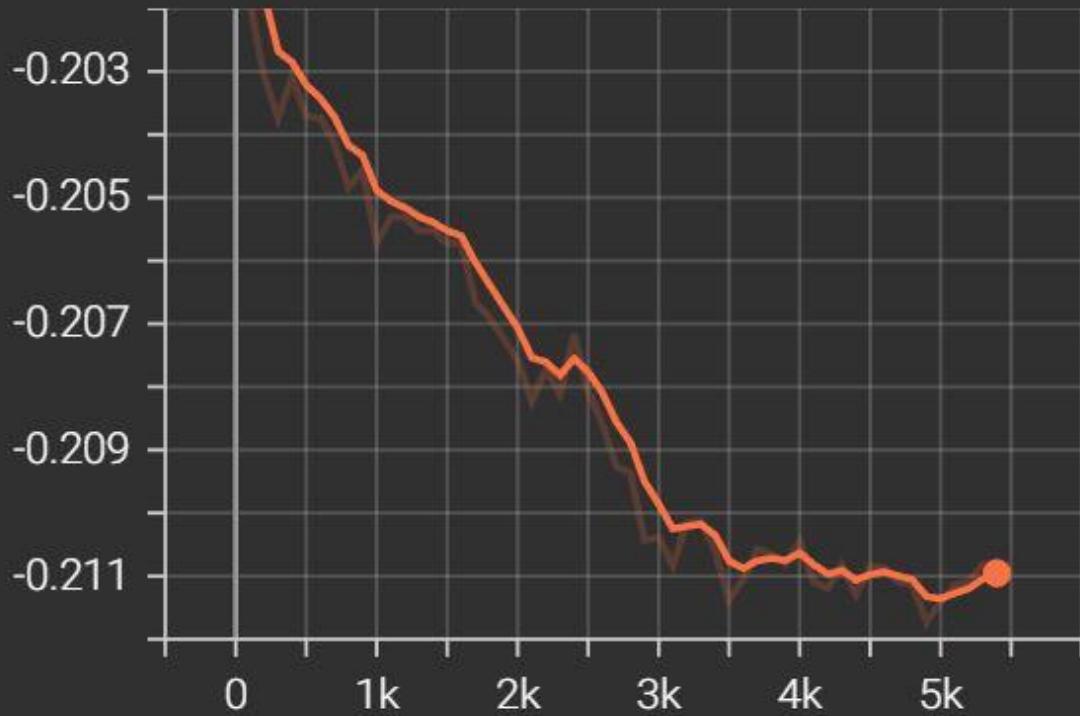
Mean_y_conv_
tag: Mean_y_conv_



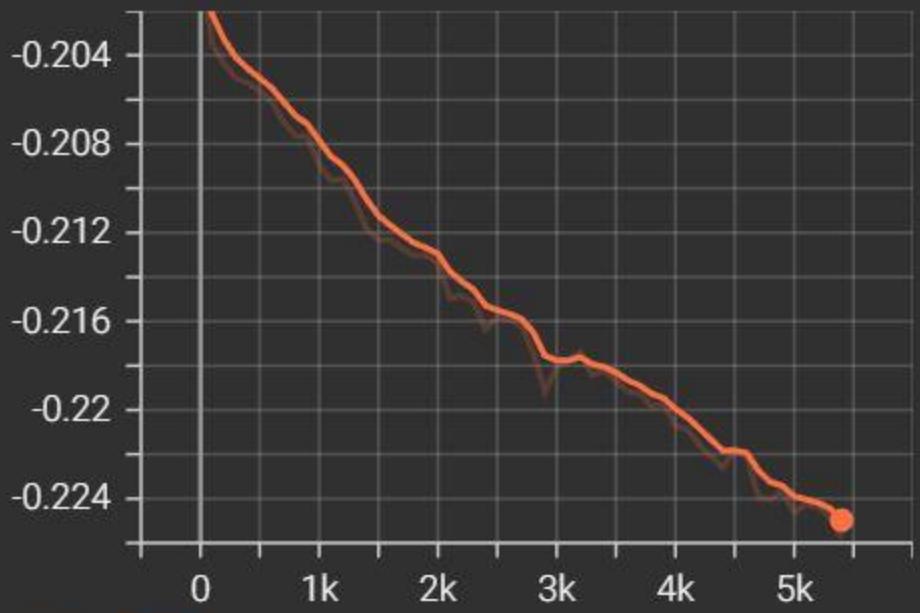
Min__h_pool1
tag: Min__h_pool1



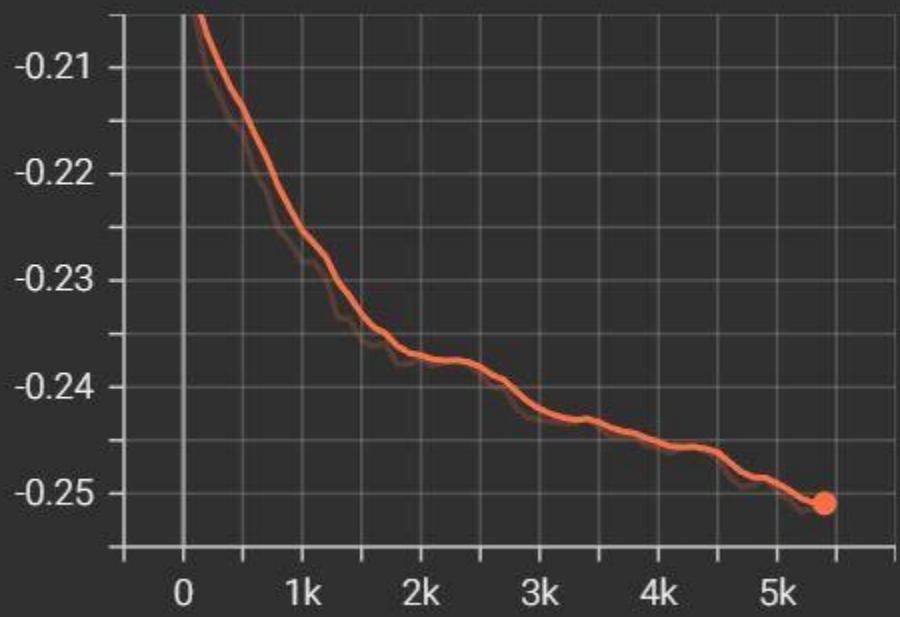
Min_W_conv1
tag: Min_W_conv1



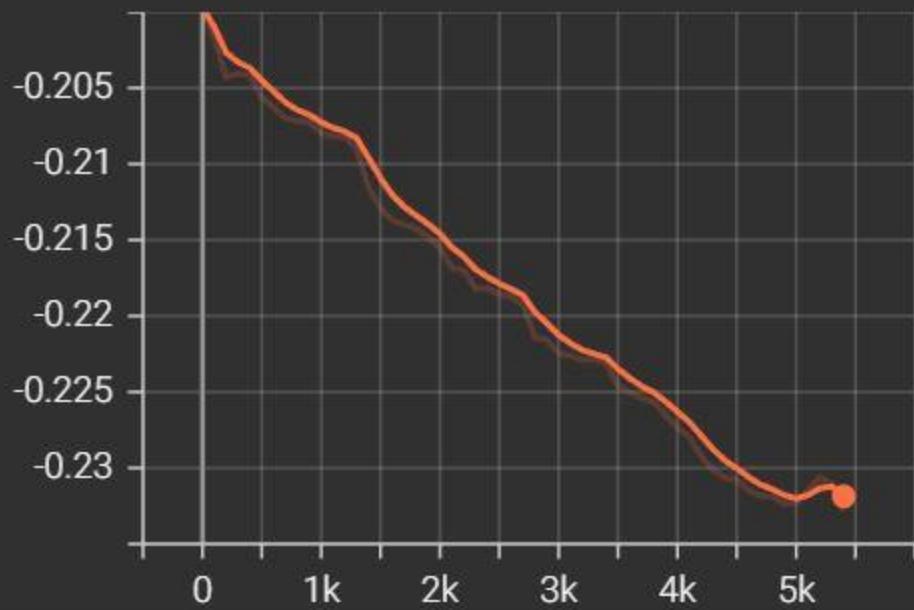
Min_W_conv2
tag: Min_W_conv2



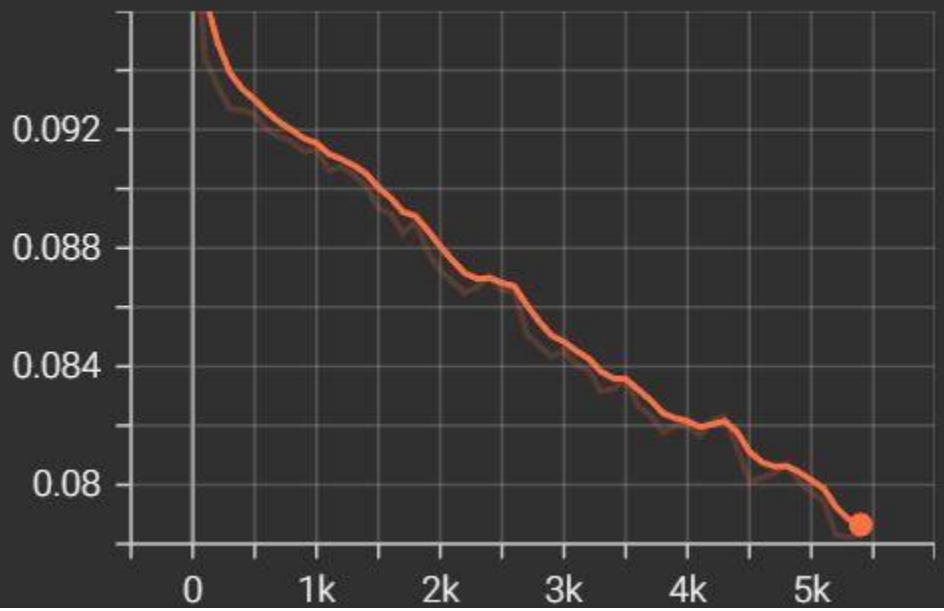
Min_W_fc1
tag: Min_W_fc1

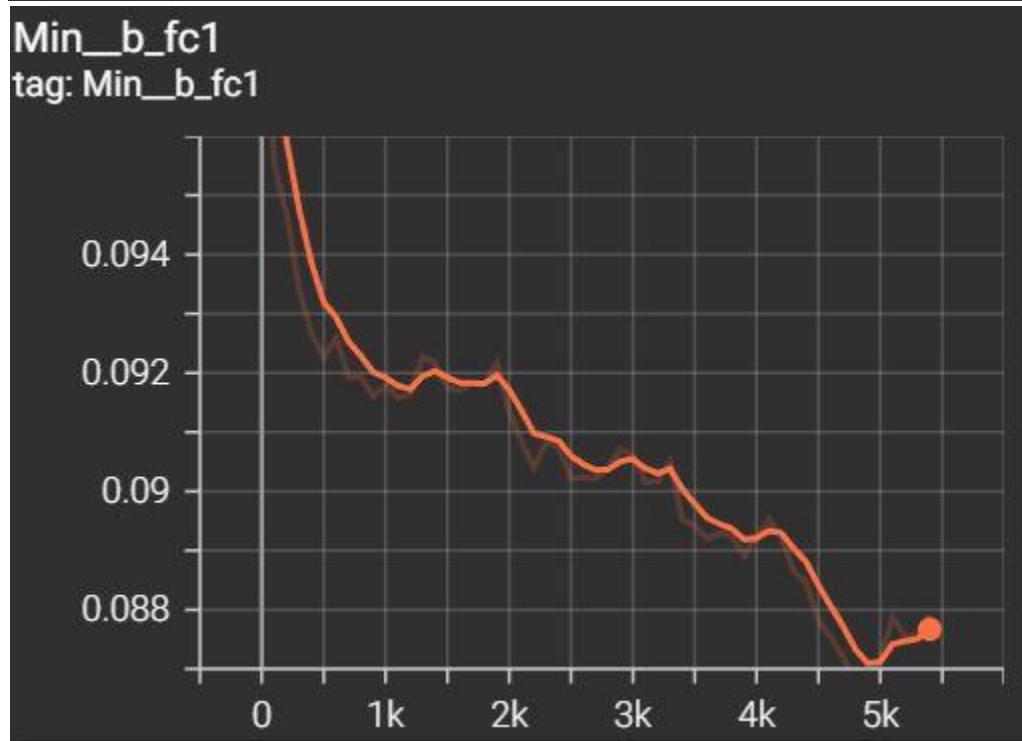
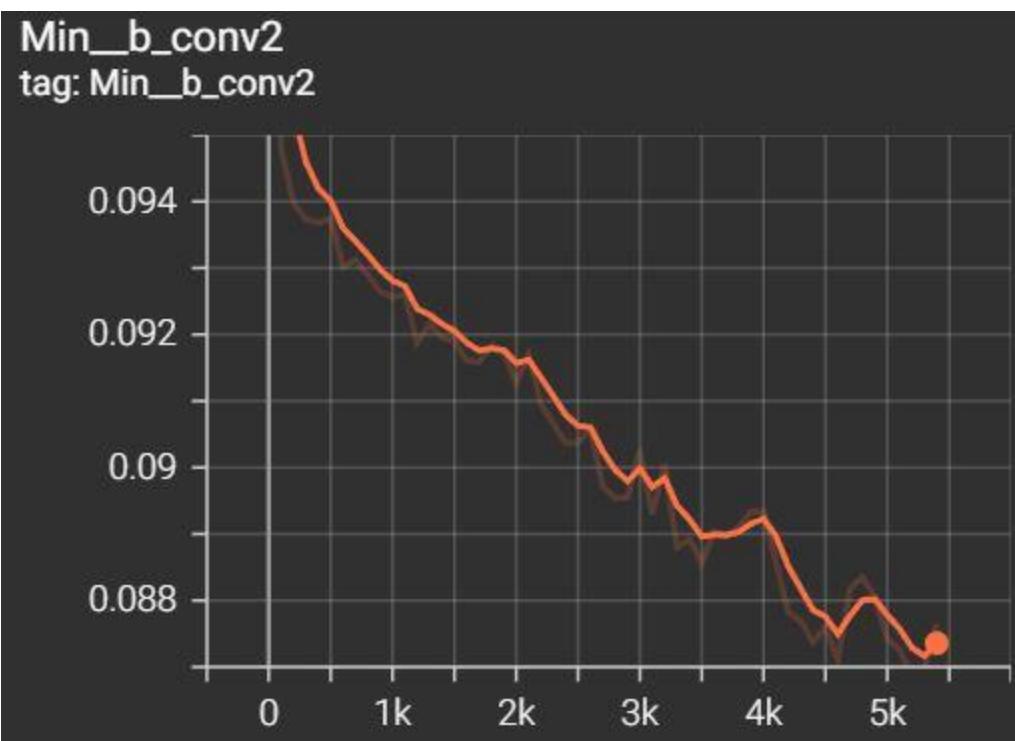


Min_W_fc2
tag: Min_W_fc2

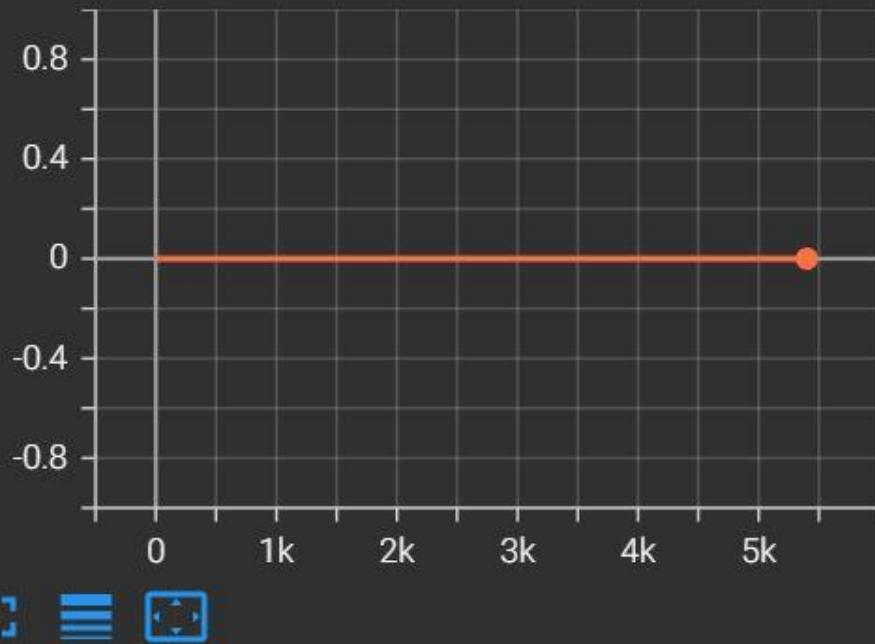


Min_b_conv1
tag: Min_b_conv1



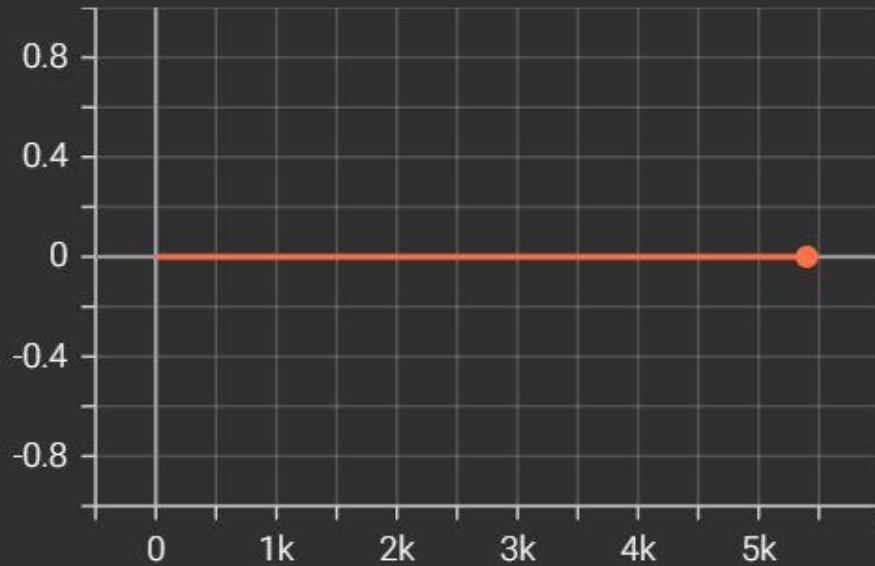


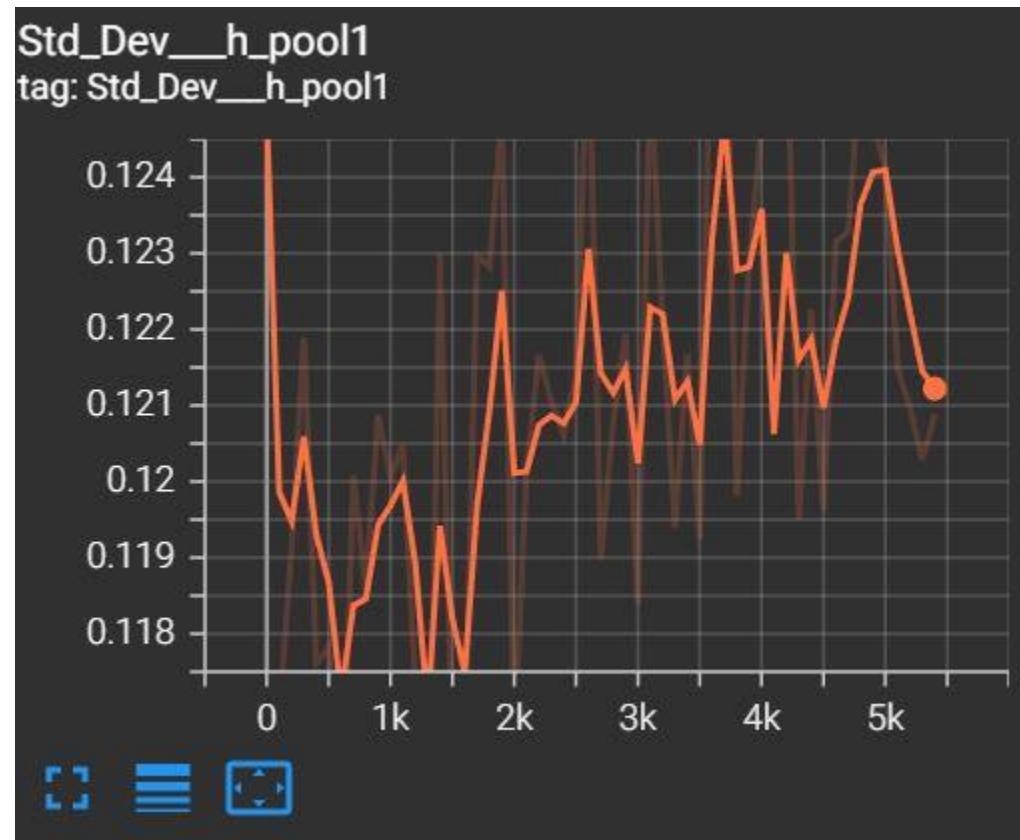
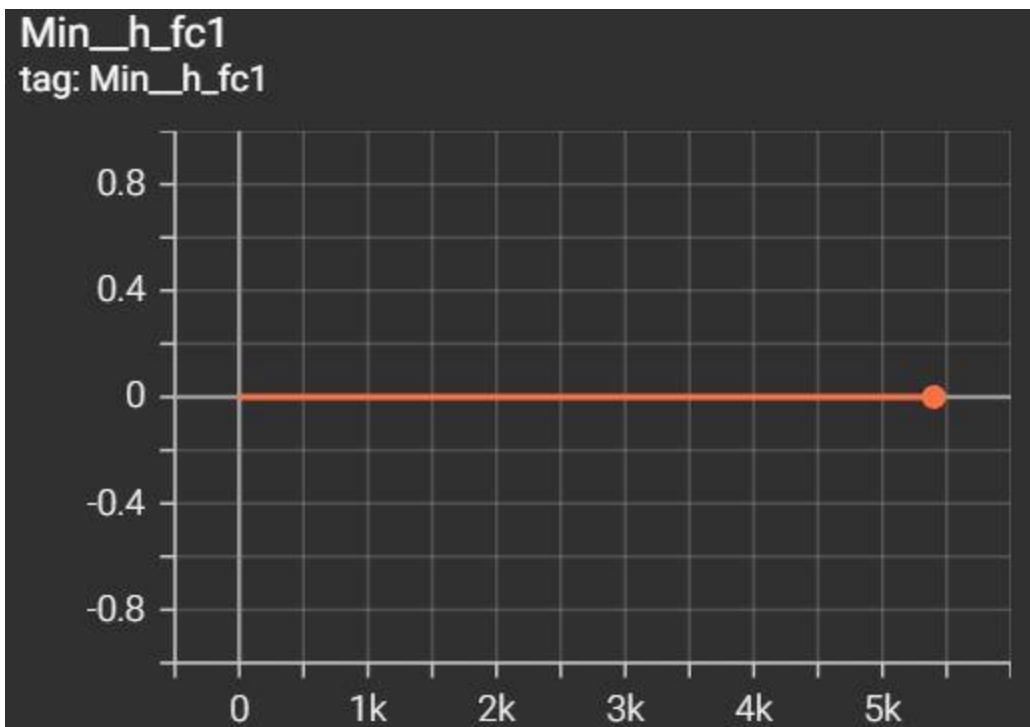
Min_h_conv1
tag: Min_h_conv1



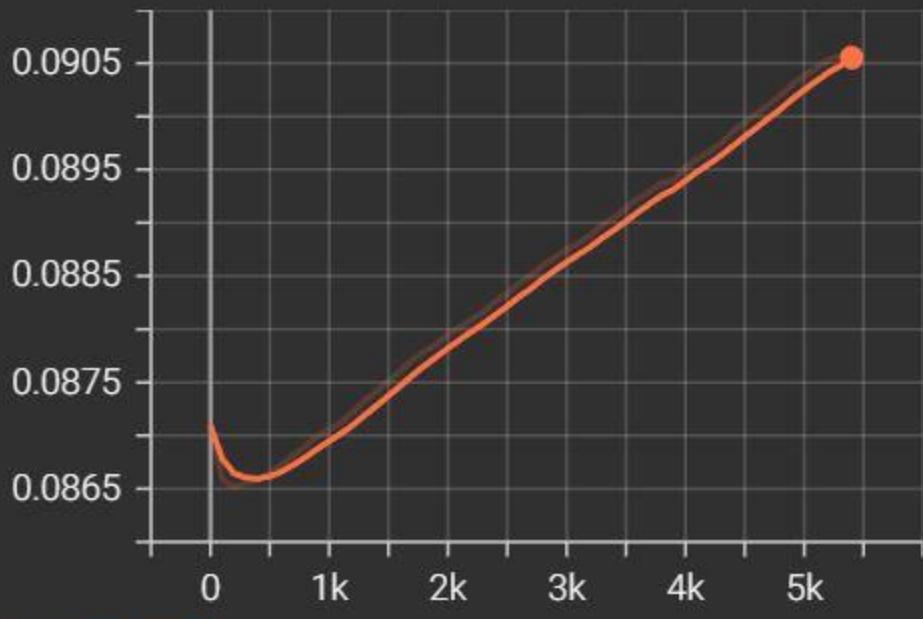
Min_h_conv2

Min_h_conv2
tag: Min_h_conv2

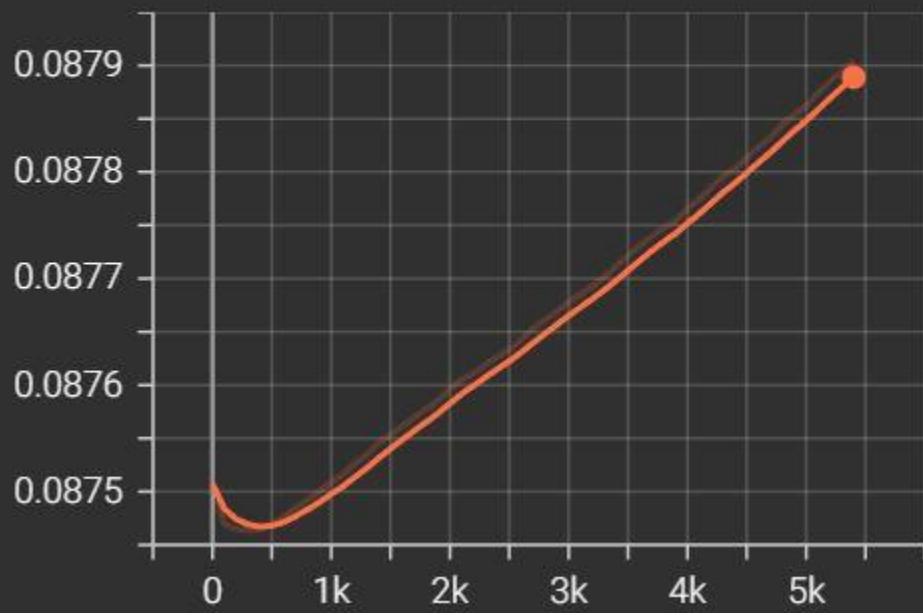




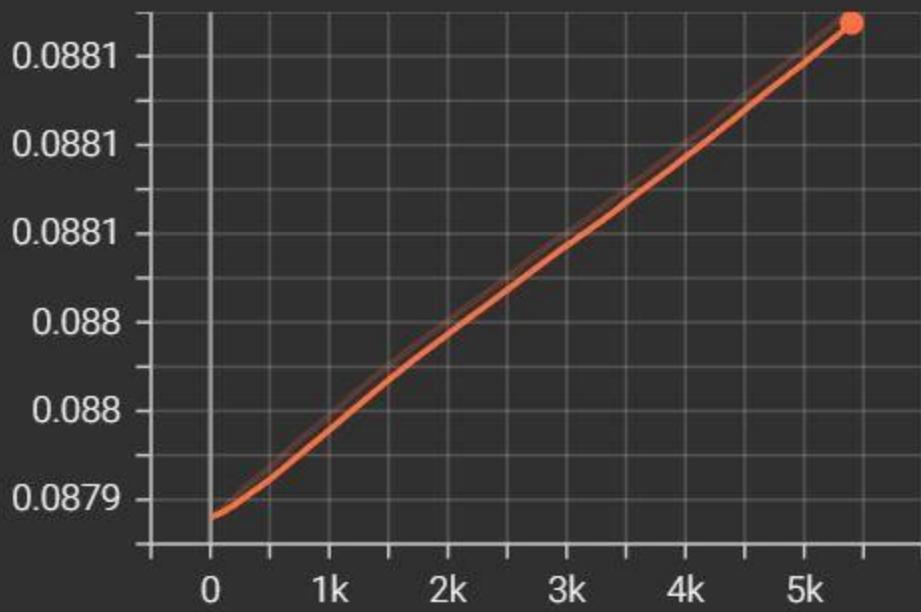
Std_Dev_W_conv1
tag: Std_Dev_W_conv1



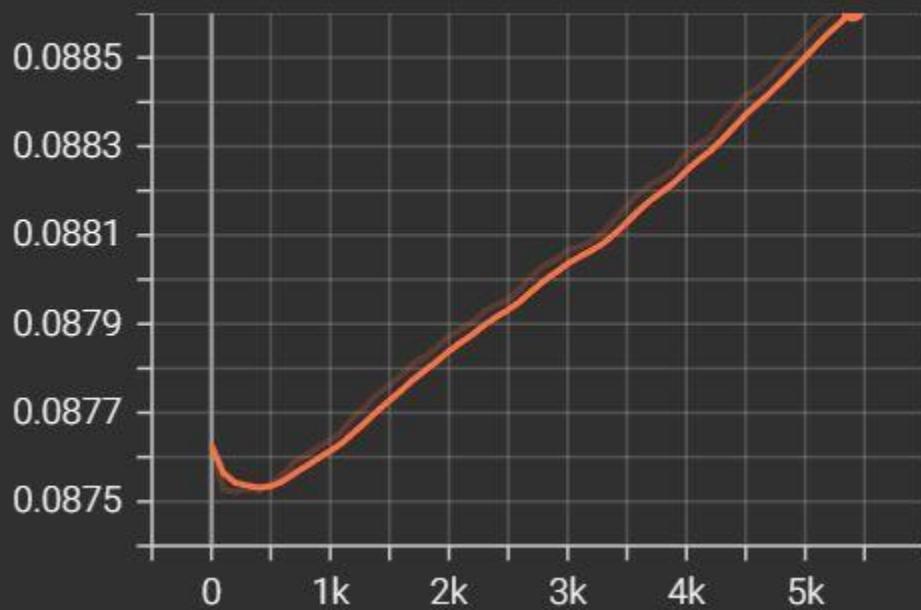
Std_Dev_W_conv2
tag: Std_Dev_W_conv2

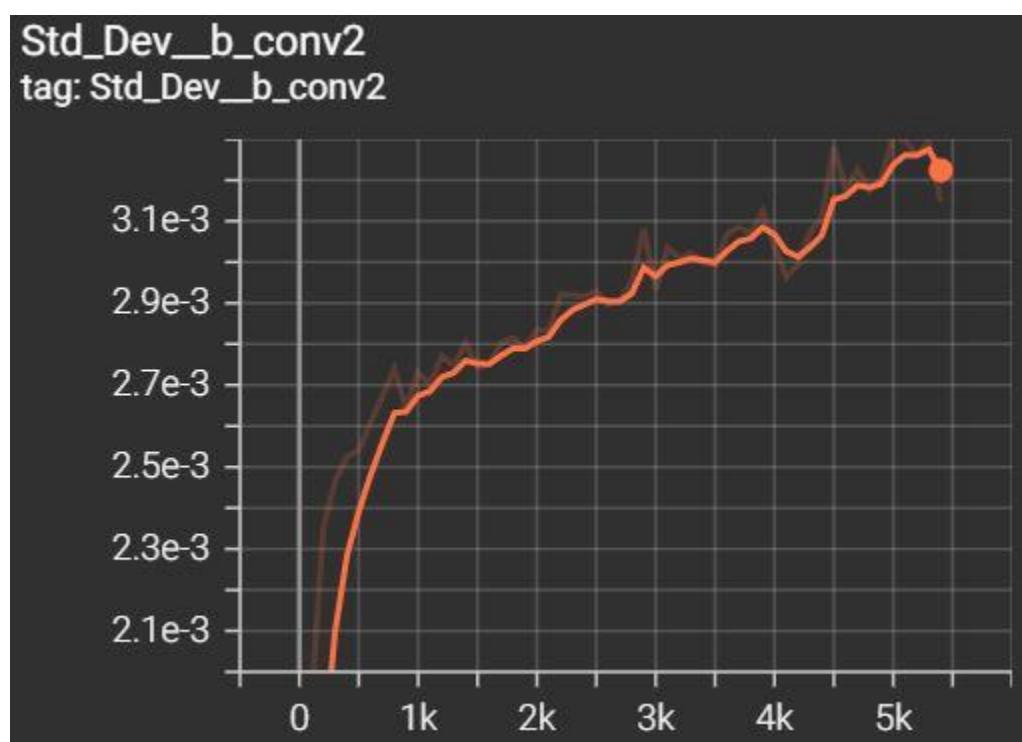
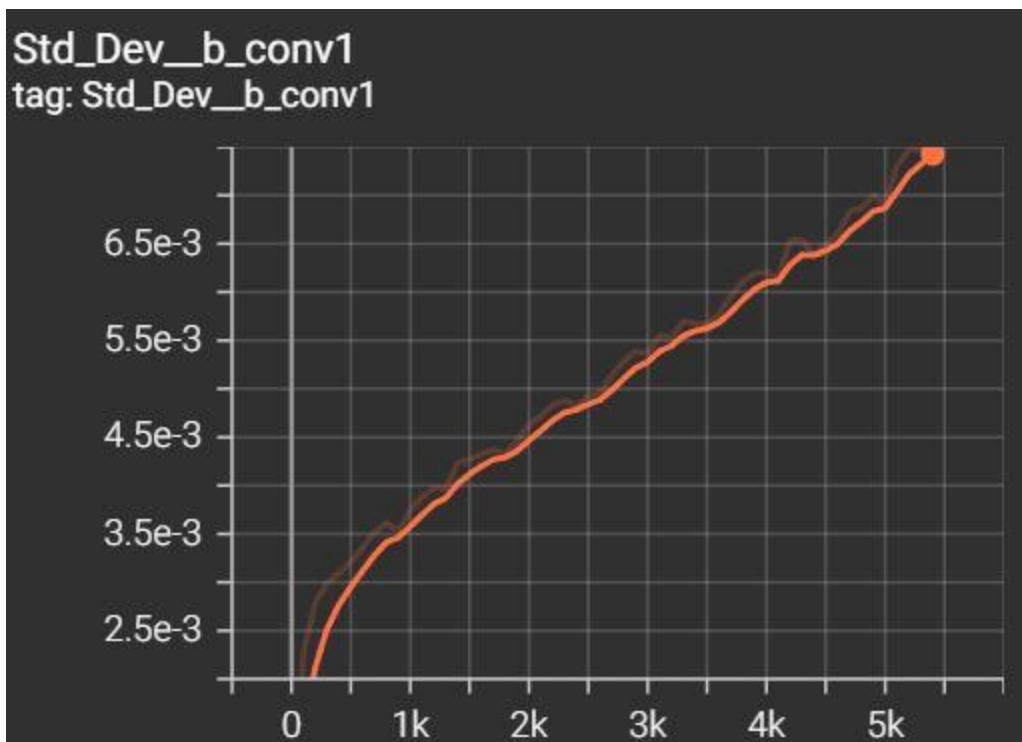


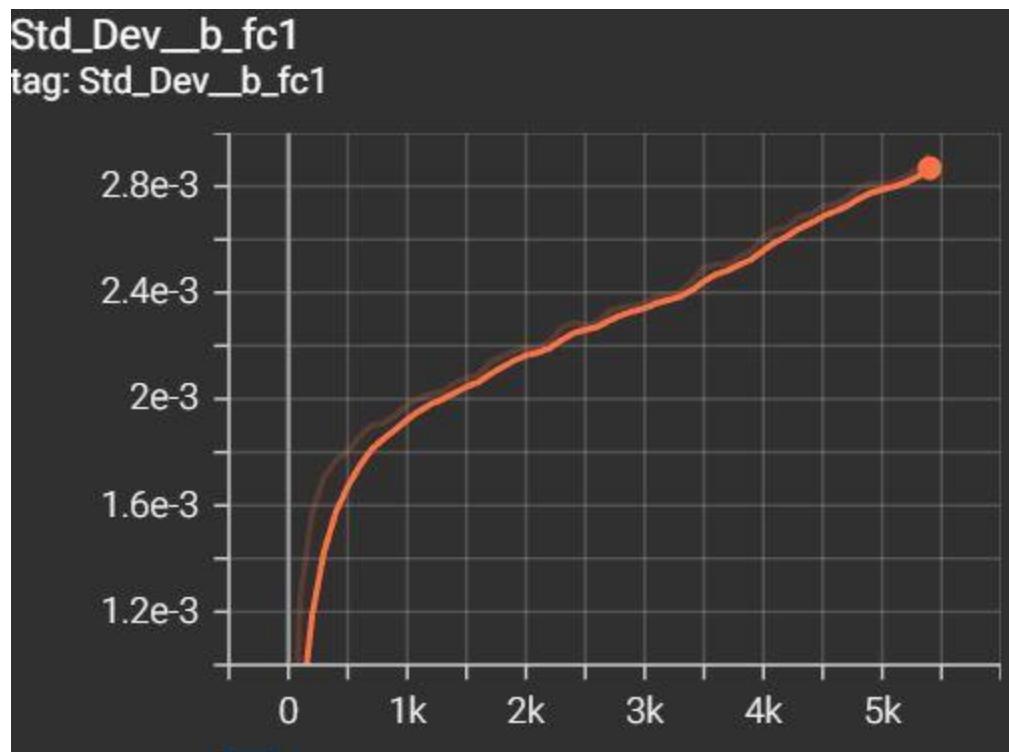
Std_Dev_W_fc1
tag: Std_Dev_W_fc1

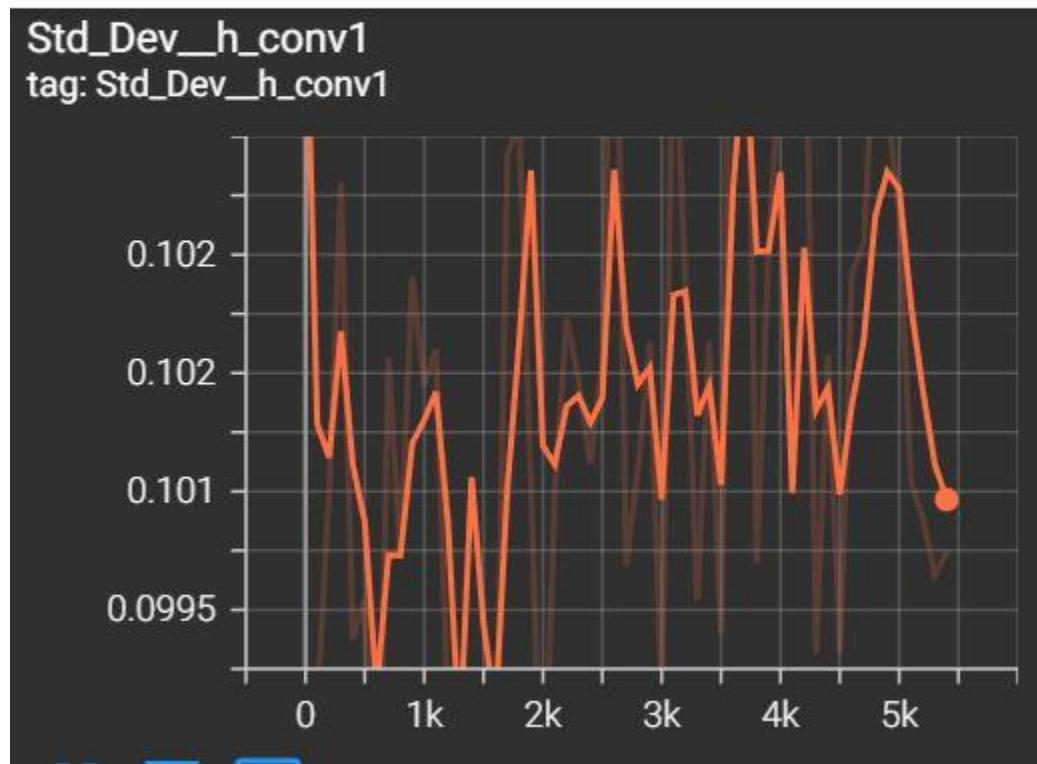
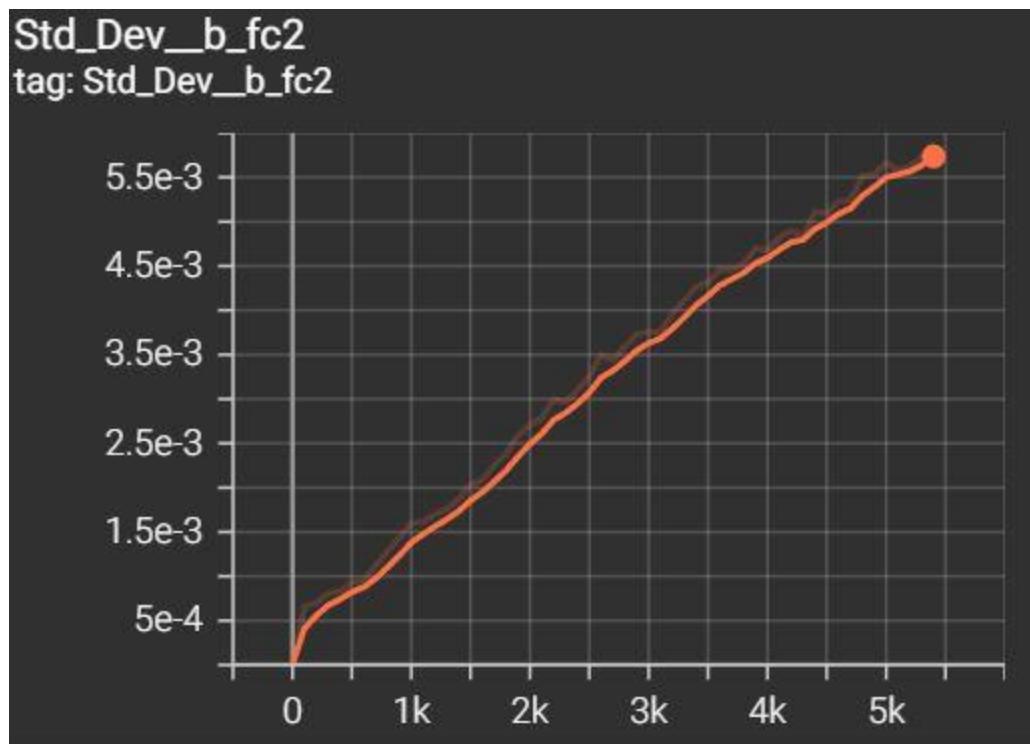


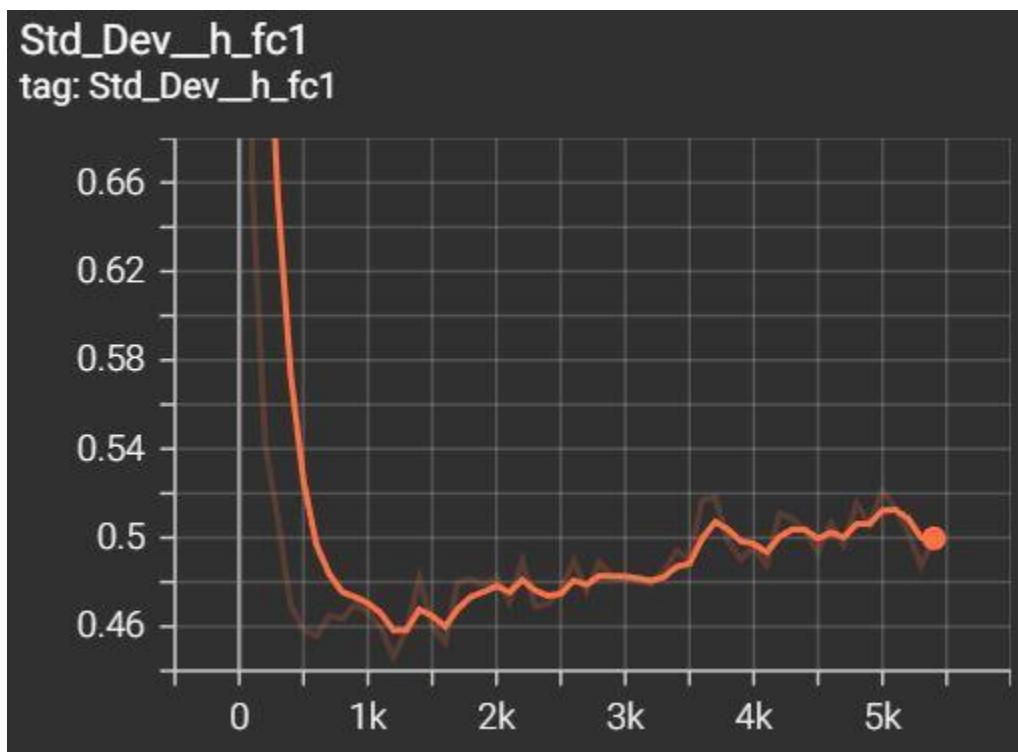
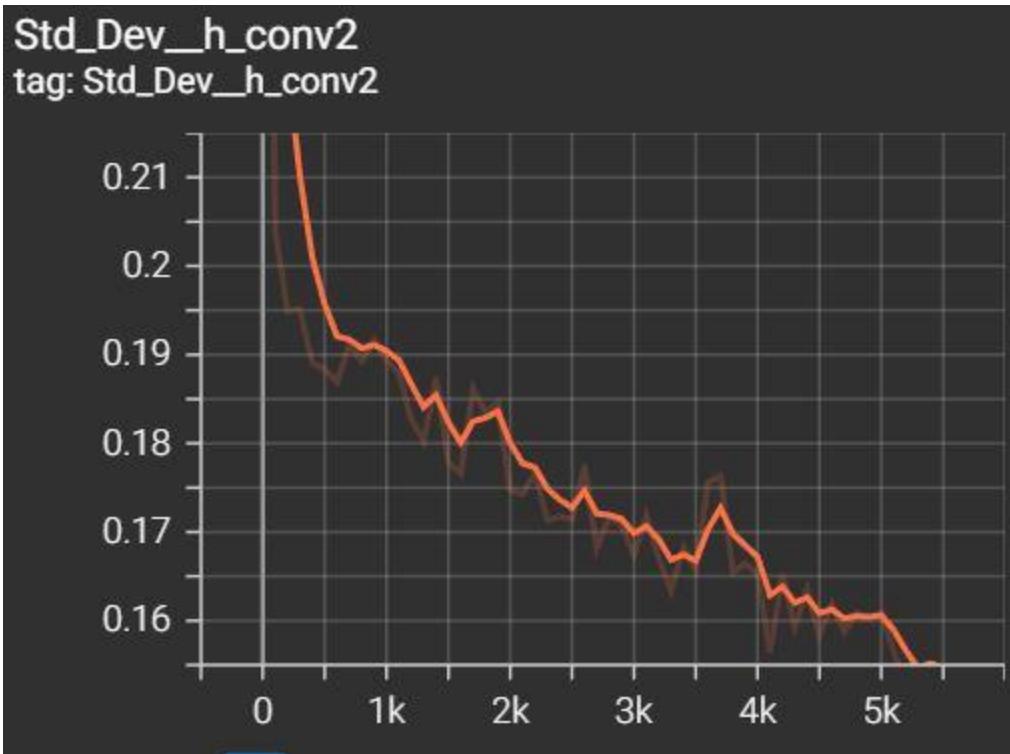
Std_Dev_W_fc2
tag: Std_Dev_W_fc2



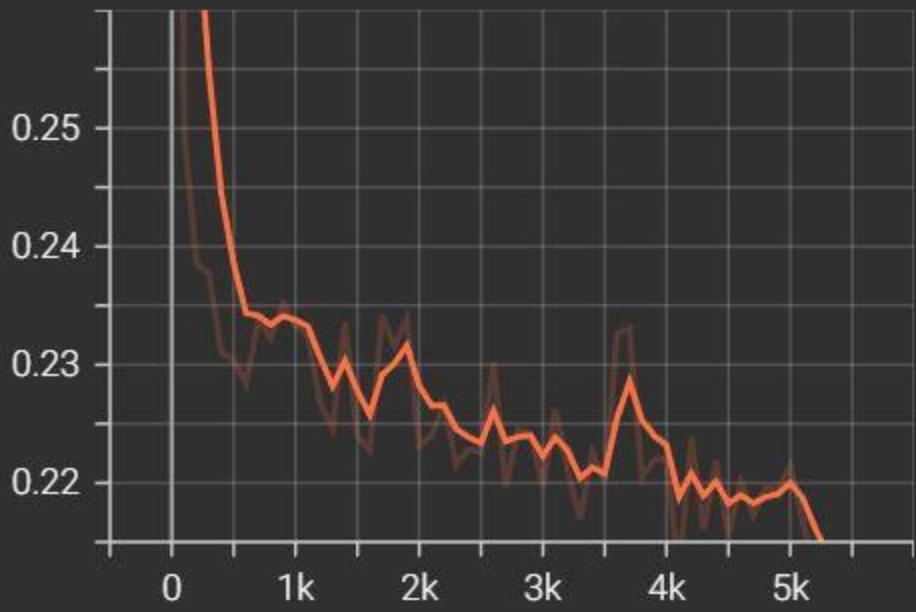




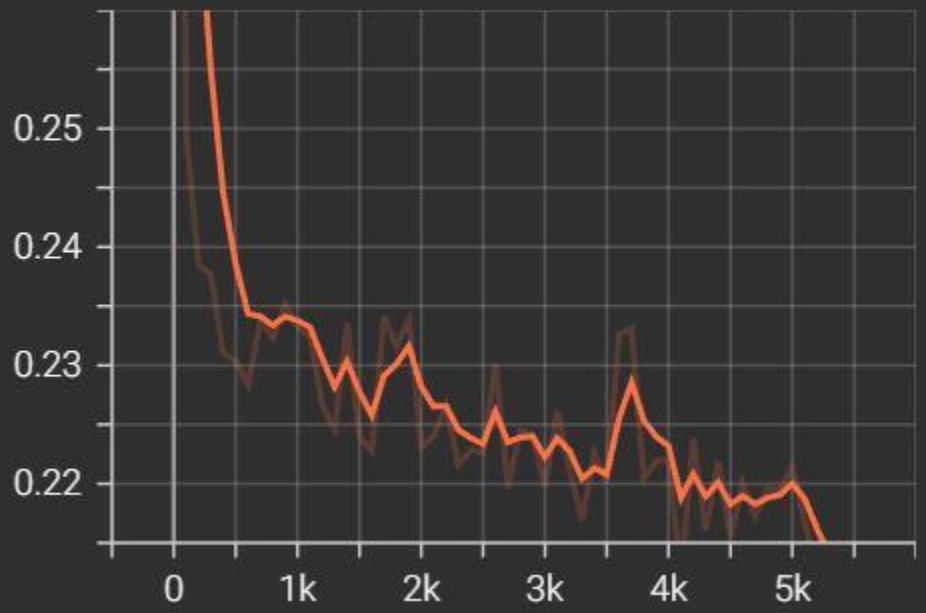




Std_Dev_h_pool2
tag: Std_Dev_h_pool2

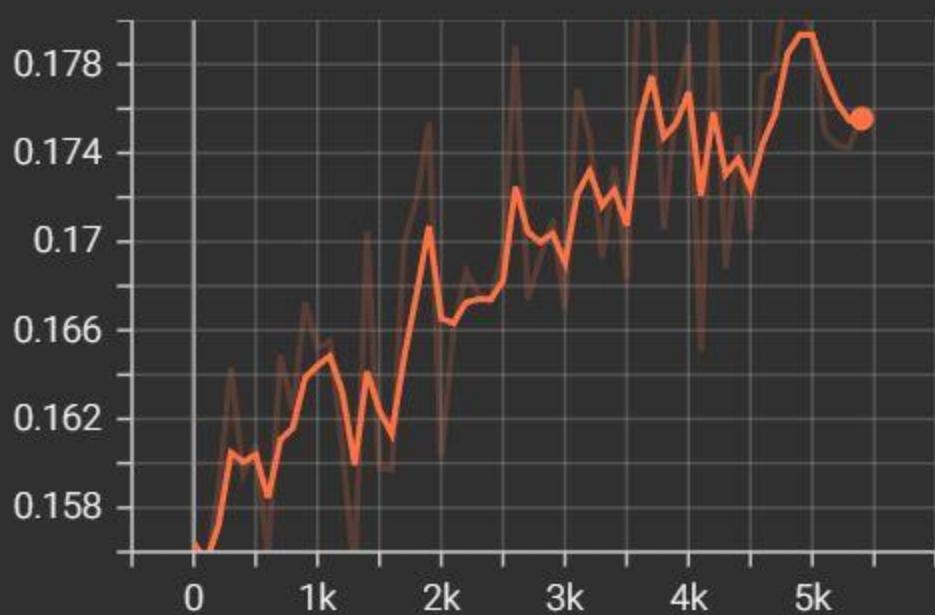


Std_Dev_h_pool2_flat
tag: Std_Dev_h_pool2_flat



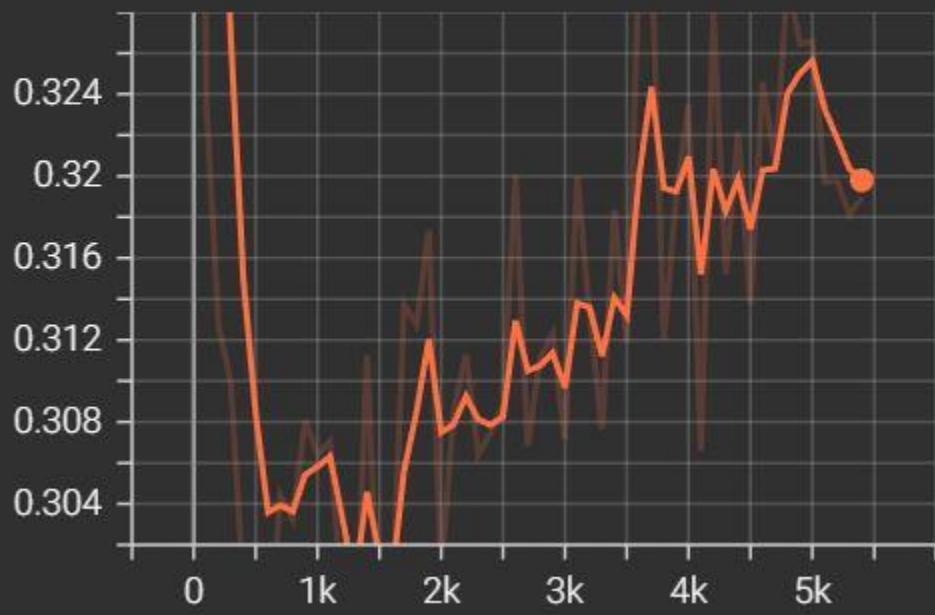
Std_Dev_input_1

tag: Std_Dev_input_1

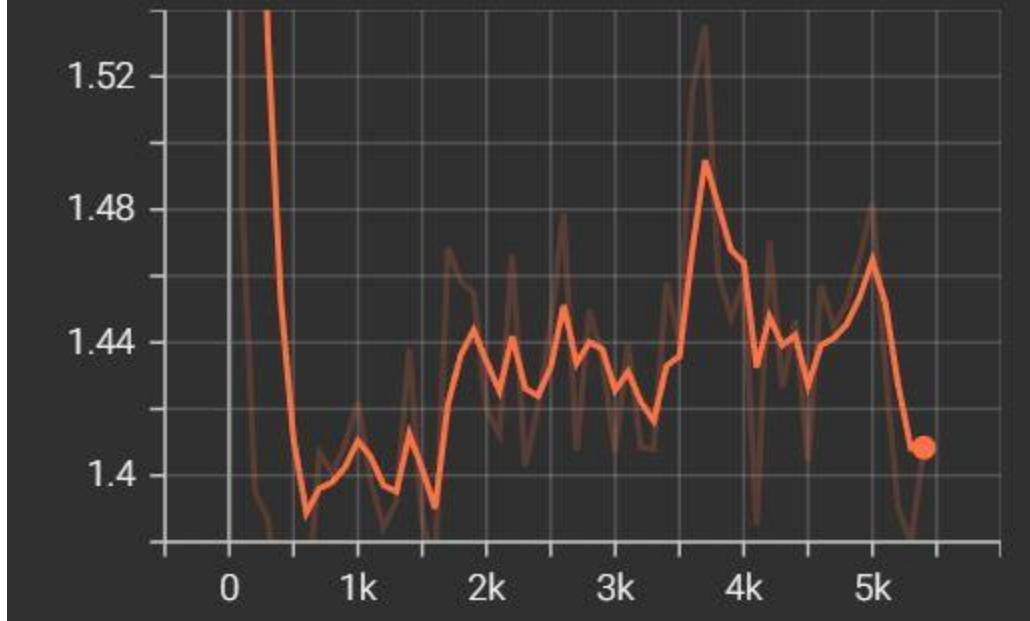


Std_Dev_input_2

tag: Std_Dev_input_2



Std_Dev_input_fc1
tag: Std_Dev_input_fc1

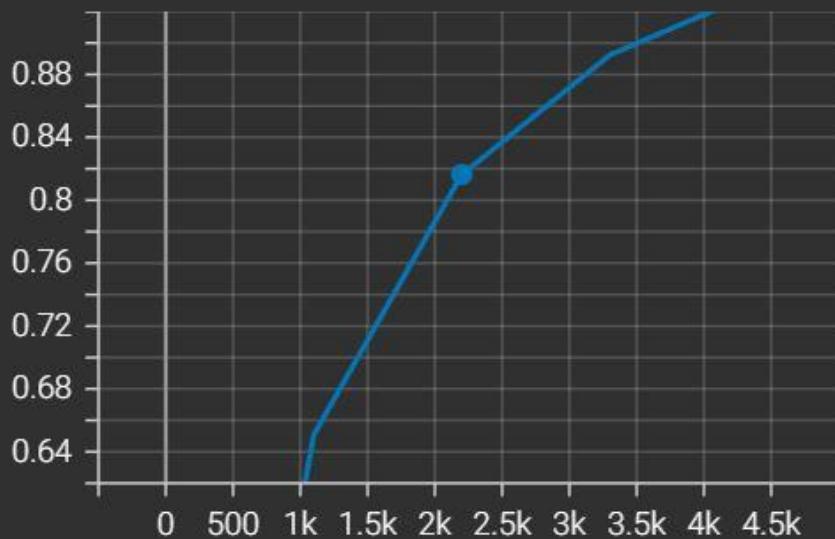


Std_Dev_y_conv_
tag: Std_Dev_y_conv_

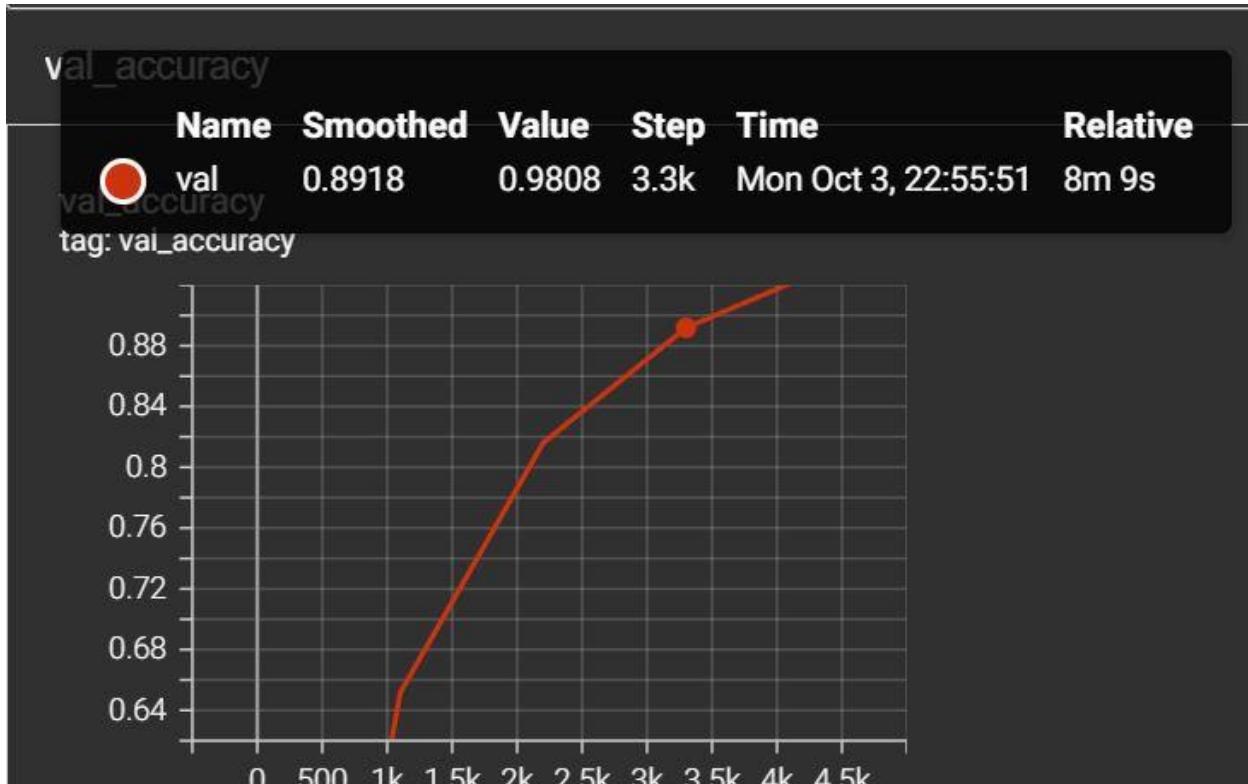


test_accuracy

tag: test_accuracy



	Name	Smoothed	Value	Step	Time	Relative
val_accuracy	test	0.8163	0.9752	2.2k	Mon Oct 3, 22:53:03	5m 24s



Problem 2 Part B:

- The approach for this part is to use the following code shown below to plot more information about the training results:

```
def variable_summaries(var):
    with tf.name_scope('summaries'):
        mean = tf.reduce_mean(var)
        tf.summary.scalar('mean', mean)
        with tf.name_scope('stddev'):
            stddev = tf.sqrt(tf.reduce_mean(tf.square(var - mean)))
            tf.summary.scalar('stddev', stddev)
            tf.summary.scalar('max', tf.reduce_max(var))
            tf.summary.scalar('min', tf.reduce_min(var))
            tf.summary.histogram('histogram', var)
```

Problem 2 Part C:

- The approach for this part is to use gradient descent instead of the Adam training algorithm this can be done by the following steps:

- Step 1 , convert ReLU to tanh function as shown in the code below:

```
# placeholders for input data and input labeles
x = tf.placeholder(tf.float32, [None, 784], name='x')
y_ = tf.placeholder(tf.float32, [None, 10], name='y_')

# Store Accuracies
val_accuracy_ = tf.placeholder(tf.float32, shape=())
test_accuracy_ = tf.placeholder(tf.float32, shape=())

# reshape the input image
x_image = tf.reshape(x, [-1, 28, 28, 1])

# first convolutional layer
W_conv1 = weight_variable([5, 5, 1, 32])
b_conv1 = bias_variable([32])
input_1 = conv2d(x_image, W_conv1) + b_conv1
h_conv1 = tf.nn.tanh(input_1)
h_pool1 = max_pool_2x2(h_conv1)

# second convolutional layer
W_conv2 = weight_variable([5, 5, 32, 64])
b_conv2 = bias_variable([64])
input_2 = conv2d(h_pool1, W_conv2) + b_conv2
h_conv2 = tf.nn.tanh(input_2)
h_pool2 = max_pool_2x2(h_conv2)

# densely connected layer
W_fc1 = weight_variable([7 * 7 * 64, 1024])
b_fc1 = bias_variable([1024])
h_pool2_flat = tf.reshape(h_pool2, [-1, 7 * 7 * 64])
input_fc1 = tf.matmul(h_pool2_flat, W_fc1) + b_fc1
h_fc1 = tf.nn.tanh(input_fc1)

# dropout
keep_prob = tf.placeholder(tf.float32)
h_fc1_drop = h_fc1_drop = tf.nn.dropout(h_fc1, keep_prob)

# softmax
W_fc2 = weight_variable([1024, 10])
b_fc2 = bias_variable([10])
y_conv = tf.nn.softmax(tf.matmul(h_fc1_drop, W_fc2) + b_fc2, name='y')
```

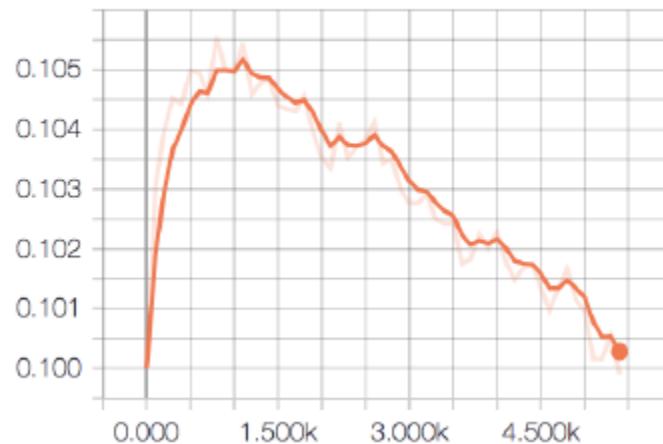
- Step 2, Change from Adam optimizer to Gradient Descent as shown in the code below:

```
# setup training
cross_entropy = tf.reduce_mean(-tf.reduce_sum(y_ * tf.log(y_conv),
reduction_indices=[1]))
#train_step = tf.train.AdamOptimizer(1e-4).minimize(cross_entropy)
train_step = tf.train.GradientDescentOptimizer(1e-4).minimize(cross_entropy)
```

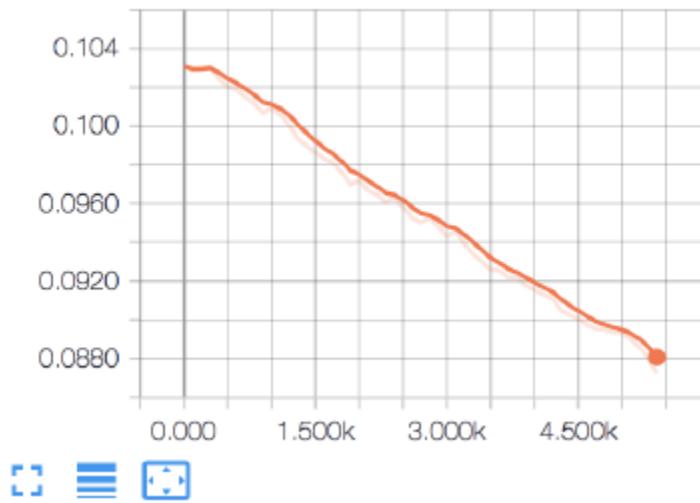
```
correct_prediction = tf.equal(tf.argmax(y_conv, 1), tf.argmax(y_, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32),
name='accuracy')
```

- The results are shown in the figures below:

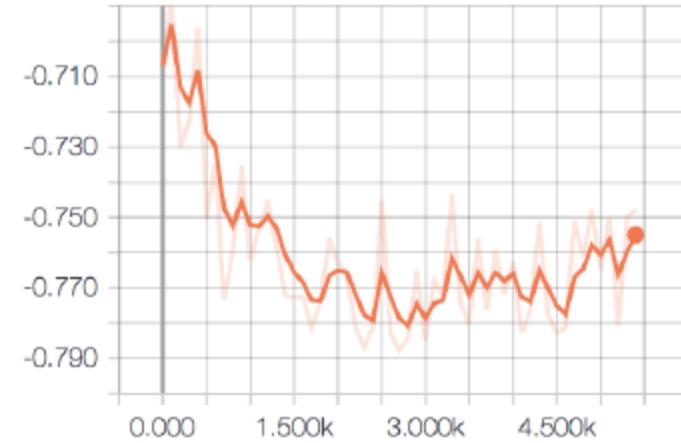
summaries_1/max_1



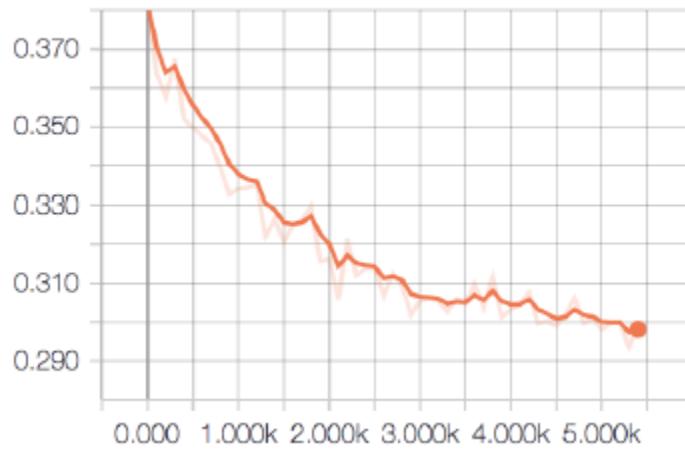
summaries_2/mean_1



summaries_3/min_1



summaries_6/stddev_1



REFERENCES:

- The code below is the entire code used for Problem 1 of the homework:

```
import numpy as np
from sklearn import datasets, linear_model
import matplotlib.pyplot as plt
from tensorflow.keras.utils import to_categorical
from scipy.special import factorial
def generate_data():
    """
    generate data
    :return: X: input data, y: given labels
    """
```

```

    """
    np.random.seed(0)
    X, y = datasets.make_moons(200, noise=0.20)
    return X, y

def plot_decision_boundary(pred_func, X, y):
    """
    plot the decision boundary
    :param pred_func: function used to predict the label
    :param X: input data
    :param y: given labels
    :return:
    """

    # Set min and max values and give it some padding
    x_min, x_max = X[:, 0].min() - .5, X[:, 0].max() + .5
    y_min, y_max = X[:, 1].min() - .5, X[:, 1].max() + .5
    h = 0.01
    # Generate a grid of points with distance h between them
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max,
    h))
    # Predict the function value for the whole grid
    Z = pred_func(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)
    # Plot the contour and training examples
    plt.contourf(xx, yy, Z, cmap=plt.cm.Spectral)
    plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.Spectral)
    plt.show()

#####
##### YOUR ASSSIGMENT STARTS HERE
##### FOLLOW THE INSTRUCTION BELOW TO BUILD AND TRAIN A 3-LAYER NEURAL NETWORK
#####

class NeuralNetwork(object):
    """
    This class builds and trains a neural network
    """
    def __init__(self, nn_input_dim, nn_hidden_dim, nn_output_dim,
                 actFun_type='tanh', reg_lambda=0.01, seed=0):
        """
        :param nn_input_dim: input dimension
        :param nn_hidden_dim: the number of hidden units
        :param nn_output_dim: output dimension
        :param actFun_type: type of activation function. 3 options: 'tanh',
        'sigmoid', 'relu'
        :param reg_lambda: regularization coefficient
        :param seed: random seed
        """
        self.nn_input_dim = nn_input_dim
        self.nn_hidden_dim = nn_hidden_dim
        self.nn_output_dim = nn_output_dim
        self.actFun_type = actFun_type
        self.reg_lambda = reg_lambda

        # initialize the weights and biases in the network
        np.random.seed(seed)
        self.W1 = np.random.randn(self.nn_input_dim, self.nn_hidden_dim) / np.sqrt(self.nn_input_dim)

```

```

        self.b1 = np.zeros((1, self.nn_hidden_dim))
        self.W2 = np.random.randn(self.nn_hidden_dim, self.nn_output_dim) /
        np.sqrt(self.nn_hidden_dim)
        self.b2 = np.zeros((1, self.nn_output_dim))

    def actFun(self, z, type):

        if type == 'tanh':
            z = (2/(1+np.exp(-2*z)))-1
            #return np.tanh(z)
            return z
        elif type == 'sigmoid':
            z = 1/(1+np.exp(-1*z))
        elif type == 'relu':
            return np.maximum(0, z)

        else:
            print("Incorrect function detected,\ \
                  please enter one of the following functions:\ \
                  ReLU, Sigmoid or Tanh")
        return z

    def diff_actFun(self, z, type):
        """
        diff_actFun compute the derivatives of the activation functions wrt
        the net input
        :param z: net input
        :param type: Tanh, Sigmoid, or ReLU
        :return: the derivatives of the activation functions wrt the net
        input
        """

        if type == 'tanh':
            z = (1/np.square(np.cosh(z))) #formula created using
            tanh^2+sech^2=1
            #return 1 - np.square(np.tanh(z))
            #f = (2/(1+np.exp(-2*z)))-1
            #z = 1 - np.square(f)
            return z
        elif type == 'sigmoid':
            z = np.exp(-z)/(1 + 2*np.exp(-z) + np.exp(-2*z))
        elif type == 'relu':
            return np.where(z > 0, 1, 0)

        else:
            print("Incorrect function detected,\ \
                  please enter one of the following functions:\ \
                  ReLU, Sigmoid or Tanh")
        return z

    def feedforward(self, X, actFun):
        """
        feedforward builds a 3-layer neural network and computes the two
        probabilities,
        one for class 0 and one for class 1
        :param X: input data
        :param actFun: activation function
        :return:
        """

        # YOU IMPLEMENT YOUR feedforward HERE

        self.z1 = np.dot(X,self.W1) + self.b1

```

```

        self.a1 = actFun(self.z1)
        self.z2 = np.dot(self.a1, self.W2) + self.b2

        #Version 1 of a2 = ^y = softmax(z2)
        #self.probs = np.exp(self.z2)/ np.sum(np.exp(self.z2),axis = 1,
        keepdims = True)

        #version 2 of a2 = ^y = softmax(z2)
        A = np.exp(self.z2 - np.max(self.z2))
        B = np.sum(A, axis = 1, keepdims = True)
        result = A / B
        self.probs = result

    return None

def calculate_loss(self, X, y):
    """
    calculate_loss compute the loss for prediction
    :param X: input data
    :param y: given labels
    :return: the loss for prediction
    """
    num_examples = len(X)
    self.feedforward(X, lambda x: self.actFun(x, type=self.actFun_type))

    # Calculating the loss
    yhat = self.probs #This is the same equation found feedforward
function
    data_loss = np.sum(-np.log(yhat[range(num_examples)], y)))

    # Add regularization term to loss (optional)
    data_loss += self.reg_lambda / 2 * (np.sum(np.square(self.W1)) +
    np.sum(np.square(self.W2))))
    return (1. / num_examples) * data_loss

def predict(self, X):
    """
    predict infers the label of a given data point X
    :param X: input data
    :return: label inferred
    """
    self.feedforward(X, lambda x: self.actFun(x, type=self.actFun_type))
    return np.argmax(self.probs, axis=1)

def backprop(self, X, y):
    """
    backprop run backpropagation to compute the gradients used to update
    the parameters in the backward step
    :param X: input data
    :param y: given labels
    :return: dL/dW1, dL/b1, dL/dW2, dL/db2
    """

    # IMPLEMENT YOUR BACKPROP HERE

    num_examples = len(X)
    delta3 = self.probs
    delta3[range(num_examples), y] -= 1
    dW2 = (self.a1.T).dot(delta3)
    db2 = np.sum(delta3, axis=0, keepdims=True)

```

```

        delta2 = delta3.dot(self.W2.T) * self.diff_actFun(self.z1,
self.actFun_type)
        dW1 = np.dot(X.T, delta2)
        db1 = np.sum(delta2, axis=0)

    return dW1, dW2, db1, db2

def fit_model(self, X, y, epsilon=0.01, num_passes=20000,
print_loss=True):
    """
    fit_model uses backpropagation to train the network
    :param X: input data
    :param y: given labels
    :param num_passes: the number of times that the algorithm runs
through the whole dataset
    :param print_loss: print the loss or not
    :return:
    """
    # Gradient descent.
    for i in range(0, num_passes):
        # Forward propagation
        self.feedforward(X, lambda x: self.actFun(x,
type=self.actFun_type))
        # Backpropagation
        dW1, dW2, db1, db2 = self.backprop(X, y)

        # Add derivatives of regularization terms (b1 and b2 don't have
regularization terms)
        dW2 += self.reg_lambda * self.W2
        dW1 += self.reg_lambda * self.W1

        # Gradient descent parameter update
        self.W1 += -epsilon * dW1
        self.b1 += -epsilon * db1
        self.W2 += -epsilon * dW2
        self.b2 += -epsilon * db2

        # Optionally print the loss.
        # This is expensive because it uses the whole dataset, so we
don't want to do it too often.
        if print_loss and i % 1000 == 0:
            print("Loss after iteration %i: %f" % (i,
self.calculate_loss(X, y)))

    def visualize_decision_boundary(self, X, y):
        """
        visualize_decision_boundary plot the decision boundary created by
the trained network
        :param X: input data
        :param y: given labels
        :return:
        """
        plot_decision_boundary(lambda x: self.predict(x), X, y)

    # # generate and visualize Make-Moons dataset
    X, y = generate_data()
    #plt.scatter(X[:, 0], X[:, 1], s=40, c=y, cmap=plt.cm.Spectral)
    #plt.show()

    model = NeuralNetwork(nn_input_dim=2, nn_hidden_dim=3, nn_output_dim=2,
actFun_type='relu')
    model.fit_model(X, y)

```

```
model.visualize_decision_boundary(X,y)
```

- The code below is the **DeepNeuralNetwork** class code used for Problem 1 of the homework:

```
import numpy as np
from sklearn import datasets, linear_model
import matplotlib.pyplot as plt
from tensorflow.keras.utils import to_categorical
from scipy.special import factorial
def generate_data():
    """
    generate data
    :return: X: input data, y: given labels
    """
    np.random.seed(0)
    X, y = datasets.make_moons(200, noise=0.20)
    return X, y

def plot_decision_boundary(pred_func, X, y):
    """
    plot the decision boundary
    :param pred_func: function used to predict the label
    :param X: input data
    :param y: given labels
    :return:
    """
    # Set min and max values and give it some padding
    x_min, x_max = X[:, 0].min() - .5, X[:, 0].max() + .5
    y_min, y_max = X[:, 1].min() - .5, X[:, 1].max() + .5
    h = 0.01
    # Generate a grid of points with distance h between them
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))
    # Predict the function value for the whole grid
    Z = pred_func(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)
    # Plot the contour and training examples
    plt.contourf(xx, yy, Z, cmap=plt.cm.Spectral)
    plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.Spectral)
    plt.show()

#####
# YOUR ASSSIGMENT STARTS HERE
# FOLLOW THE INSTRUCTION BELOW TO BUILD AND TRAIN A 3-LAYER NEURAL NETWORK
#####
class NeuralNetwork(object):
    """
    This class builds and trains a neural network
    """
```

```

    def __init__(self, nn_input_dim, nn_hidden_dim , nn_output_dim,
actFun_type='tanh', reg_lambda=0.01, seed=0):
    """
        :param nn_input_dim: input dimension
        :param nn_hidden_dim: the number of hidden units
        :param nn_output_dim: output dimension
        :param actFun_type: type of activation function. 3 options: 'tanh',
'sigmoid', 'relu'
        :param reg_lambda: regularization coefficient
        :param seed: random seed
    """
    self.nn_input_dim = nn_input_dim
    self.nn_hidden_dim = nn_hidden_dim
    self.nn_output_dim = nn_output_dim
    self.actFun_type = actFun_type
    self.reg_lambda = reg_lambda

    # initialize the weights and biases in the network
    np.random.seed(seed)
    self.W1 = np.random.randn(self.nn_input_dim, self.nn_hidden_dim) / 
np.sqrt(self.nn_input_dim)
    self.b1 = np.zeros((1, self.nn_hidden_dim))
    self.W2 = np.random.randn(self.nn_hidden_dim, self.nn_output_dim) / 
np.sqrt(self.nn_hidden_dim)
    self.b2 = np.zeros((1, self.nn_output_dim))

def actFun(self, z, type):

    if type == 'tanh':
        z = (2/(1+np.exp(-2*z)))-1
        #return np.tanh(z)
        return z
    elif type == 'sigmoid':
        z = 1/(1+np.exp(-1*z))
    elif type == 'relu':
        return np.maximum(0, z)

    else:
        print("Incorrect function detected,\n"
              "please enter one of the following functions:\n"
              "ReLU, Sigmoid or Tanh")
    return z

def diff_actFun(self, z, type):
    """
        diff_actFun compute the derivatives of the activation functions wrt
the net input
        :param z: net input
        :param type: Tanh, Sigmoid, or ReLU
        :return: the derivatives of the activation functions wrt the net
input
    """

    if type == 'tanh':
        z = (1/np.square(np.cosh(z))) #formula created using
tanh^2+sech^2=1
        #return 1 - np.square(np.tanh(z))
        #f = (2/(1+np.exp(-2*z)))-1
        #z = 1 - np.square(f)
        return z
    elif type == 'sigmoid':
        z = np.exp(-z)/(1 + 2*np.exp(-z) + np.exp(-2*z))
    elif type == 'relu':

```

```

        return np.where(z > 0, 1, 0)

    else:
        print("Incorrect function detected,\n"
              "please enter one of the following functions:\\"
              "ReLU, Sigmoid or Tanh")
    return z

def feedforward(self, X, actFun):
    """
    feedforward builds a 3-layer neural network and computes the two
    probabilities,
    one for class 0 and one for class 1
    :param X: input data
    :param actFun: activation function
    :return:
    """

    # YOU IMPLEMENT YOUR feedforward HERE

    self.z1 = np.dot(X, self.W1) + self.b1
    self.a1 = actFun(self.z1)
    self.z2 = np.dot(self.a1, self.W2) + self.b2

    #Version 1 of a2 = ^y = softmax(z2)
    #self.probs = np.exp(self.z2)/ np.sum(np.exp(self.z2),axis = 1,
    keepdims = True)

    #version 2 of a2 = ^y = softmax(z2)
    A = np.exp(self.z2 - np.max(self.z2))
    B = np.sum(A, axis = 1, keepdims = True)
    result = A / B
    self.probs = result

    return None

def calculate_loss(self, X, y):
    """
    calculate_loss compute the loss for prediction
    :param X: input data
    :param y: given labels
    :return: the loss for prediction
    """

    num_examples = len(X)
    self.feedforward(X, lambda x: self.actFun(x, type=self.actFun_type))

    # Calculating the loss
    yhat = self.probs #This is the same equation found feedforward
function
    data_loss = np.sum(-np.log(yhat[range(num_examples)], y)))

    # Add regularization term to loss (optional)
    data_loss += self.reg_lambda / 2 * (np.sum(np.square(self.W1)) +
    np.sum(np.square(self.W2)))
    return (1. / num_examples) * data_loss

def predict(self, X):
    """
    predict infers the label of a given data point X
    :param X: input data
    """

```

```

:return: label inferred
'''
self.feedforward(X, lambda x: self.actFun(x, type=self.actFun_type))
return np.argmax(self.probs, axis=1)

def backprop(self, X, y):
    '''
        backprop run backpropagation to compute the gradients used to update
        the parameters in the backward step
        :param X: input data
        :param y: given labels
        :return: dL/dW1, dL/b1, dL/dW2, dL/db2
    '''

    # IMPLEMENT YOUR BACKPROP HERE

    num_examples = len(X)
    delta3 = self.probs
    delta3[range(num_examples), y] -= 1
    dW2 = (self.a1.T).dot(delta3)
    db2 = np.sum(delta3, axis=0, keepdims=True)
    delta2 = delta3.dot(self.W2.T) * self.diff_actFun(self.z1,
    self.actFun_type)
    dW1 = np.dot(X.T, delta2)
    db1 = np.sum(delta2, axis=0)

    return dW1, dW2, db1, db2

def fit_model(self, X, y, epsilon=0.01, num_passes=20000,
print_loss=True):
    '''
        fit_model uses backpropagation to train the network
        :param X: input data
        :param y: given labels
        :param num_passes: the number of times that the algorithm runs
        through the whole dataset
        :param print_loss: print the loss or not
        :return:
    '''

    # Gradient descent.
    for i in range(0, num_passes):
        # Forward propagation
        self.feedforward(X, lambda x: self.actFun(x,
        type=self.actFun_type))
        # Backpropagation
        dW1, dW2, db1, db2 = self.backprop(X, y)

        # Add derivatives of regularization terms (b1 and b2 don't have
        regularization terms)
        dW2 += self.reg_lambda * self.W2
        dW1 += self.reg_lambda * self.W1

        # Gradient descent parameter update
        self.W1 += -epsilon * dW1
        self.b1 += -epsilon * db1
        self.W2 += -epsilon * dW2
        self.b2 += -epsilon * db2

        # Optionally print the loss.
        # This is expensive because it uses the whole dataset, so we
        don't want to do it too often.
        if print_loss and i % 1000 == 0:

```

```

        print("Loss after iteration %i: %f" % (i,
self.calculate_loss(X, y)))

    def visualize_decision_boundary(self, X, y):
        """
        visualize_decision_boundary plot the decision boundary created by
the trained network
        :param X: input data
        :param y: given labels
        :return:
        """
        plot_decision_boundary(lambda x: self.predict(x), X, y)

class DeepNeuralNetwork(object):
    """
    This class builds and trains a deep neural network
    """

    def __init__(self, nn_dims, actFun_type='tanh', reg_lambda=0.01,
seed=0):

        self.nn_dims = nn_dims
        self.actFun_type = actFun_type
        self.reg_lambda = reg_lambda

        # initialize the weights and biases in the network
        self.W = []
        self.b = []

        np.random.seed(seed)
        for i in range(len(nn_dims)-1):
            self.W.append(np.random.randn(self.nn_dims[i],
self.nn_dims[i+1])
                           / np.sqrt(self.nn_dims[i]))
            self.b.append(np.zeros((1, self.nn_dims[i+1])))

    def actFun(self, z, type):

        if type == 'tanh':
            z = (2/(1+np.exp(-2*z)))-1
            #return np.tanh(z)
            return z
        elif type == 'sigmoid':
            z = 1/(1+np.exp(-1*z))
        elif type == 'relu':
            return np.maximum(0, z)

        else:
            print("Incorrect function detected,\ \
                  please enter one of the following functions:\ \
                  ReLU, Sigmoid or Tanh")
        return z

    def diff_actFun(self, z, type):

        if type == 'tanh':
            z = (1/np.square(np.cosh(z))) #formula created using
tanh^2+sech^2=1
            #return 1 - np.square(np.tanh(z))

```

```

#f = (2/(1+np.exp(-2*z)))-1
#z = 1 - np.square(f)
return z
elif type == 'sigmoid':
    z = np.exp(-z)/(1 + 2*np.exp(-z) + np.exp(-2*z))
elif type == 'relu':
    return np.where(z > 0, 1, 0)

else:
    print("Incorrect function detected,\\
          please enter one of the following functions:\\
          ReLU, Sigmoid or Tanh")
return z

def feedforward(self, X, actFun):

    self.a = []
    self.z = []
    for i in range(len(self.W)):
        if i == 0:
            self.z.append(np.dot(X, self.W[i]) + self.b[i])
        else:
            self.z.append(np.dot(self.a[i-1], self.W[i]) + self.b[i])
    if i != len(self.W) - 1:
        self.a.append(actFun(self.z[i]))
    C = np.exp(self.z[len(self.z)-1])
    self.probs = C / np.sum(C, axis=1, keepdims=True)
    return None

def calculate_loss(self, X, y):

    num_examples = len(X)
    self.feedforward(X, lambda x: self.actFun(x, type=self.actFun_type))
    # Calculating the loss

    probs = np.exp(self.z[len(self.z)-1]) / \
           np.sum(np.exp(self.z[len(self.z)-1]), axis=1, keepdims=True)
    loss = -np.log(probs[range(num_examples), y])
    final_data_loss = np.sum(loss)

    # Add regularization term
    W_sum = 0
    for i in len(self.W):
        W_sum += np.sum(np.square(self.W[i]))
    final_data_loss += self.reg_lambda / 2 * W_sum
    return (1. / num_examples) * final_data_loss

def predict(self, X):

    self.feedforward(X, lambda x: self.actFun(x, type=self.actFun_type))
    return np.argmax(self.probs, axis=1)

def backprop(self, X, y):

    num_examples = len(X)
    delta = self.probs
    delta[range(num_examples), y] -= 1

    db = []
    dW = []
    for num in range(len(self.z)):
        i = len(self.z) - num - 1

```

```

        if i != 0:
            dW.insert(0, np.dot(self.a[i - 1].T, delta))
            db.insert(0, np.sum(delta, axis=0, keepdims=True))
            delta = np.dot(delta, self.W[i].T) * \
                self.diff_actFun(self.z[i-1], type=self.actFun_type)
        else:
            db.insert(0, np.sum(delta, axis=0, keepdims=False))
            dW.insert(0, np.dot(X.T, delta))

    return dW, db

def fit_model(self, X, y, epsilon=0.01, num_passes=20000,
print_loss=True):

    # Gradient descent.
    for i in range(0, num_passes):
        # Forward propagation
        self.feedforward(X, lambda x: self.actFun(x,
type=self.actFun_type))
        # Backpropagation
        dW, db = self.backprop(X, y)

        # Add regularization terms
        for i in range(len(dW)):
            dW[i] += self.reg_lambda * self.W[i]

        # Gradient descent parameter update
        for i in range(len(self.W)):
            self.W[i] += -epsilon * dW[i]
            self.b[i] += -epsilon * db[i]

        # print the loss.
        if print_loss and i % 1000 == 0:
            print("Loss after iteration %i: %f" % (i,
self.calculate_loss(X, y)))

    def visualize_decision_boundary(self, X, y):
        plot_decision_boundary(lambda x: self.predict(x), X, y)

def main():

    # # generate and visualize Make-Moons dataset
    #X, y = generate_data()
    # # generate and visualize Make-Circles dataset
    X, y = datasets.make_circles(n_samples=100, shuffle=True, noise=None,
random_state=None, factor=0.8)

    #Train the Neural Network Model
    model = DeepNeuralNetwork(nn_dims=[2,10,8,8,5], actFun_type='sigmoid')
    model.fit_model(X,y)
    model.visualize_decision_boundary(X,y)

if __name__ == "__main__":
    main()

```

- The code below is the entire code used for **problem 2** of the homework:

```

import os
import time

import tensorflow as tf

if (tf.__version__.split('.')[0] == '2'):
    import tensorflow.compat.v1 as tf
    tf.disable_v2_behavior()

# Load MNIST dataset
# import input_data

import input_data
mnist = input_data.read_data_sets('MNIST_data', one_hot=True)

sess = tf.InteractiveSession()

def weight_variable(shape):
    """
    Initialize weights
    :param shape: shape of weights, e.g. [w, h ,Cin, Cout] where
    w: width of the filters
    h: height of the filters
    Cin: the number of the channels of the filters
    Cout: the number of filters
    :return: a tensor variable for weights with initial values
    """

    # IMPLEMENT YOUR WEIGHT_VARIABLE HERE
    initial = tf.truncated_normal(shape, stddev=0.1)
    W = tf.Variable(initial)
    return W

def bias_variable(shape):
    """
    Initialize biases
    :param shape: shape of biases, e.g. [Cout] where
    Cout: the number of filters
    :return: a tensor variable for biases with initial values
    """

    # IMPLEMENT YOUR BIAS_VARIABLE HERE
    initial = tf.constant(0.1, shape=shape)
    b = tf.Variable(initial)
    return b

def conv2d(x, W):
    """
    Perform 2-D convolution
    :param x: input tensor of size [N, W, H, Cin] where
    N: the number of images
    W: width of images
    H: height of images
    Cin: the number of channels of images
    :param W: weight tensor [w, h, Cin, Cout]
    w: width of the filters
    """

```

```

    h: height of the filters
    Cin: the number of the channels of the filters = the number of channels
of images
    Cout: the number of filters
    :return: a tensor of features extracted by the filters, a.k.a. the
results after convolution
    '''

    # IMPLEMENT YOUR CONV2D HERE
    h_conv = tf.nn.conv2d(x, W, strides=[1, 1, 1, 1], padding='SAME')

    return h_conv


def max_pool_2x2(x):
    '''
    Perform non-overlapping 2-D maxpooling on 2x2 regions in the input data
    :param x: input data
    :return: the results of maxpooling (max-marginalized + downsampling)
    '''

    # IMPLEMENT YOUR MAX POOL 2X2 HERE
    h_max = tf.nn.max_pool(x, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1],
padding='SAME')

    return h_max


def stats_summary(var_name, var):
    ##monitor the statistics (min, max, mean, standard deviation, histogram)
    tf.summary.scalar("Min_" + var_name, tf.reduce_min(var))
    tf.summary.scalar("Max_" + var_name, tf.reduce_max(var))
    tf.summary.scalar("Mean_" + var_name, tf.reduce_mean(var))
    tf.summary.scalar("Std Dev_" + var_name, tf.math.reduce_std(var))
    tf.summary.histogram("Histo_" + var_name, var)


def main():
    # Specify training parameters
    result_dir = './results/' # directory where the results from the
training are saved
    max_step = 5500 # the maximum iterations. After max_step iterations,
the training will stop no matter what
    test_dir = result_dir + 'test/'
    val_dir = result_dir + 'val/'

    start_time = time.time() # start timing

    # FILL IN THE CODE BELOW TO BUILD YOUR NETWORK

    # placeholders for input data and input labeles
    x = tf.placeholder(tf.float32, [None, 784], name='x')
    y_ = tf.placeholder(tf.float32, [None, 10], name='y_')

    # Store Accuracies
    val_accuracy_ = tf.placeholder(tf.float32, shape=())
    test_accuracy_ = tf.placeholder(tf.float32, shape=())

    # reshape the input image
    x_image = tf.reshape(x, [-1, 28, 28, 1])

    # first convolutional layer
    W_conv1 = weight_variable([5, 5, 1, 32])

```

```

b_conv1 = bias_variable([32])
input_1 = conv2d(x_image, W_conv1) + b_conv1
h_conv1 = tf.nn.relu(input_1)
h_pool1 = max_pool_2x2(h_conv1)

# second convolutional layer
W_conv2 = weight_variable([5, 5, 32, 64])
b_conv2 = bias_variable([64])
input_2 = conv2d(h_pool1, W_conv2) + b_conv2
h_conv2 = tf.nn.relu(input_2)
h_pool2 = max_pool_2x2(h_conv2)

# densely connected layer
W_fc1 = weight_variable([7 * 7 * 64, 1024])
b_fc1 = bias_variable([1024])
h_pool2_flat = tf.reshape(h_pool2, [-1, 7 * 7 * 64])
input_fc1 = tf.matmul(h_pool2_flat, W_fc1) + b_fc1
h_fc1 = tf.nn.relu(input_fc1)

# dropout
keep_prob = tf.placeholder(tf.float32)
h_fc1_drop = h_fc1_drop = tf.nn.dropout(h_fc1, keep_prob)

# softmax
W_fc2 = weight_variable([1024, 10])
b_fc2 = bias_variable([10])
y_conv = tf.nn.softmax(tf.matmul(h_fc1_drop, W_fc2) + b_fc2, name='y')

# FILL IN THE FOLLOWING CODE TO SET UP THE TRAINING

# setup training
cross_entropy = tf.reduce_mean(-tf.reduce_sum(y_ * tf.log(y_conv),
reduction_indices=[1]))
train_step = tf.train.AdamOptimizer(1e-4).minimize(cross_entropy)
correct_prediction = tf.equal(tf.argmax(y_conv, 1), tf.argmax(y_, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32),
name='accuracy')

# Add a scalar summary for the snapshot loss.
tf.summary.scalar(cross_entropy.op.name, cross_entropy)
# Build the summary operation based on the TF collection of Summaries.
# Layer-1
stats_summary("W_conv1", W_conv1)
stats_summary("b_conv1", b_conv1)
stats_summary("input_1", input_1)
stats_summary("h_conv1", h_conv1)
stats_summary("h_pool1", h_pool1)
# Layer-2
stats_summary("W_conv2", W_conv2)
stats_summary("b_conv2", b_conv2)
stats_summary("input_2", input_2)
stats_summary("h_conv2", h_conv2)
stats_summary("h_pool2", h_pool2)
# densely connected layer
stats_summary("W_fc1", W_fc1)
stats_summary("b_fc1", b_fc1)
stats_summary("h_pool2_flat", h_pool2_flat)
stats_summary("input_fc1", input_fc1)
stats_summary("h_fc1", h_fc1)
# Output Layer - Softmax
stats_summary("W_fc2", W_fc2)
stats_summary("b_fc2", b_fc2)
stats_summary("y_conv ", y_conv)

```

```

summary_op = tf.summary.merge_all()
summary_op_test = tf.summary.scalar('test_accuracy', test_accuracy_)
summary_op_val = tf.summary.scalar('val_accuracy', val_accuracy_)

# Add the variable initializer Op.
init = tf.initialize_all_variables()

# Create a saver for writing training checkpoints.
saver = tf.train.Saver()

# Instantiate a SummaryWriter to output summaries and the Graph.
summary_writer = tf.summary.FileWriter(result_dir, sess.graph)
# Val test
summary_writer_test = tf.summary.FileWriter(test_dir, sess.graph)
summary_writer_val = tf.summary.FileWriter(val_dir, sess.graph)

# Run the Op to initialize the variables.
sess.run(init)

# run the training
for i in range(max_step):
    batch = mnist.train.next_batch(50) # make the data batch, which is
    # used in the training iteration.
    # the batch size is 50
    if i % 100 == 0:
        # output the training accuracy every 100 iterations
        train_accuracy = accuracy.eval(feed_dict={
            x: batch[0], y_: batch[1], keep_prob: 1.0})
        print("step %d, training accuracy %g" % (i, train_accuracy))

        # Update the events file which is used to monitor the training
        # (in this case,
        # only the training loss is monitored)
        summary_str = sess.run(summary_op, feed_dict={x: batch[0], y_:
batch[1], keep_prob: 0.5})
        summary_writer.add_summary(summary_str, i)
        summary_writer.flush()

    # save the checkpoints every 1100 iterations
    if i % 1100 == 0 or i == max_step:
        checkpoint_file = os.path.join(result_dir, 'checkpoint')
        saver.save(sess, checkpoint_file, global_step=i)

        feed_dict_test = {x: mnist.test.images, y_: mnist.test.labels,
keep_prob: 1.0}
        test_accuracy = accuracy.eval(feed_dict=feed_dict_test)
        print('step %d, test accuracy %g' % (i, test_accuracy))
        summary_str_test = sess.run(summary_op_test,
feed_dict={test_accuracy : test_accuracy})
        summary_writer_test.add_summary(summary_str_test, i)
        summary_writer_test.flush()

        feed_dict_val = {x: mnist.validation.images, y_:
mnist.validation.labels, keep_prob: 1.0}
        val_accuracy = accuracy.eval(feed_dict=feed_dict_val)
        print('step %d, validation accuracy %g' % (i, val_accuracy))
        summary_str_val = sess.run(summary_op_val,
feed_dict={val_accuracy_ : val_accuracy})
        summary_writer_val.add_summary(summary_str_val, i)
        summary_writer_val.flush()

```

```
    train_step.run(feed_dict={x: batch[0], y_: batch[1], keep_prob: 0.5}) # run one train_step

    # print validation error
    print('validation accuracy %g' % accuracy.eval(feed_dict={
        x: mnist.validation.images, y_: mnist.validation.labels, keep_prob: 1.0}))

    # print test error
    print("test accuracy %g" % accuracy.eval(feed_dict={
        x: mnist.test.images, y_: mnist.test.labels, keep_prob: 1.0}))

stop_time = time.time()
print('The training takes %f second to finish' % (stop_time - start_time))

if __name__ == "__main__":
    main()
```